

Image Filtering (Frequency Domain)

Dr. Sander Ali Khowaja,

Assistant Professor, Department of Telecommunication Engineering
Faculty of Engineering and Technology, University of Sindh, Pakistan

Senior Member, IEEE – Member, ACM

<https://sander-ali.github.io>

Computer Vision & Image Processing



Some AI related news

Augmented reality headset enables users to see hidden objects

The device could help workers locate objects for fulfilling e-commerce orders or identify parts for assembling products.

 Watch Video

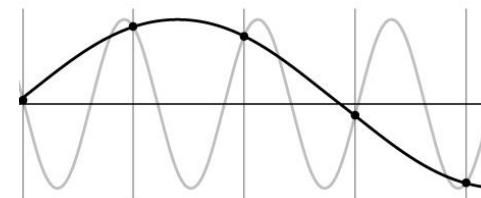
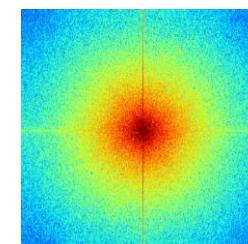
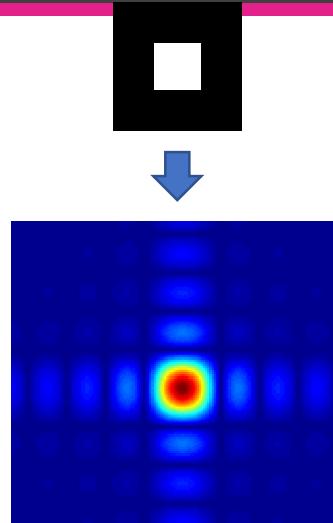
Adam Zewe | MIT News Office

February 27, 2023

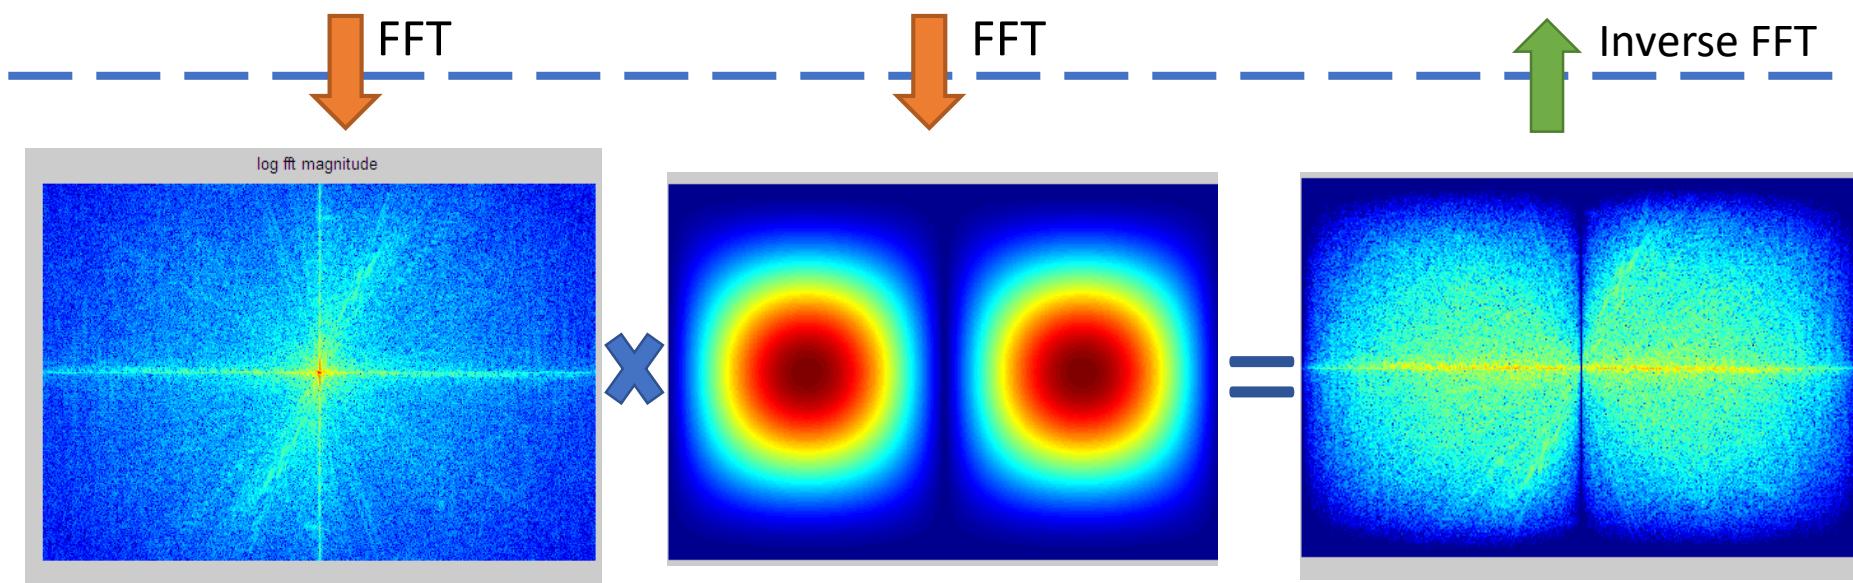
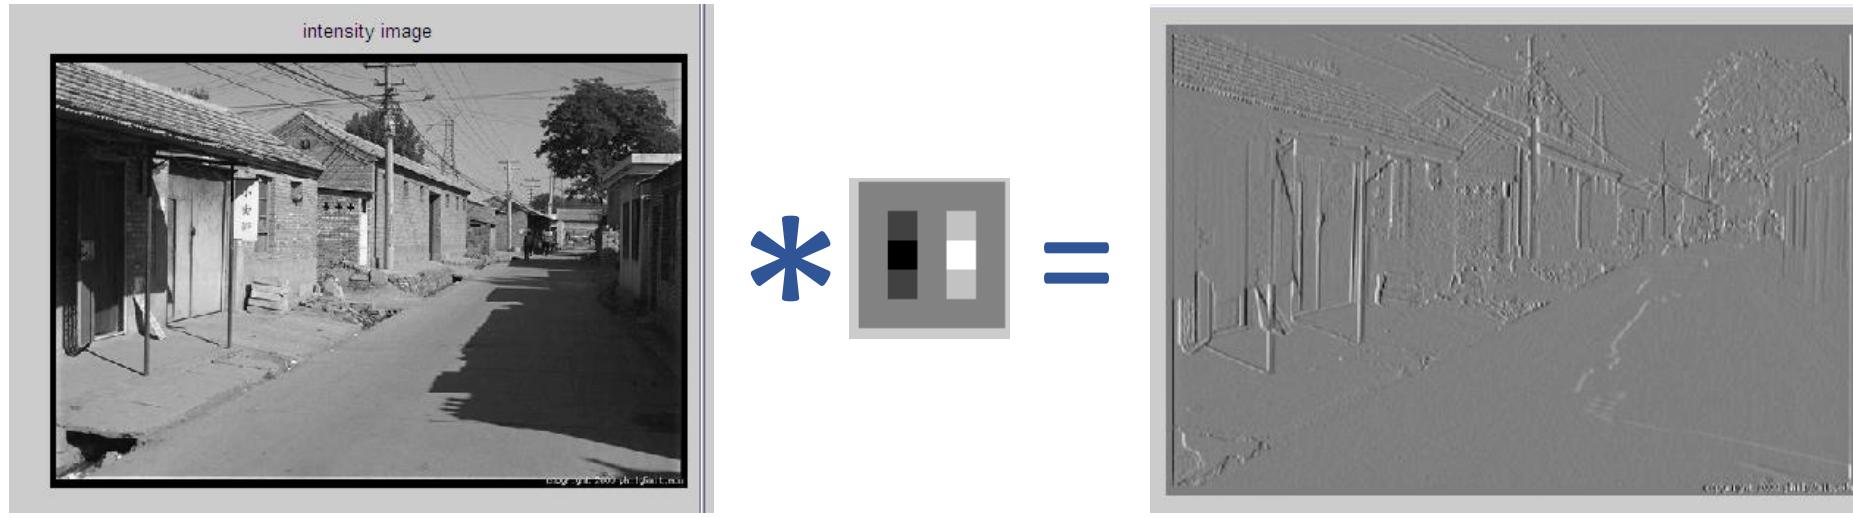


Previous Class

- Sometimes it makes sense to think of images and filtering in the frequency domain
 - Fourier analysis
- Can be faster to filter using FFT for large images ($N \log N$ vs. N^2 for auto-correlation)
- Images are mostly smooth
 - Basis for compression
- Remember to low-pass before sampling



Spatial domain



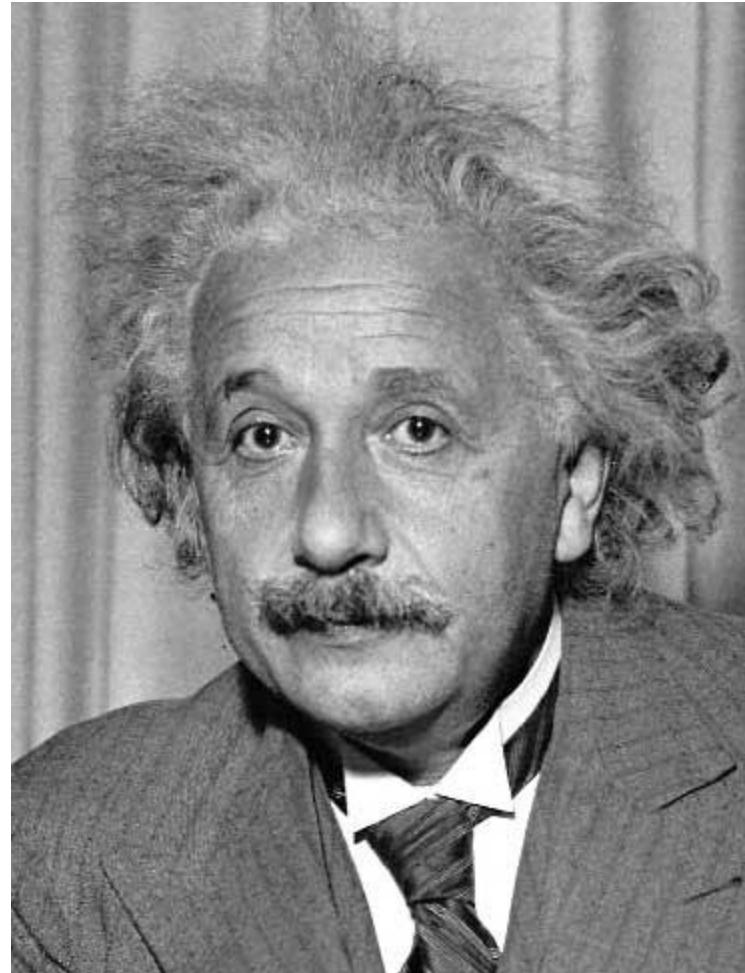
Today's Class

- Template matching
- Image Pyramids
- Compression



Template Matching

- Goal: find  in image
- Main challenge: What is a good similarity or distance measure between two patches? $D(\square, \square)$
 - Correlation
 - Zero-mean correlation
 - Sum Square Difference
 - Normalized Cross Correlation



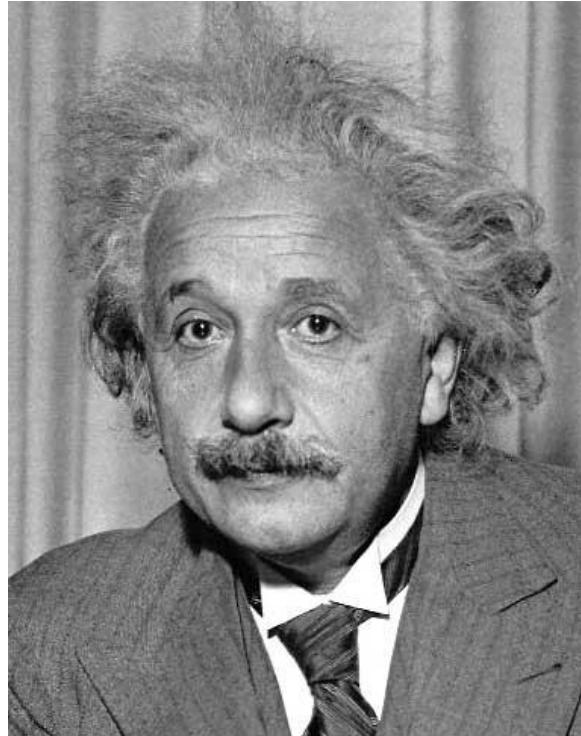
Dr. Sander Ali Khowaja

Matching with Filters

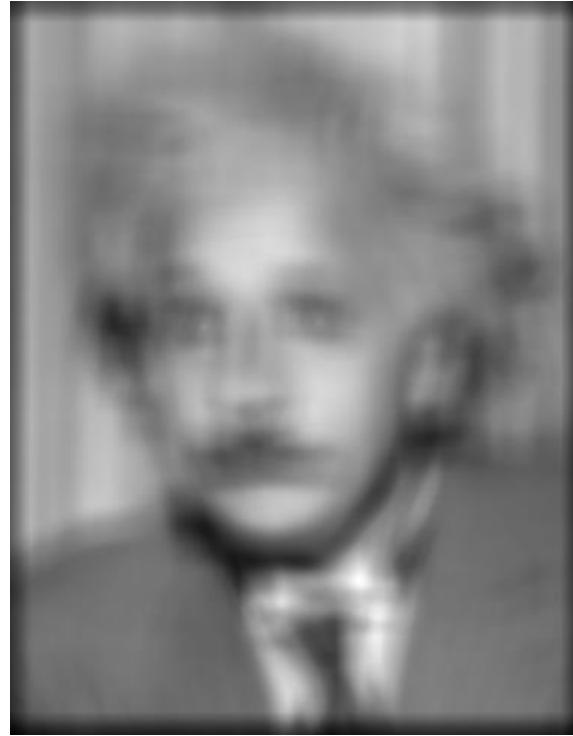
- Goal: find  in image
- Method 0: filter the image with eye patch

$$h[m,n] = \sum_{k,l} g[k,l] f[m+k, n+l]$$

g = filter f = image



Input



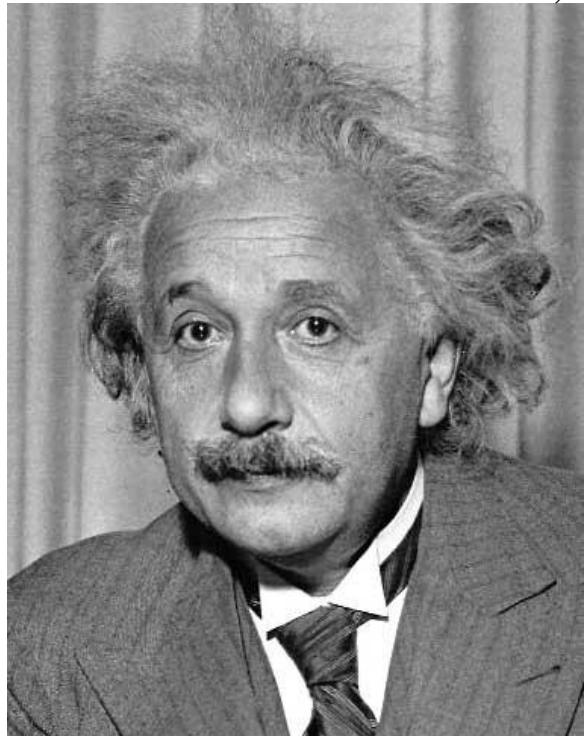
Filtered Image

What went wrong?

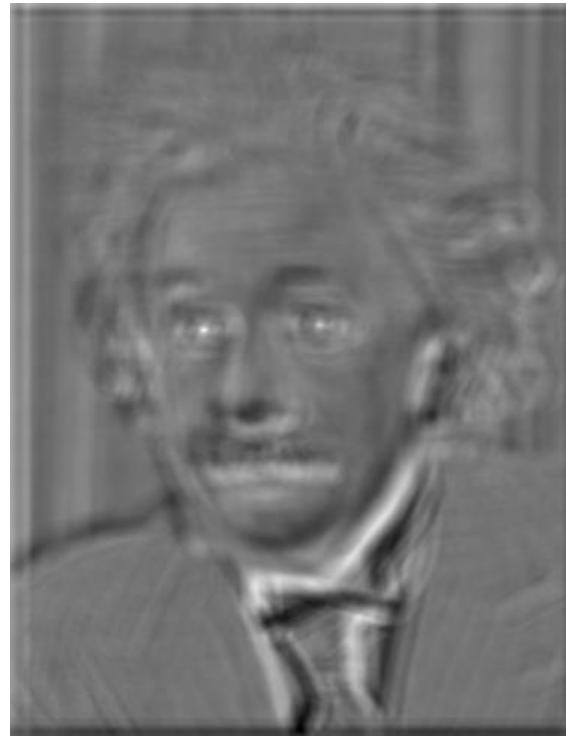
Matching with Filters

- Goal: find  in image
- Method 1: filter the image with zero-mean eye

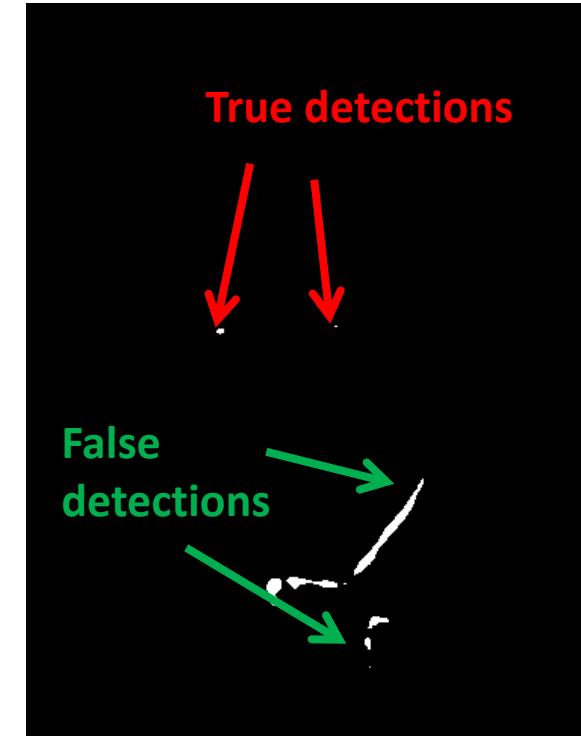
$$h[m,n] = \sum_{k,l} (g[k,l] - \bar{g}) \underbrace{(f[m+k, n+l])}_{\text{mean of template } g}$$



Input



Filtered Image (scaled)

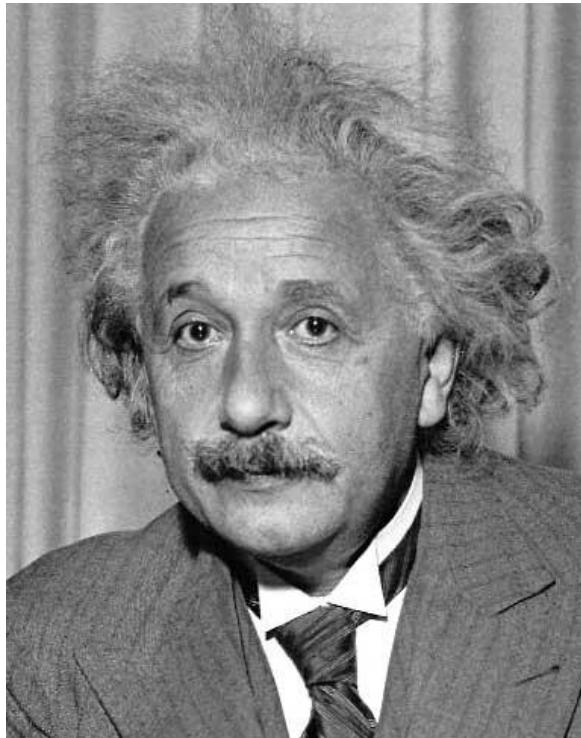


Thresholded Image

Matching with Filters

- Goal: find  in image
- Method 2: Sum of squared differences (SSD)

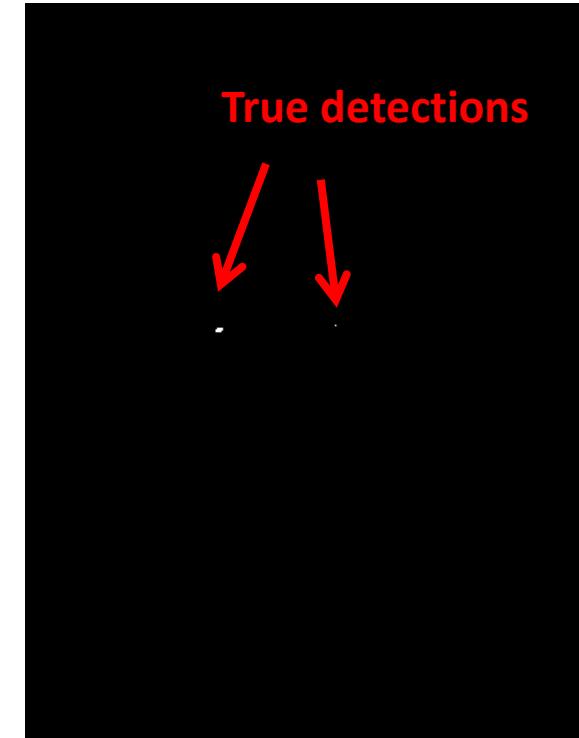
$$h[m,n] = \sum_{k,l} (g[k,l] - f[m+k, n+l])^2$$



Input



1 - sqrt(SSD)



Thresholded Image

Matching with Filters

Can SSD be implemented with linear filters?

$$h[m, n] = \sum_{k, l} (g[k, l] - f[m + k, n + l])^2$$

$$h[m, n] = \sum_{k, l} (g[k, l])^2 - 2 \sum_{k, l} g[k, l]f[m + k, n + l] + \sum_{k, l} (f[m + k, n + l])^2$$

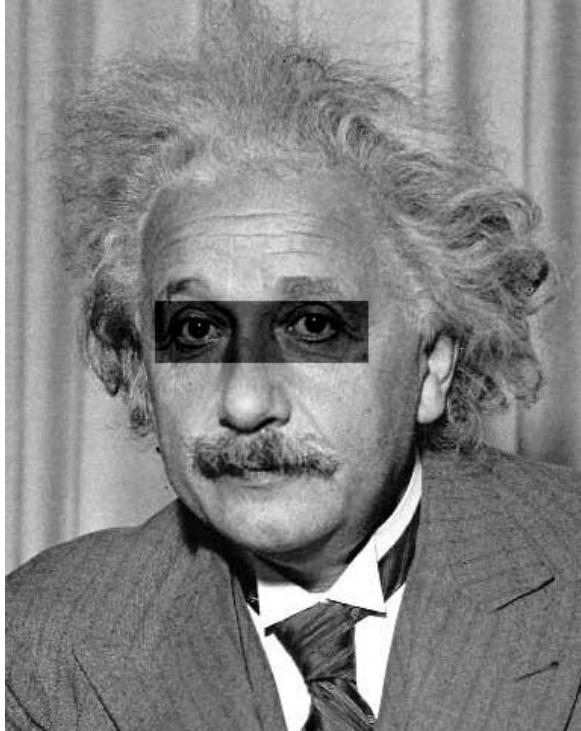


Matching with SSD

- Goal: find  in image
- Method 2: SSD

What's the potential downside of SSD?

$$h[m,n] = \sum_{k,l} (g[k,l] - f[m+k, n+l])^2$$



Input



1- sqrt(SSD)

Dr. Sander Ali Khowaja

Template Matching



Maximize:

$$R_{tf}[i,j] = \sum_m \sum_n f[m,n]t[m-i,n-j] = t \otimes f$$

(Cross-Correlation)

How do we locate the template in the image?

Minimize:

$$E[i,j] = \sum_m \sum_n (f[m,n] - t[m-i,n-j])^2$$

Sum of Squared Differences (SSD)

$$E[i,j] = \sum_m \sum_n (f^2[m,n] + t^2[m-i,n-j] - 2f[m,n]t[m-i,n-j])$$

Maximize



Convolution vs Correlation (Template Matching)

Convolution:

$$g[i,j] = \sum_m \sum_n f[m,n] t[i-m, j-n] = t * f$$

Correlation:

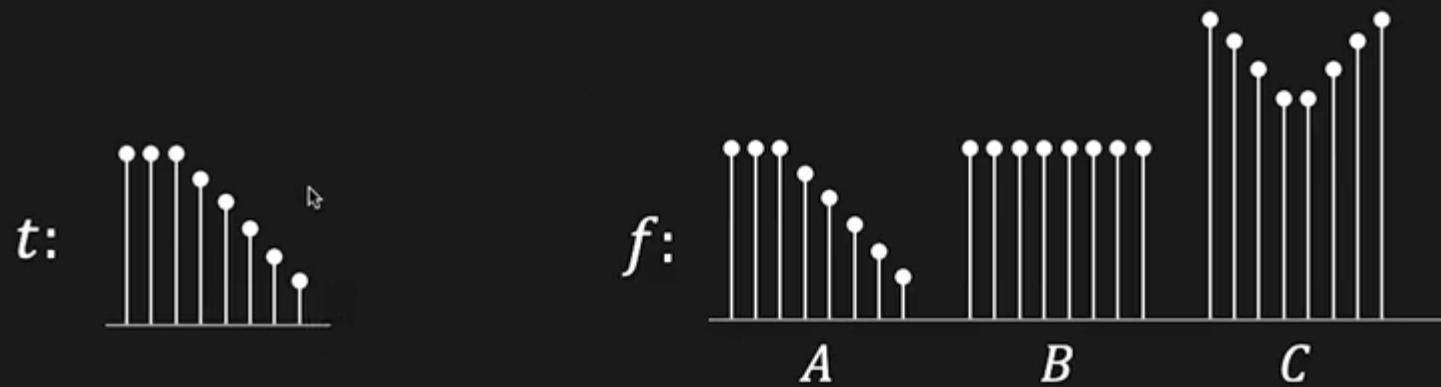
$$R_{tf}[i,j] = \sum_m \sum_n f[m,n] t[m-i, n-j] = t \otimes f$$

No Flipping in Correlation



Problem with Cross-Correlation

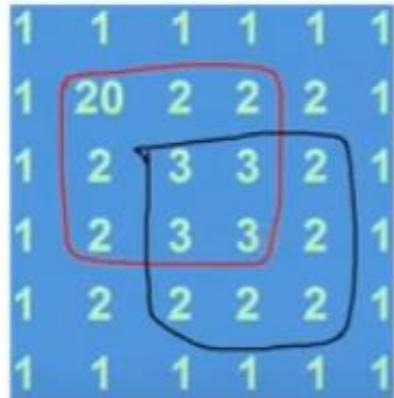
$$R_{tf}[i, j] = \sum_m \sum_n f[m, n]t[m - i, n - j] = t \otimes f$$



$$R_{tf}(C) > R_{tf}(B) > R_{tf}(A)$$

We need $R_{tf}(A)$ to be the maximum!

Cross Correlation Example



| | | | | | |
|----|----|-----|----|----|----|
| 47 | 54 | 56 | 20 | 18 | 12 |
| 54 | 87 | 94 | 40 | 34 | 21 |
| 56 | 94 | 107 | 54 | 43 | 24 |
| 20 | 40 | 54 | 56 | 44 | 24 |
| 18 | 34 | 43 | 44 | 37 | 22 |
| 12 | 21 | 24 | 24 | 22 | 15 |

```
clc
clear all
close all
warning off
a=[1 1 1 1 1 1;
   1 20 2 2 2 1;
   1 2 3 3 2 1;
   1 2 3 3 2 1;
   1 2 2 2 2 1;
   1 1 1 1 1 1];
w1=[3 3 2;
    3 3 2;
    2 2 2];
newA = zeros(size(a)+2);
newA(2:end-1,2:end-1)=a;
```

```
a=newA;
[row col]=size(a);
a1=zeros(row,col);
for x=2:1:row-1
    for y=2:1:col-1
        a1(x,y)=w1(1)*a(x-1,y-1)+w1(2)*a(x-1,y)+w1(3)*...
                   a(x-1,y+1)+w1(4)*a(x,y-1)+w1(5)*a(x,y)+w1(6)*...
                   a(x,y+1)+w1(7)*a(x+1,y-1)+w1(8)*a(x+1,y)+w1(9)*...
                   a(x+1,y+1);
    end
end
a1=a1(2:end-1,2:end-1);
```

Matching with Filters

- Goal: find  in image
- Method 3: Normalized cross-correlation

$$h[m,n] = \frac{\sum_{k,l} (g[k,l] - \bar{g})(f[m+k, n+l] - \bar{f}_{m,n})}{\left(\sum_{k,l} (g[k,l] - \bar{g})^2 \sum_{k,l} (f[m+k, n+l] - \bar{f}_{m,n})^2 \right)^{0.5}}$$

mean template mean image patch

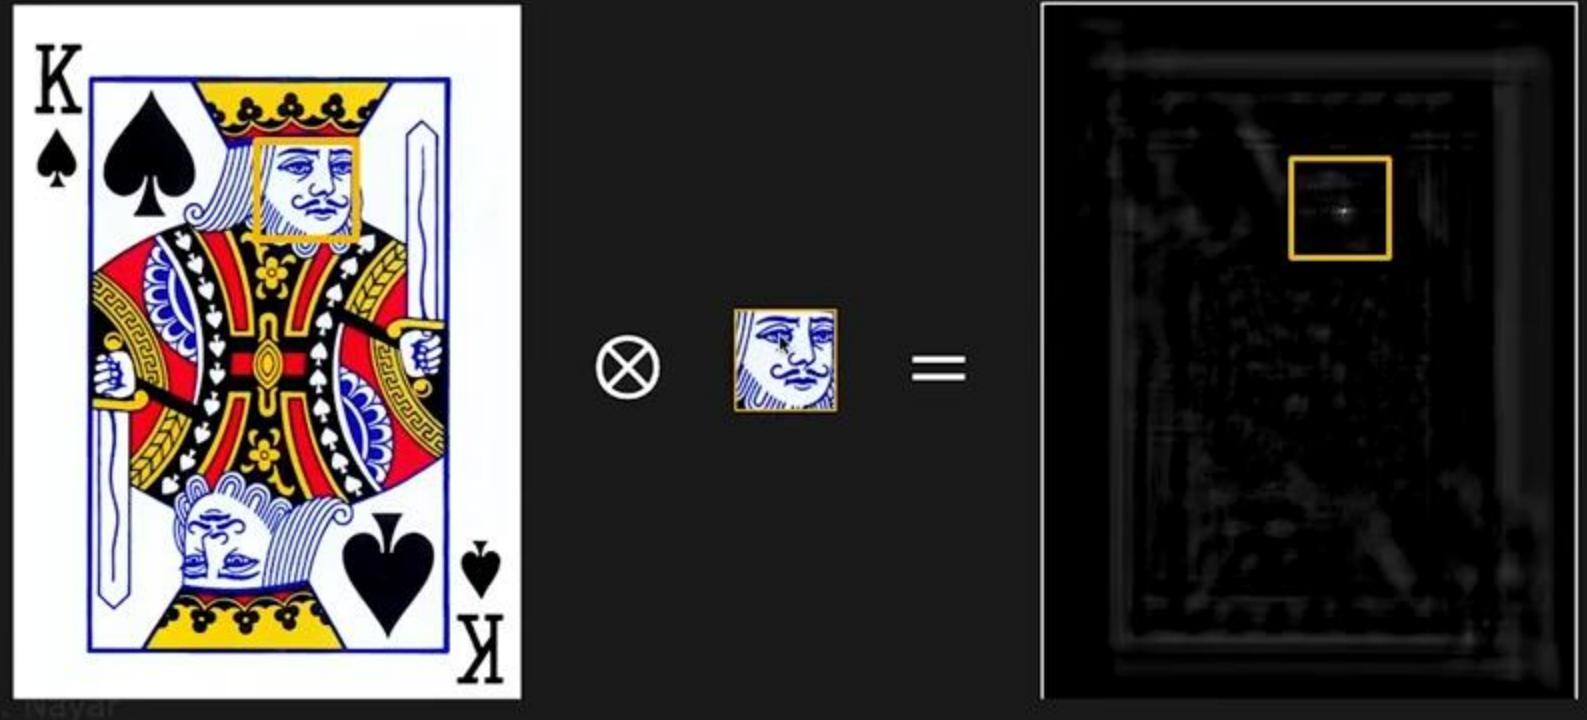
Matlab: `normxcorr2(template, im)`

Normalized Cross Correlation

Account for energy differences

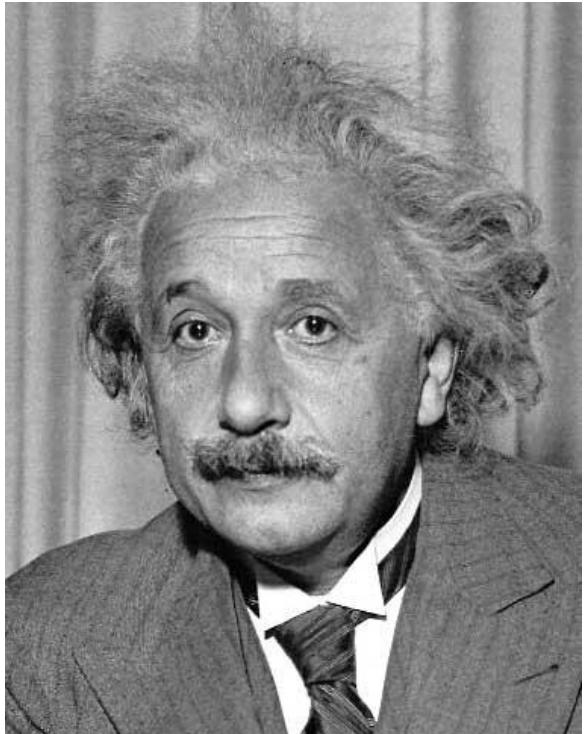
$$N_{tf}[i,j] = \frac{\sum_m \sum_n f[m,n]t[m-i,n-j]}{\sqrt{\sum_m \sum_n f^2[m,n]} \sqrt{\sum_m \sum_n t^2[m-i,n-j]}}$$

In insensitive to changes in Brightness



Matching with Filters

- Goal: find  in image
- Method 3: Normalized cross-correlation

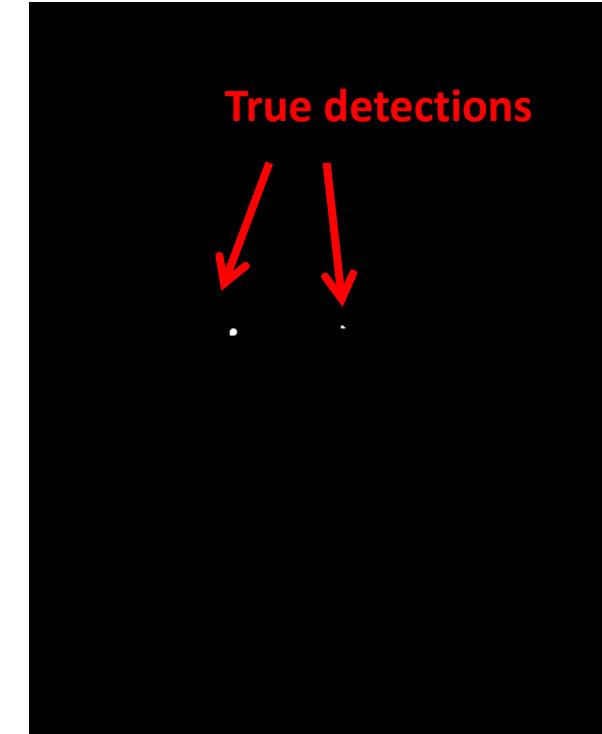


Input



Normalized X-Correlation

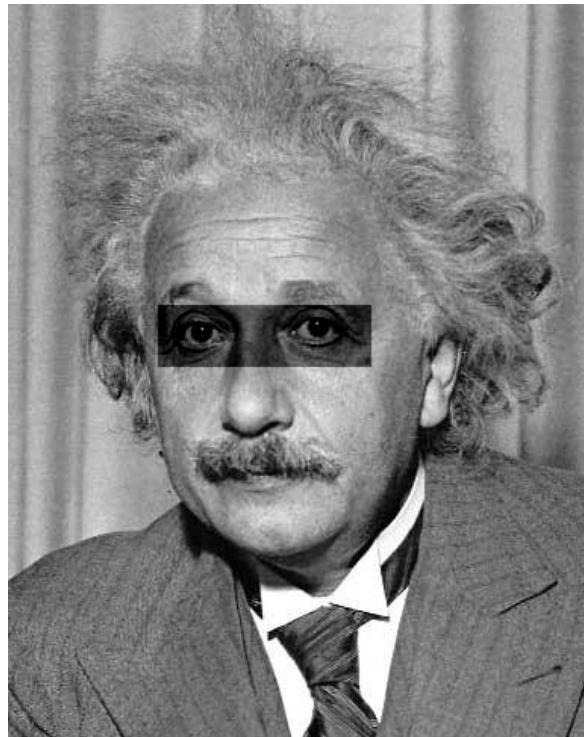
Dr. Sander Ali Khowaja



Thresholded Image

Matching with Filters

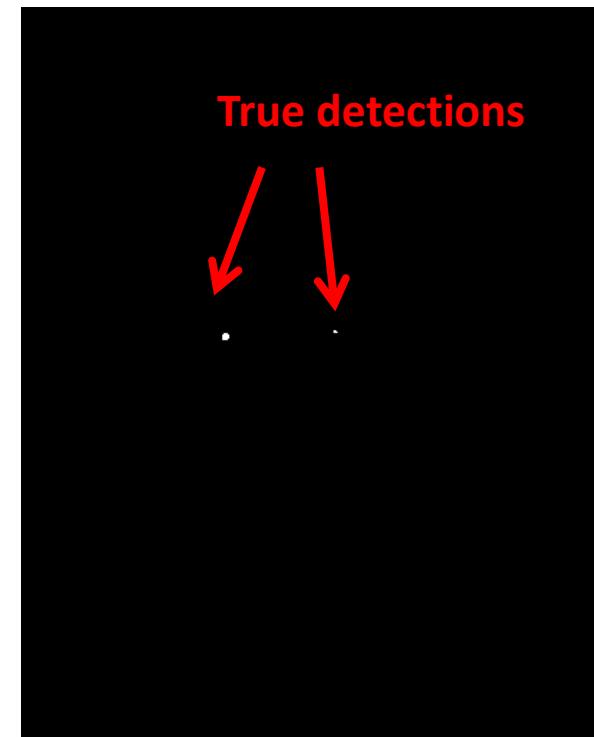
- Goal: find  in image
- Method 3: Normalized cross-correlation



Input



Normalized X-Correlation



Thresholded Image

What is the best method to use then?

A: Depends

- Zero-mean filter: fastest but not a great matcher
- SSD: next fastest, sensitive to overall intensity
- Normalized cross-correlation: slowest, invariant to local average intensity and contrast



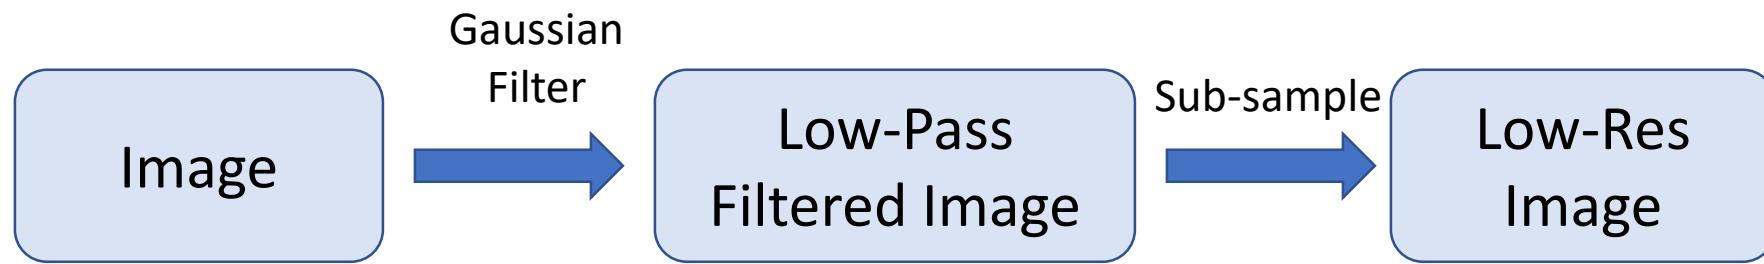
What about scaling

Q: What if we want to find larger or smaller eyes?

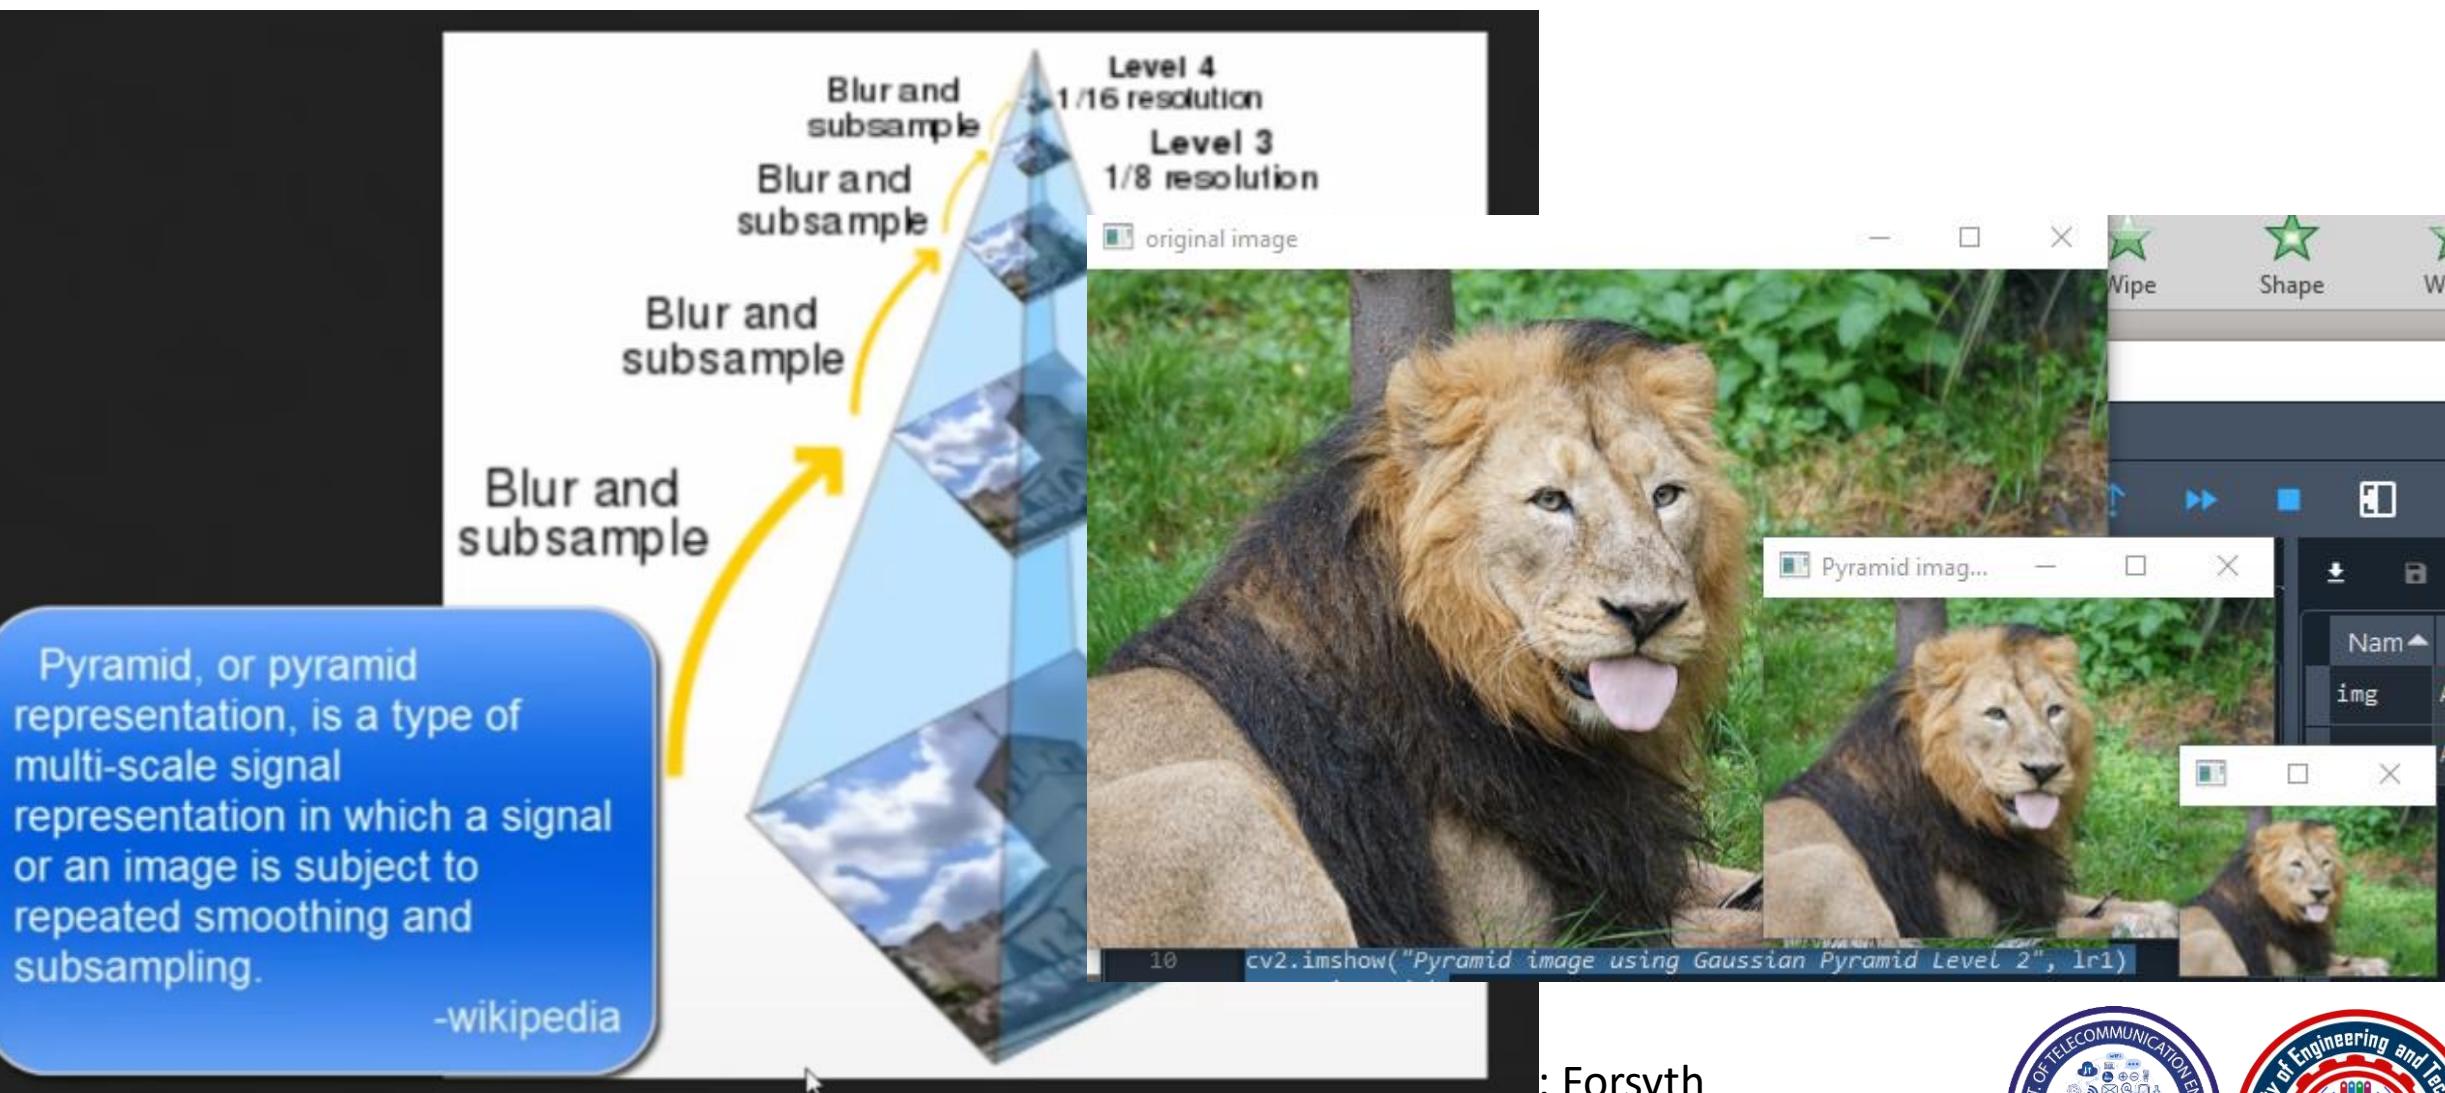
A: Image Pyramid



Review of Sampling



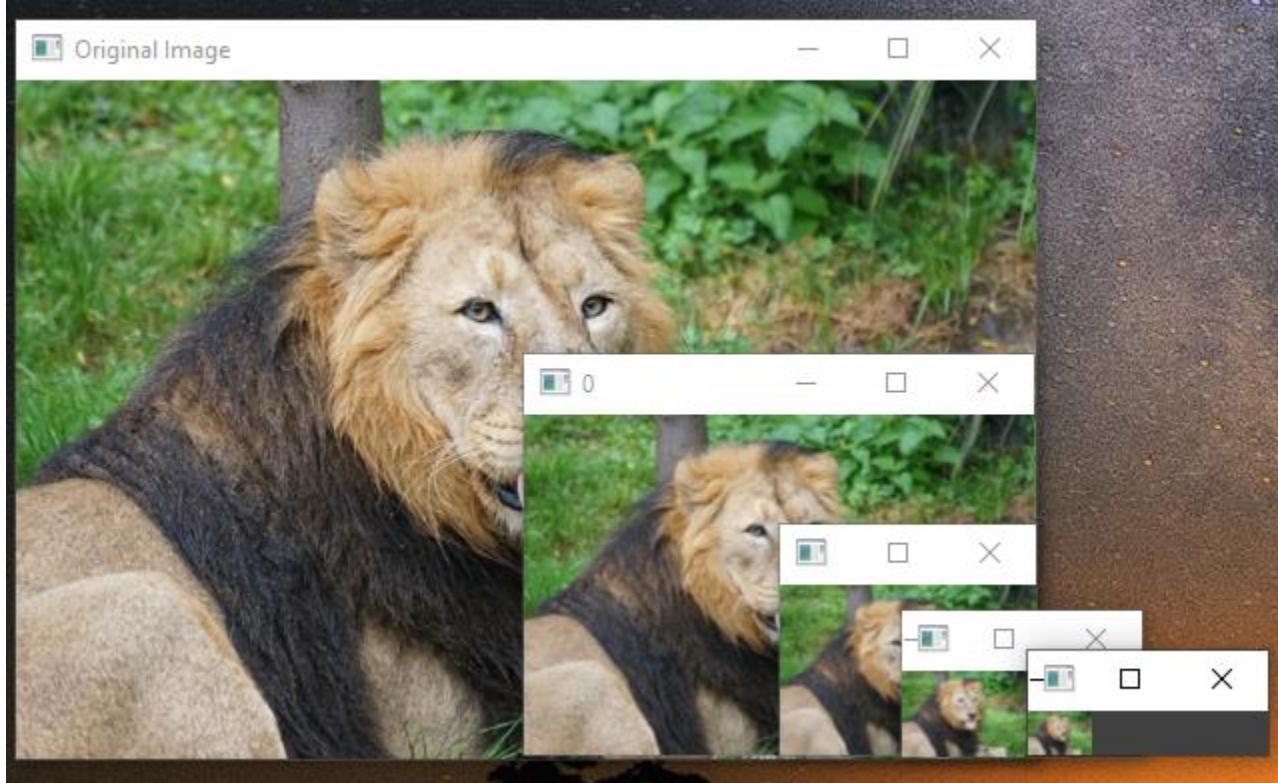
Gaussian Pyramid



Gaussian Pyramid OpenCV

```
1 import cv2
2 import numpy as np
3
4 img = cv2.imread("0809x4.png")
5 lr = cv2.pyrDown(img)
6 lr1 = cv2.pyrDown(lr)
7 hr1 = cv2.pyrUp(lr1)
8 hr = cv2.pyrUp(hr1)
9
10 cv2.imshow("original image", img)
11 cv2.imshow("Pyramid image using Gaussian Pyramid", lr)
12 cv2.imshow("Pyramid image using Gaussian Pyramid Level 2", lr1)
13 cv2.imshow("Pyramid image using Gaussian Pyramid", hr1)
14 cv2.imshow("Pyramid image using Gaussian Pyramid Level 2", hr)
15 cv2.waitKey(0)
16 cv2.destroyAllWindows()
```

```
1 import cv2
2 import numpy as np
3
4 img = cv2.imread("0809x4.png")
5 layer = img.copy()
6 gp = [layer]
7
8 for i in range(6):
9     layer = cv2.pyrDown(layer)
10    gp.append(layer)
11    cv2.imshow(str(i), layer)
12
13 cv2.imshow("Original Image", img)
14 cv2.waitKey(0)
15 cv2.destroyAllWindows()
```

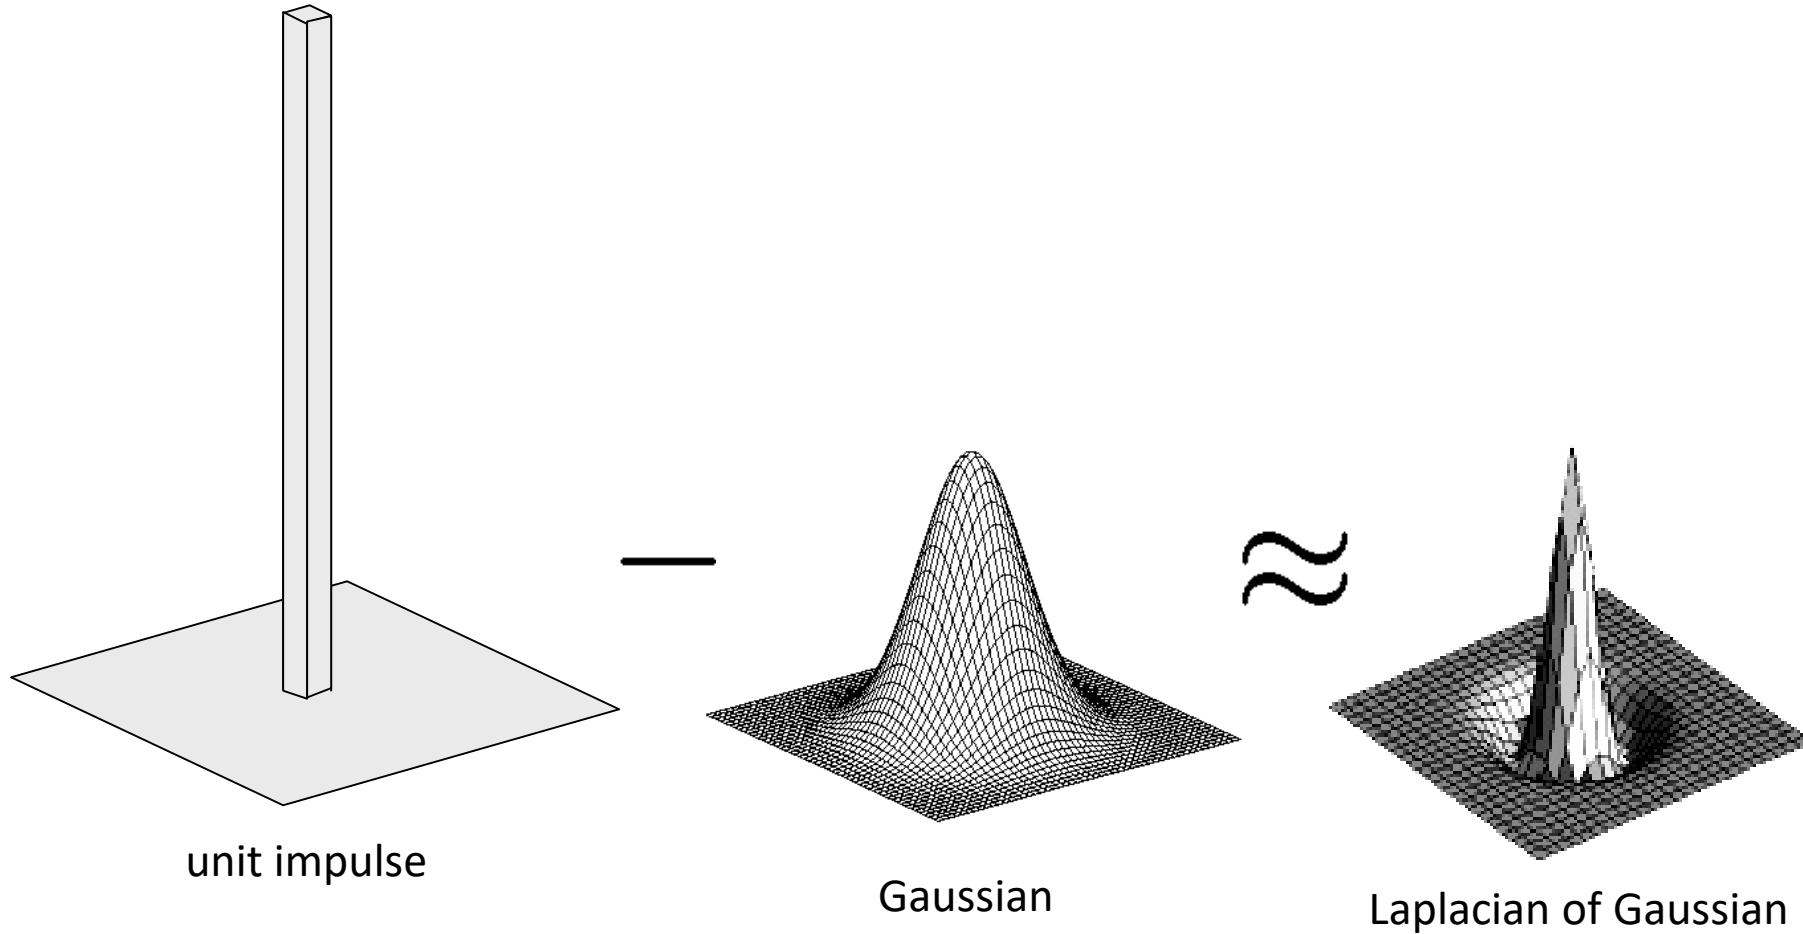


Template matching with Gaussian Pyramid

Input: Image, Template

1. Match template at current scale
2. Downsample image
In practice, scale step of 1.1 to 1.2
3. Repeat 1-2 until image is very small
4. Take responses above some threshold, perhaps with non-maxima suppression

Laplacian Filter

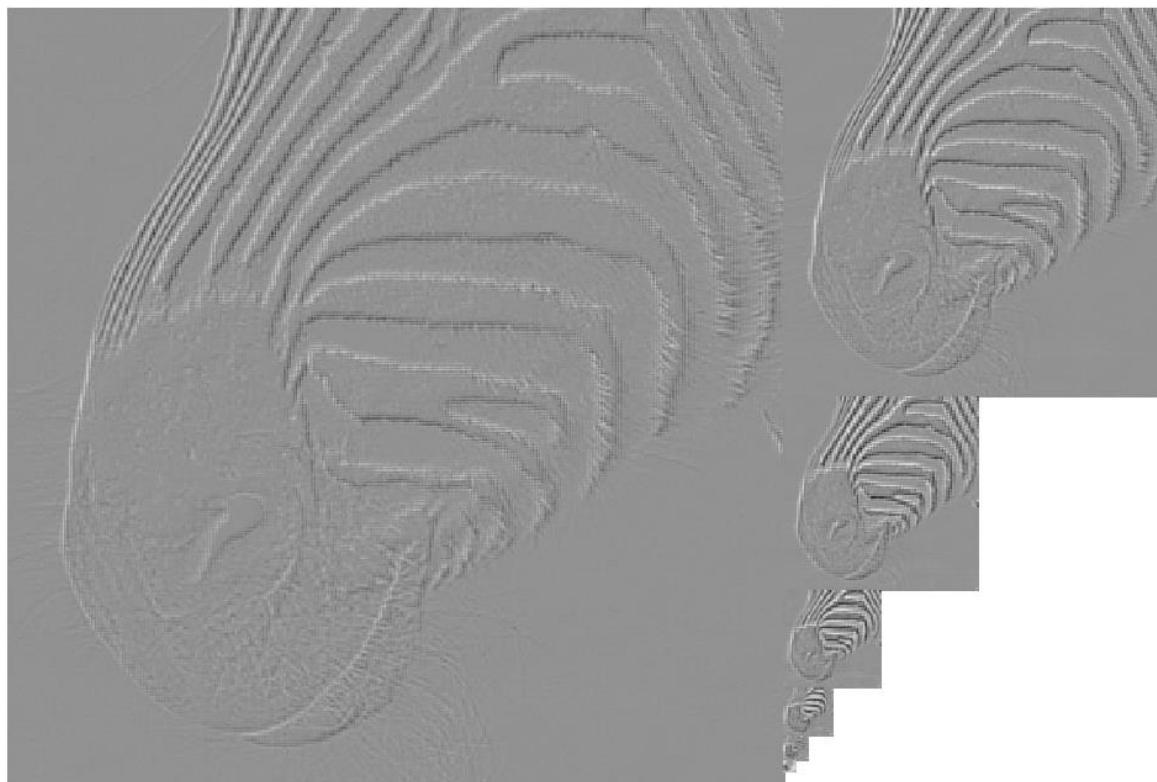


Source: Lazebnik

Laplacian Pyramid



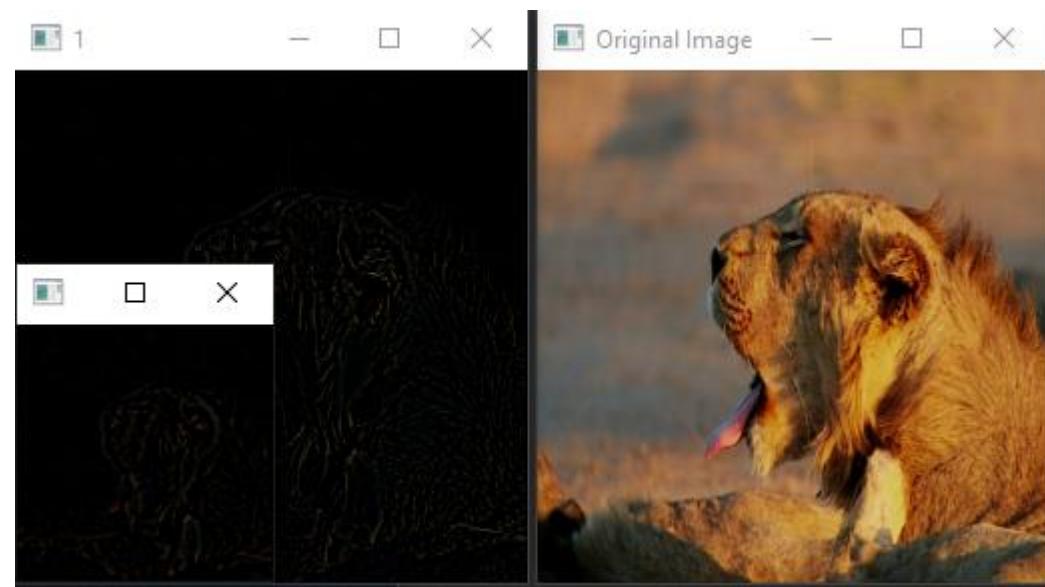
512 256 128 64 32 16 8



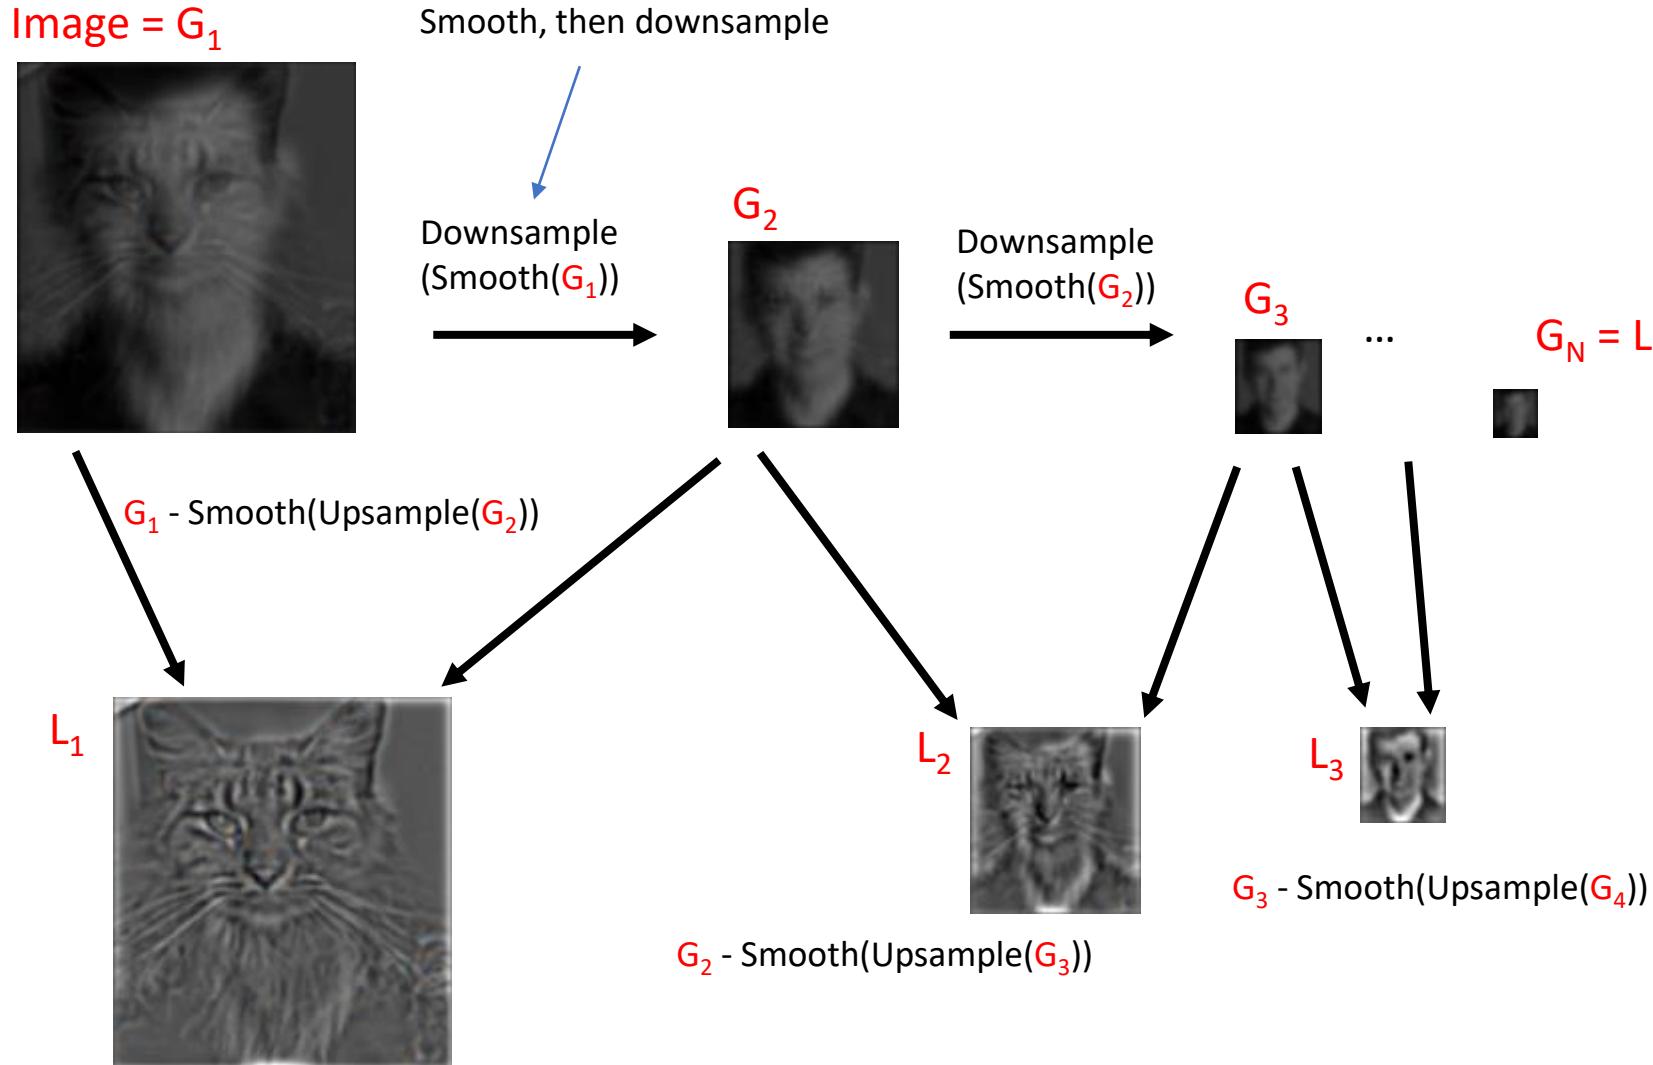
Source: Forsyth

Python Example (Laplacian Pyramid)

```
1 import cv2
2 import numpy as np
3
4 img = cv2.imread("09011.png")
5 layer = img.copy()
6 gp = [layer]
7
8 for i in range(6):
9     layer = cv2.pyrDown(layer)
10    gp.append(layer)
11    #cv2.imshow(str(i), layer)
12
13 layer = gp[5]
14 cv2.imshow('Upper Level gaussian pyramid', layer)
15
16 lp = [layer]
17
18 for i in range(5, 0, -1):
19     gaussian_extend = cv2.pyrUp(gp[i])
20     laplacian = cv2.subtract(gp[i-1], gaussian_extend)
21     cv2.imshow(str(i), laplacian)
22
23
24 cv2.imshow("Original Image", img)
25 cv2.waitKey(0)
26 cv2.destroyAllWindows()
```



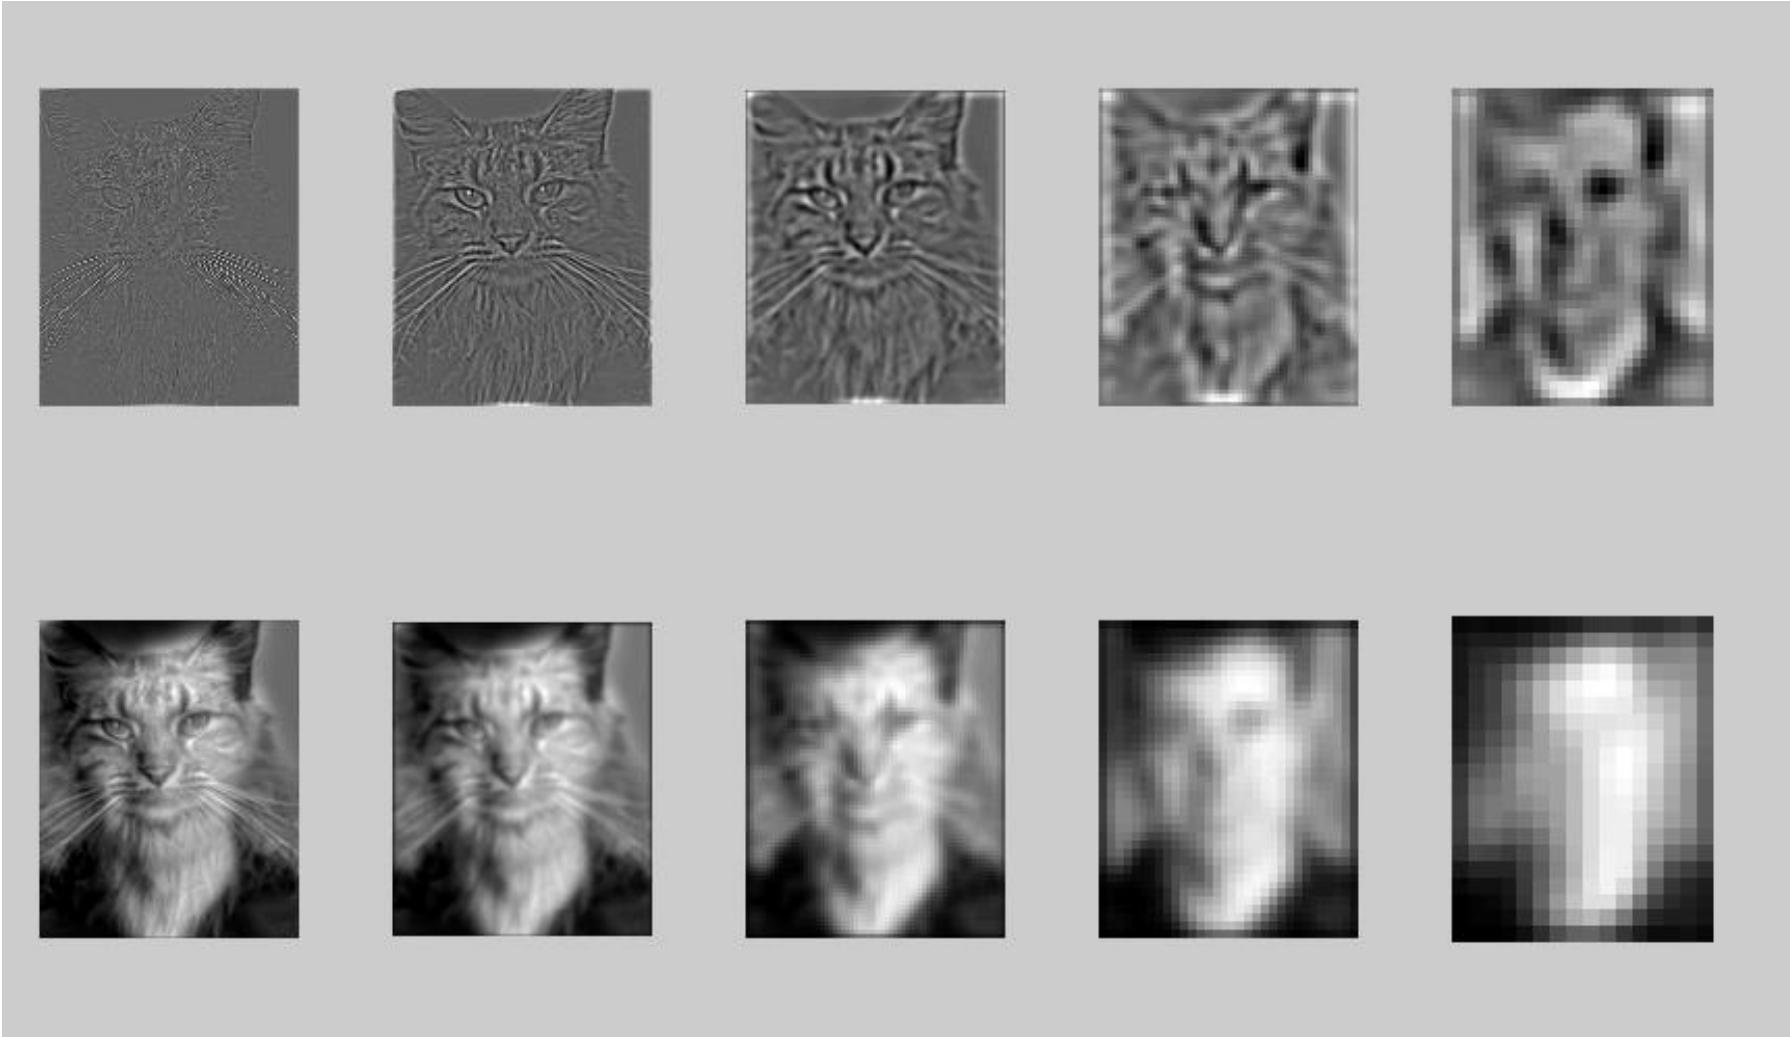
Creating the Gaussian/Laplacian Pyramid



- Use same filter for smoothing in each step (e.g., Gaussian with $\sigma = 2$)
- Downsample/upsample with “nearest” interpolation

Hybrid image in Laplacian Pyramids

High frequency → Low frequency



Reconstructing Image from Laplacian Pyramid

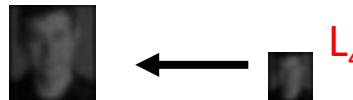
$$\text{Image} = L_1 + \text{Smooth}(\text{Upsample}(G_2))$$



$$G_2 = L_2 + \text{Smooth}(\text{Upsample}(G_3))$$



$$G_3 = L_3 + \text{Smooth}(\text{Upsample}(L_4))$$

 L_4  L_1

- Use same filter for smoothing as in deconstruction
- Upsample with “nearest” interpolation
- Reconstruction will be lossless

Dr. Sander Ali Khowaja



Major Uses of Image Pyramid

- Object detection
 - Scale search
 - Features
- Detecting stable interest points
- Course-to-fine registration
- Compression







Template matching in Python

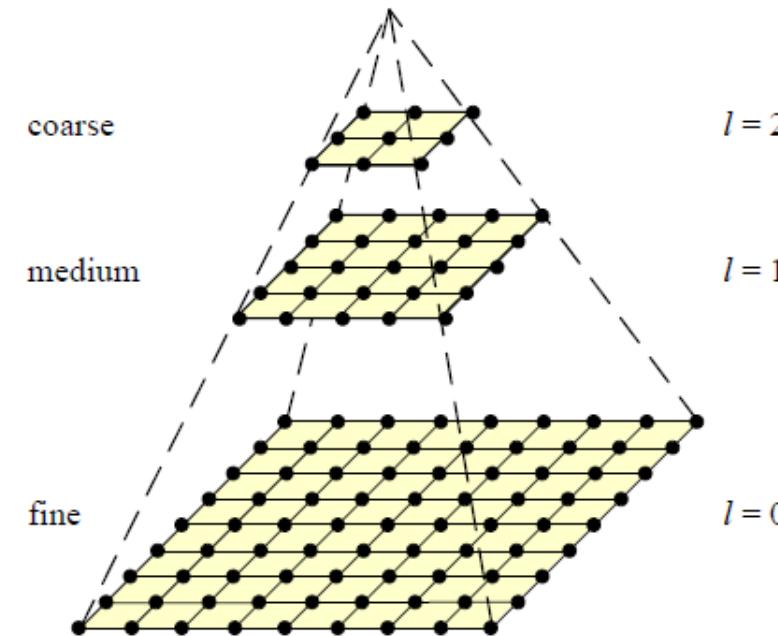
```
1 # Python program to illustrate
2 # template matching
3 import cv2
4 import numpy as np
5
6 # Read the main image
7 img_rgb = cv2.imread('Image Path').
8
9 # Convert it to grayscale
10 img_gray = cv2.cvtColor(img_rgb, cv2.COLOR_BGR2GRAY)
11
12 # Read the template
13 template = cv2.imread('template', 0)
14
15 # Store width and height of template in w and h
16 w, h = template.shape[::-1]
17
18 # Perform match operations.
19 res = cv2.matchTemplate(img_gray, template, cv2.TM_CCOEFF_NORMED)
20
21 # Specify a threshold
22 threshold = 0.8
23
24 # Store the coordinates of matched area in a numpy array
25 loc = np.where(res >= threshold)
26
27 # Draw a rectangle around the matched region.
28 for pt in zip(*loc[::-1]):
29     cv2.rectangle(img_rgb, pt, (pt[0] + w, pt[1] + h), (0, 255, 255), 2)
30
31 # Show the final image with the matched area.
32 cv2.imshow('Detected', img_rgb)
```

```
1 # Python program to illustrate
2 # template matching
3 import cv2
4 import numpy as np
5
6 # Read the main image
7 img_rgb = cv2.imread('Image Path').
8
9 # Convert it to grayscale
10 img_gray = cv2.cvtColor(img_rgb, cv2.COLOR_BGR2GRAY)
11
12 # Read the template
13 template = cv2.imread('template', 0)
14
15 # Store width and height of template in w and h
16 w, h = template.shape[::-1]
17
18 # Resize the image according to scale and
19 # keeping track of ratio of resizing
20 resize = imutils.resize(img_gray, width=int(shape[0]), height=int(img_gray.shape[1]*scale))
21
22 # If resize image is smaller than that of template
23 # break the loop
24 # Detect edges in the resized, grayscale image and apply template
25 # Matching to find the template in image edged
26 # If we have found a new maximum correlation value, update
27 # the found variable if
28 # found = null/maxVal > found][0]
29 if resized.shape[0] < h or resized.shape[1] < w:
30     break
31 found=(maxVal, maxLoc, r)
32
33 # Unpack the found variables and compute(x,y) coordinates
34 # of the bounding box
35 (__, maxLoc, r)=found
36 (startX, startY)=(int(maxLoc[0]*r), int maxLoc[1]*r)
37 (endX, endY)=(int((maxLoc[0]+tw)*r), int(maxLoc[1]+tH)*r)
38
39 # Draw a bounding box around the detected result and display the image
40 cv2.rectangle(image, (startX, startY), (endX, endY), (0, 0, 255), 2)
41 cv2.imshow("Image", image)
42 cv2.waitKey(0)
```

Pyramid Application: Coarse to Fine Image Registration

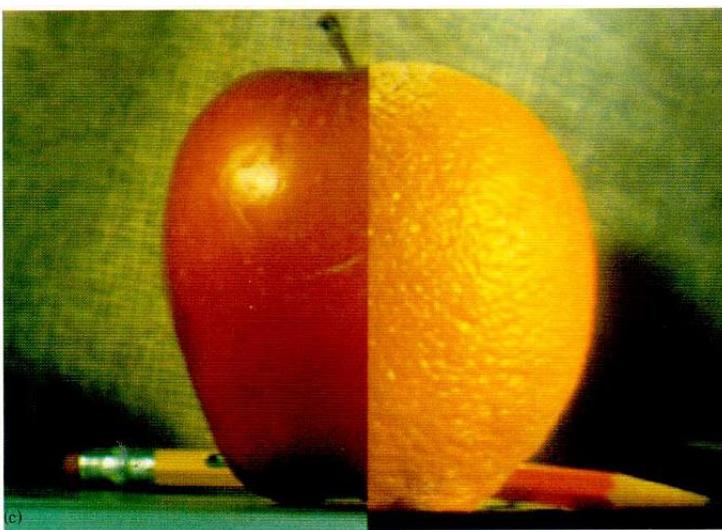
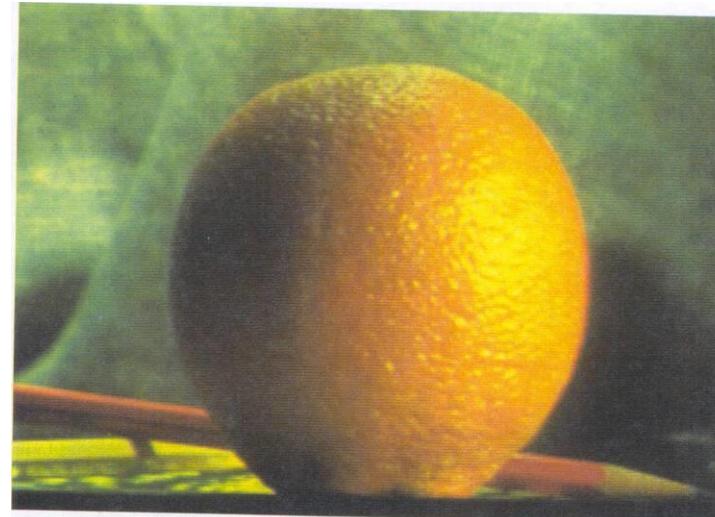
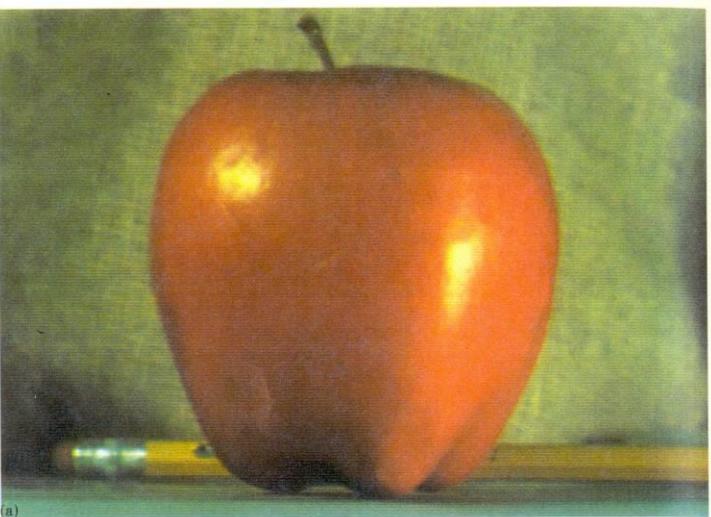
1. Compute Gaussian pyramid
2. Align with coarse pyramid
3. Successively align with finer pyramids
 - Search smaller range

Why is this faster?

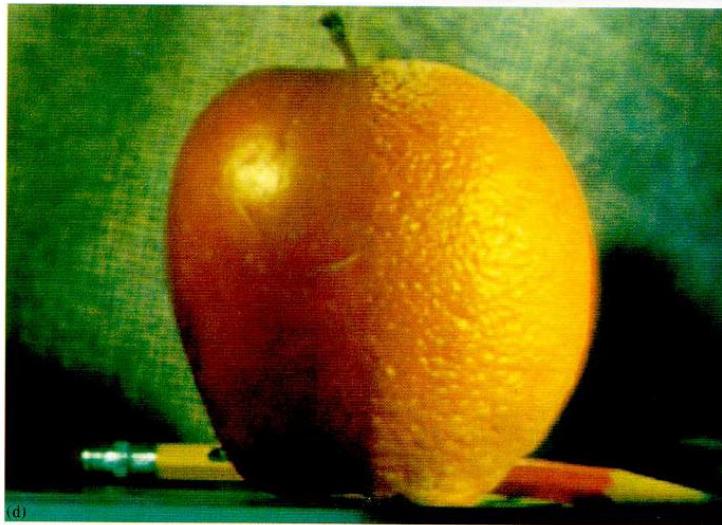
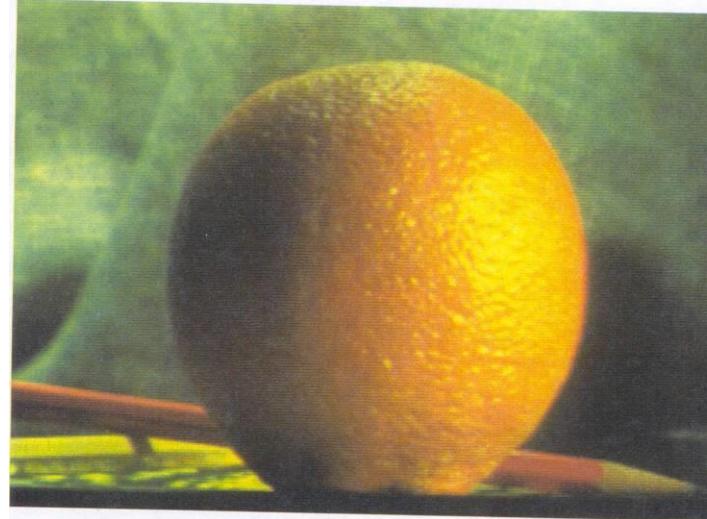
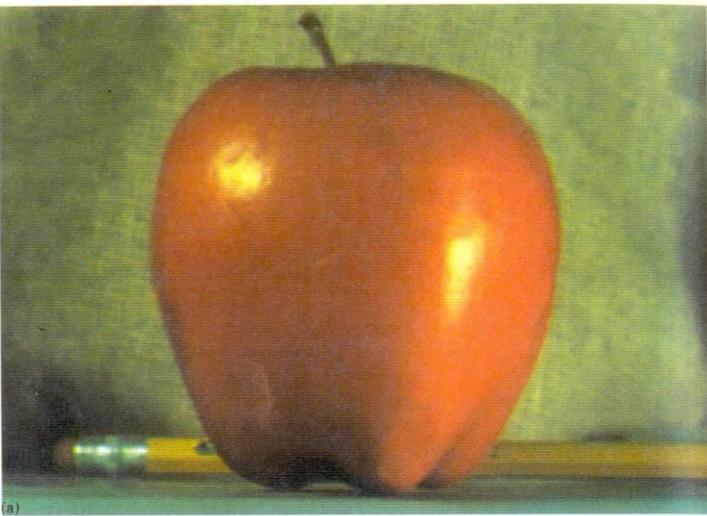


Are we guaranteed to get the same result?

Applications: Pyramid Blending

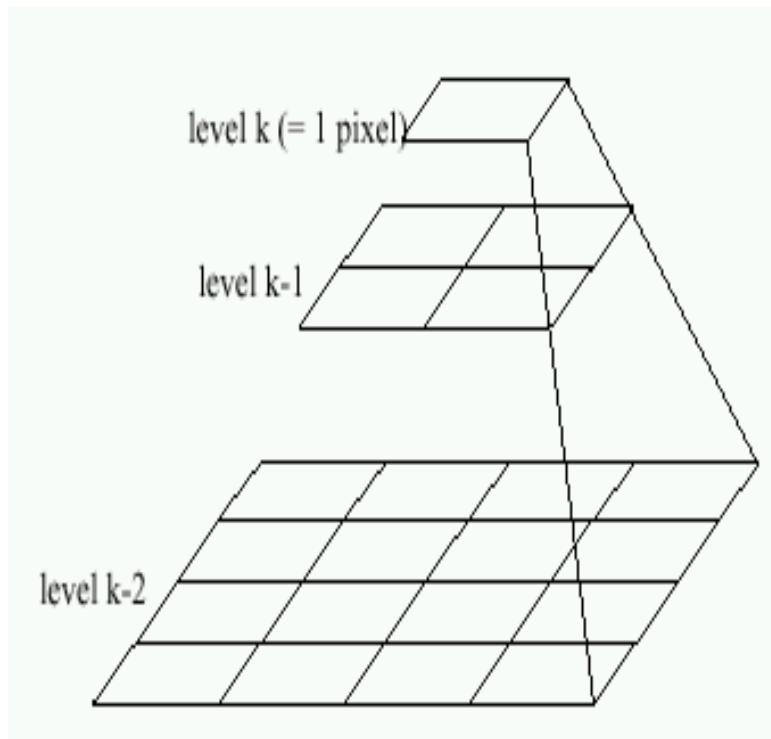


Applications: Pyramid Blending

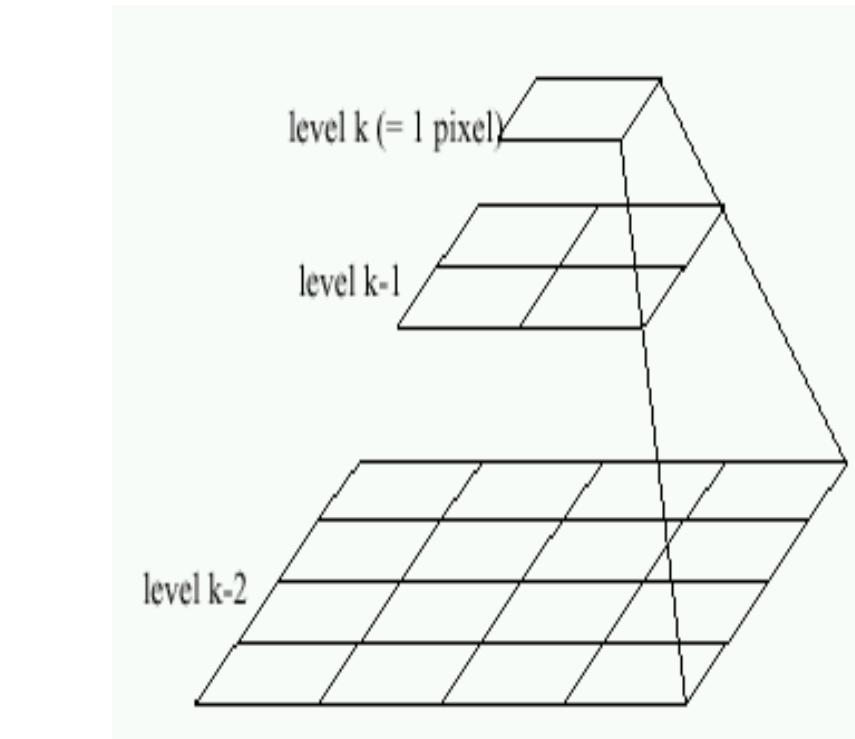
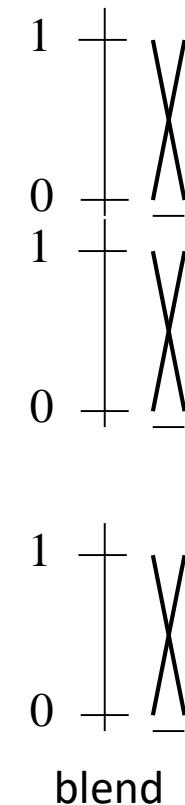


Pyramid Blending

- At low frequencies, blend slowly
- At high frequencies, blend quickly



Left pyramid



Right pyramid

Pyramid Blending: Examples

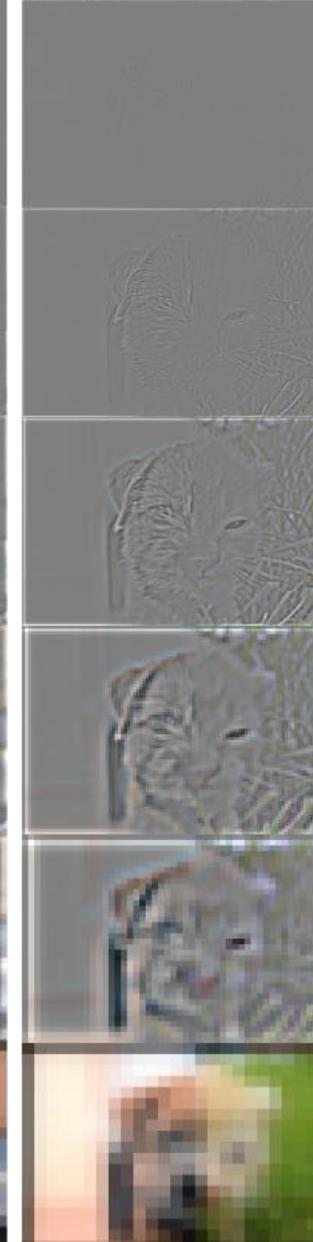
Gaussian Pyramids



Laplacian Pyramids



Billeted Laplacian Pyramids



Blended Output Image



ja



Pyramid: Blending Example

Gaussian Pyramids

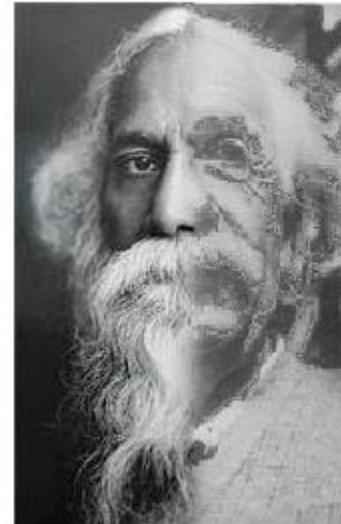


Laplacian Pyramids



Blended Laplacian Pyramids

Blended Output Image

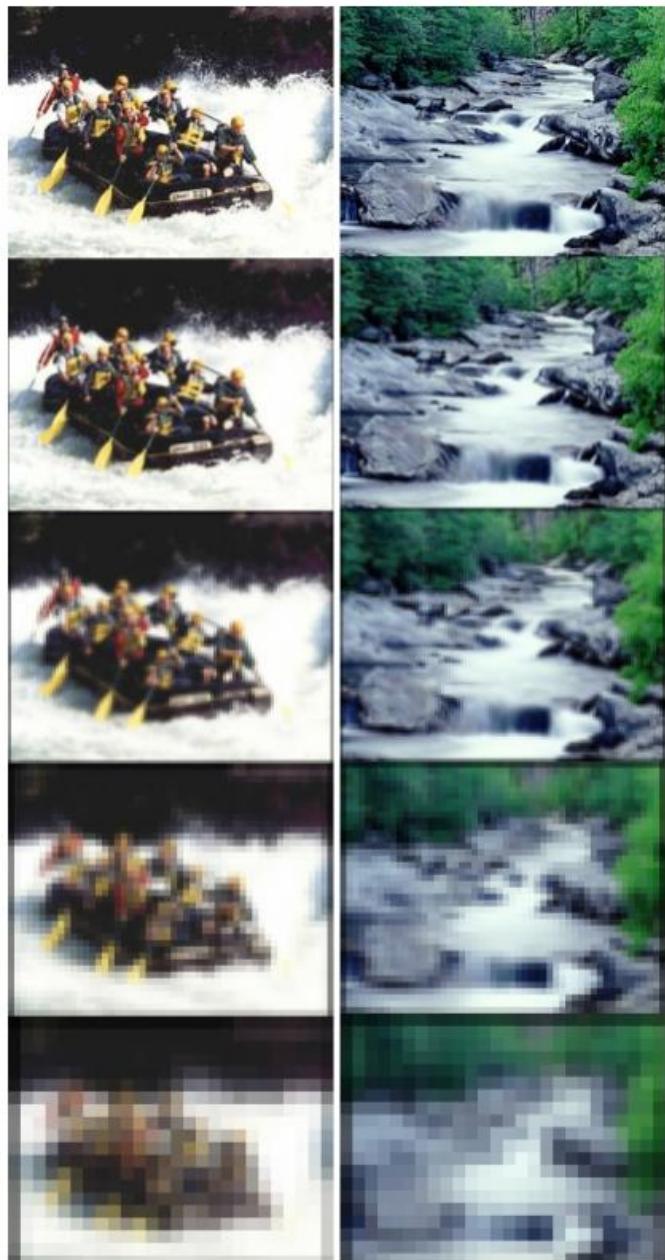


Sohder Ali Khowaja

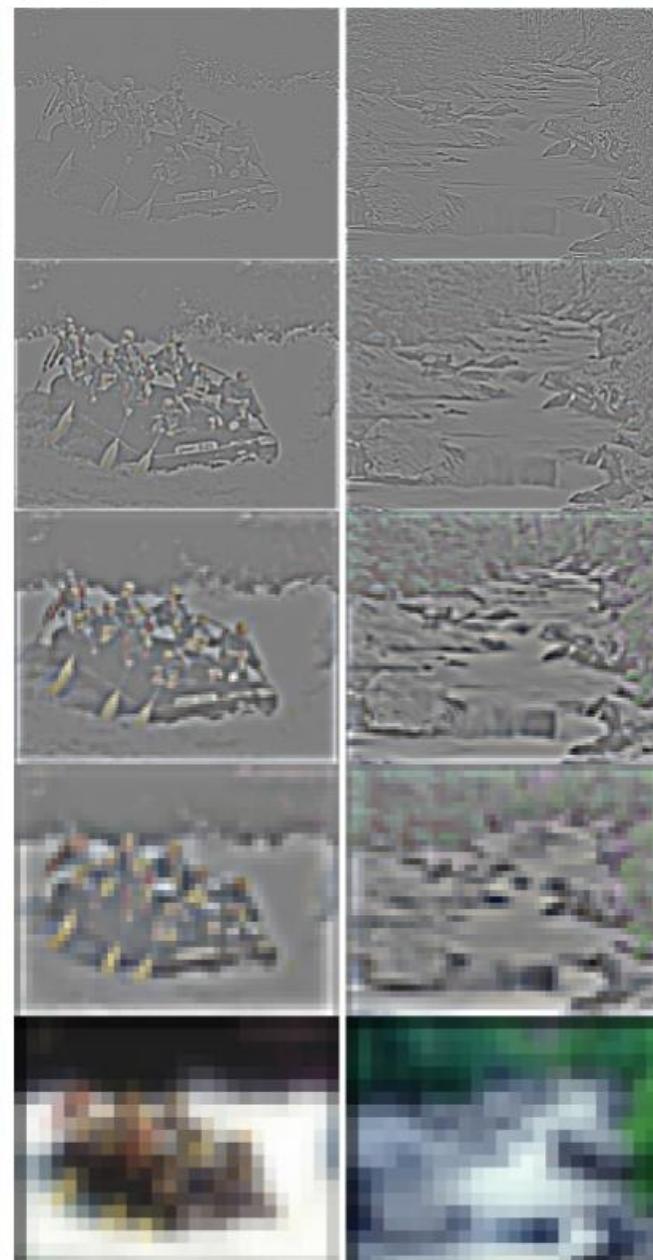


Pyramid: Blending Example

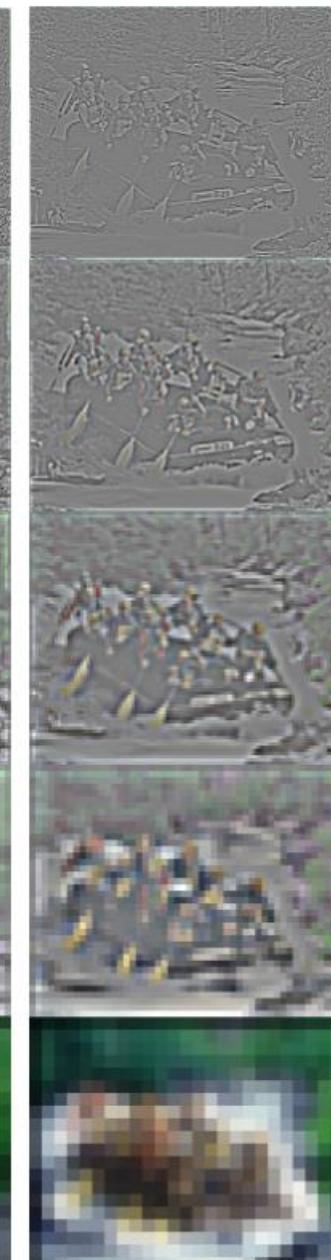
Gaussian Pyramids



Laplacian Pyramids



Blended Laplacian Pyramids



Blended Output Image

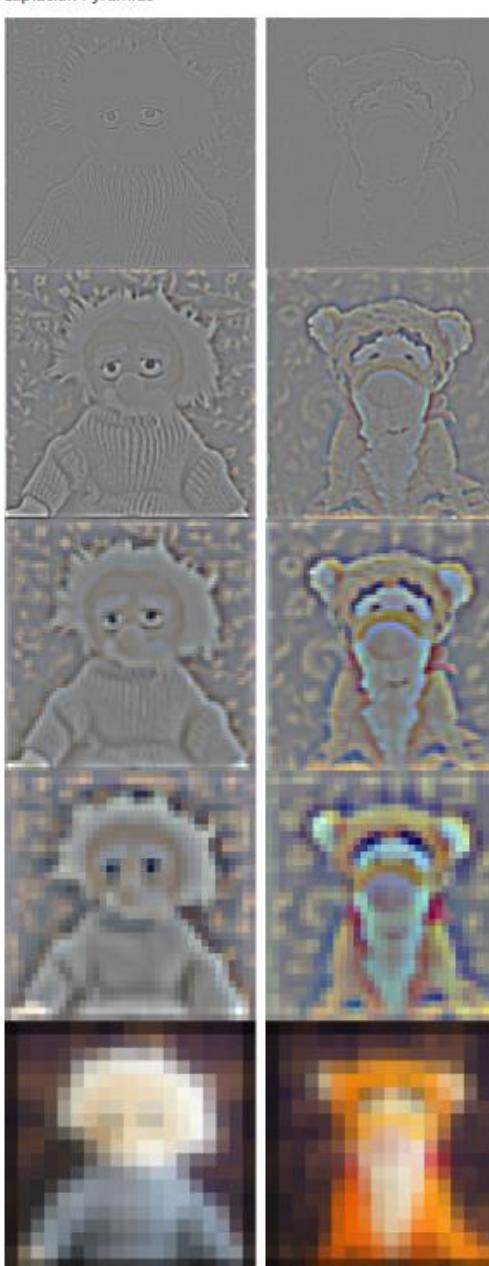


Pyramid: Blending Example

Gaussian Pyramids



Laplacian Pyramids



Blended Laplacian Pyramids



Blended Output Image



Khawaja



Python code for Image Blending

```
1 import cv2
2 import numpy as np
3
4 img1 = cv2.imread('0809.png')
5 img2 = cv2.imread('09011.png')
6 print(img1.shape)
7 print(img2.shape)
8 blend_lion = np.hstack((img1[:, :256], img2[:, 256:]))
9 ## generate Gaussian pyramid for lion1
10 img1_copy = img1.copy()
11 gp_img1 = [img1_copy]
12 for i in range(6):
13     img1_copy = cv2.pyrDown(img1_copy)
14     gp_img1.append(img1_copy)
15 ## generate Gaussian pyramid for lion2
16 img2_copy = img2.copy()
17 gp_img2 = [img2_copy]
18 for i in range(6):
19     img2_copy = cv2.pyrDown(img2_copy)
20     gp_img2.append(img2_copy)
21 ## generate Laplacian Pyramid for apple
22 img1_copy = gp_img1[5]
23 lp_img1 = [img1_copy]
24 for i in range(5, 0, -1):
25     gaussian_expanded = cv2.pyrUp(gp_img1[i])
26     laplacian = cv2.subtract(gp_img1[i-1], gaussian_expanded)
27     lp_img1.append(laplacian)
28 ## generate Laplacian Pyramid for orange
29 img2_copy = gp_img2[5]
30 lp_img2 = [img2_copy]
31 for i in range(5, 0, -1):
32     gaussian_expanded = cv2.pyrUp(gp_img2[i])
33     laplacian = cv2.subtract(gp_img2[i-1], gaussian_expanded)
34     lp_img2.append(laplacian)
35 ## Now add left and right halves of images in each level
36 img1_img2_pyramid = []
37 n = 0
38 for img1_lap, img2_lap in zip(lp_img1, lp_img2):
39     n += 1
40     cols, rows, ch = img1_lap.shape
41     laplacian = np.hstack((img1_lap[:, 0:int(cols/2)], img2_lap[:, int(cols/2):]))
42     img1_img2_pyramid.append(laplacian)
43 # now reconstruct
44 img1_img2_reconstruct = img1_img2_pyramid[0]
45 for i in range(1, 6):
46     img1_img2_reconstruct = cv2.pyrUp(img1_img2_reconstruct)
47     img1_img2_reconstruct = cv2.add(img1_img2_pyramid[i], img1_img2_reconstruct)
48 cv2.imshow("img1", img1)
49 cv2.imshow("img2", img2)
50 cv2.imshow("img1 img2", img1_img2_pyramid)
51 cv2.imshow("img1_img2_reconstruct", img1_img2_reconstruct)
52 cv2.waitKey(0)
53 cv2.destroyAllWindows()
```



MATLAB code for Image Blending

```
1 A = im2double(imread("0855x4.png"));
2 A = imresize(A,[256 256]);
3 B = im2double(imread("0880x4.png"));
4 B = imresize(B,[256,256]);
5 subplot(1,2,1)
6 imshow(A)
7 subplot(1,2,2)
8 imshow(B)
9
10 mask_A = zeros(256,256);
11 mask_A(:,1:128) = 1;
12
13 clf
14 imshow(mask_A)
15 xticks([])
16 yticks([])
17 axis on
18
19 C = (A .* mask_A) + (B .* (1 - mask_A));
20 clf|
21 imshow(C)
22 xticks([])
23 yticks([])
24
25 mrp_A = multiresolutionPyramid(A);
26 mrp_B = multiresolutionPyramid(B);
27 mrp_mask_A = multiresolutionPyramid(mask_A);
28
29 lap_A = laplacianPyramid(mrp_A);
30 lap_B = laplacianPyramid(mrp_B);
```

```
31
32 for k = 1:length(lap_A)
33     lap_blend{k} = (lap_A{k} .* mrp_mask_A{k}) + ...
34                 (lap_B{k} .* (1 - mrp_mask_A{k}));
35 end
36
37 C_blended = reconstructFromLaplacianPyramid(lap_blend);
38 imshow(C_blended)
39
40 x = linspace(-1,1,256);
41 y = x';
42 mask2_A = hypot(x,y) <= 0.5;
43
44 mrp_mask2_A = multiresolutionPyramid(mask2_A);
45
46 for k = 1:length(lap_A)
47     lap_blend2{k} = (lap_A{k} .* mrp_mask2_A{k}) + ...
48                 (lap_B{k} .* (1 - mrp_mask2_A{k}));
49 end
50
51 C2_blended = reconstructFromLaplacianPyramid(lap_blend2);
52 imshow(C2_blended)
```



- Pixels:
 - great for spatial resolution, poor access to frequency
- Fourier transform:
 - great for frequency, not for spatial info
- Pyramids/filter banks:
 - balance between spatial and frequency information

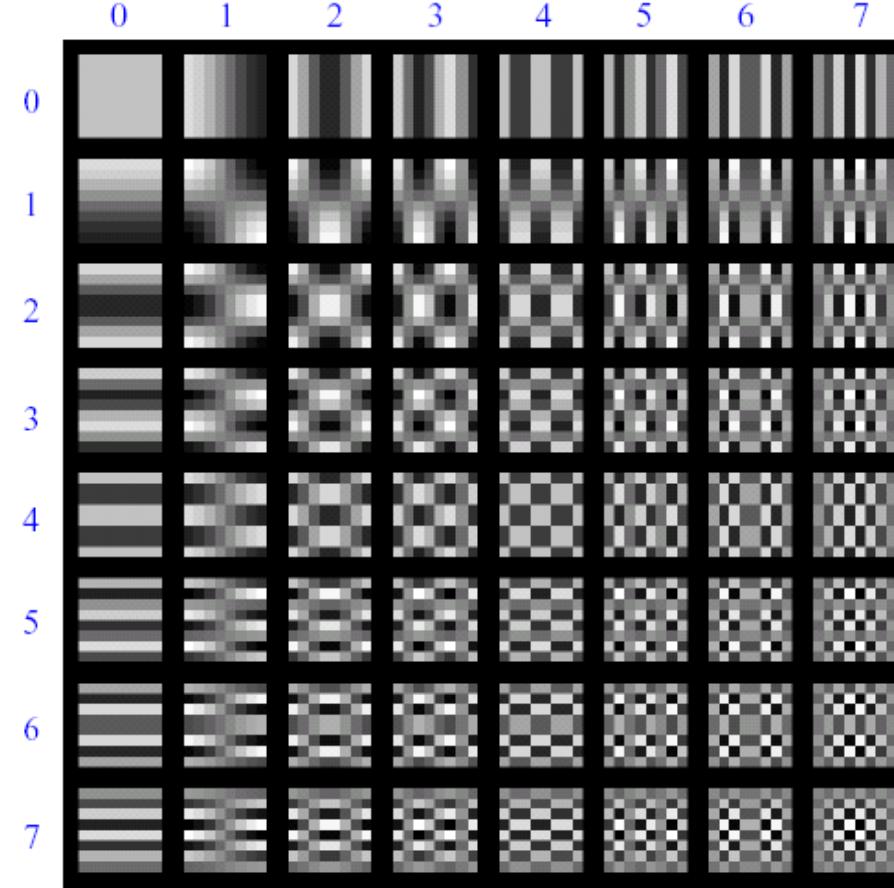
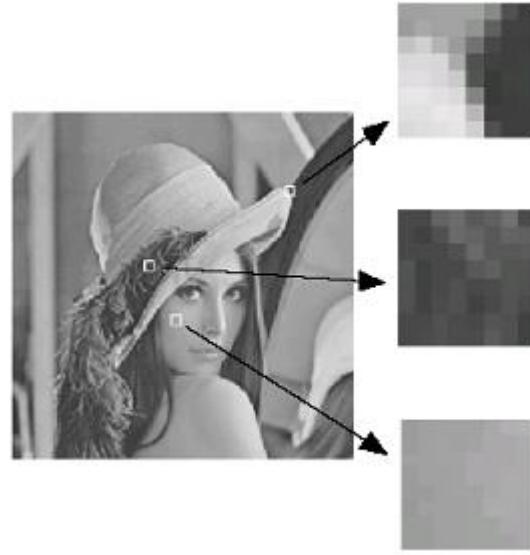


Compression

How is it that a 4MP image (12000KB) can be compressed to 400KB without a noticeable change?



Lossy Image Compression (JPEG)



Block-based Discrete Cosine Transform (DCT)

Slides: Efros

Using DCT in JPEG

- The first coefficient $B(0,0)$ is the DC component, the average intensity
- The top-left coeffs represent low frequencies, the bottom right – high frequencies

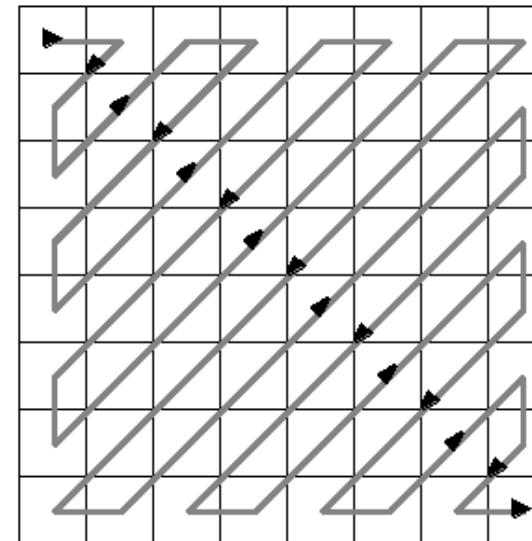
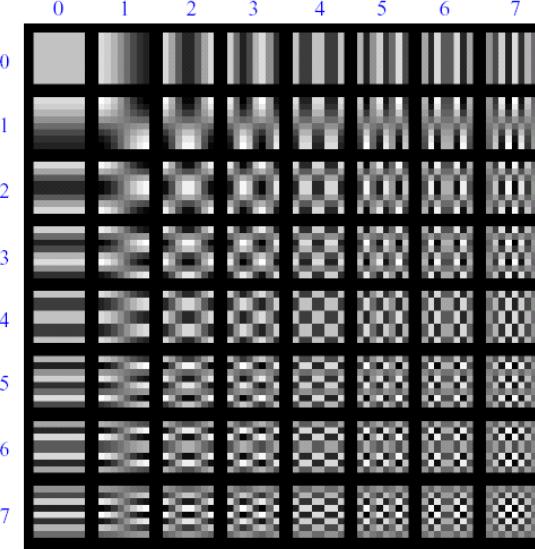


Image Compression using DCT

- Quantize
 - More coarsely for high frequencies (which also tend to have smaller values)
 - Many quantized high frequency values will be zero
- Encode
 - Can decode with inverse dct

Filter responses

$$G = \begin{bmatrix} -415.38 & -30.19 & -61.20 & 27.24 & 56.13 & -20.10 & -2.39 & 0.46 \\ 4.47 & -21.86 & -60.76 & 10.25 & 13.15 & -7.09 & -8.54 & 4.88 \\ -46.83 & 7.37 & 77.13 & -24.56 & -28.91 & 9.93 & 5.42 & -5.65 \\ -48.53 & 12.07 & 34.10 & -14.76 & -10.24 & 6.30 & 1.83 & 1.95 \\ 12.12 & -6.55 & -13.20 & -3.95 & -1.88 & 1.75 & -2.79 & 3.14 \\ -7.73 & 2.91 & 2.38 & -5.94 & -2.38 & 0.94 & 4.30 & 1.85 \\ -1.03 & 0.18 & 0.42 & -2.42 & -0.88 & -3.02 & 4.12 & -0.66 \\ -0.17 & 0.14 & -1.07 & -4.19 & -1.17 & -0.10 & 0.50 & 1.68 \end{bmatrix}$$

Quantized values

$$B = \begin{bmatrix} -26 & -3 & -6 & 2 & 2 & -1 & 0 & 0 \\ 0 & -2 & -4 & 1 & 1 & 0 & 0 & 0 \\ -3 & 1 & 5 & -1 & -1 & 0 & 0 & 0 \\ -3 & 1 & 2 & -1 & 0 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

→



Quantization table

$$Q = \begin{bmatrix} 16 & 11 & 10 & 16 & 24 & 40 & 51 & 61 \\ 12 & 12 & 14 & 19 & 26 & 58 & 60 & 55 \\ 14 & 13 & 16 & 24 & 40 & 57 & 69 & 56 \\ 14 & 17 & 22 & 29 & 51 & 87 & 80 & 62 \\ 18 & 22 & 37 & 56 & 68 & 109 & 103 & 77 \\ 24 & 35 & 55 & 64 & 81 & 104 & 113 & 92 \\ 49 & 64 & 78 & 87 & 103 & 121 & 120 & 101 \\ 72 & 92 & 95 & 98 & 112 & 100 & 103 & 99 \end{bmatrix}$$



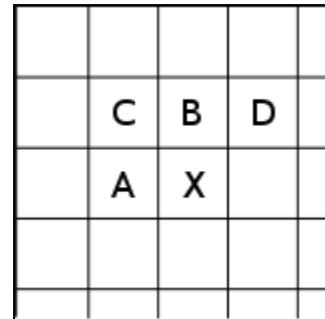
JPEG Compression Summary

1. Convert image to YCrCb
2. Subsample color by factor of 2
 - People have bad resolution for color
3. Split into blocks (8x8, typically), subtract 128
4. For each block
 - a. Compute DCT coefficients
 - b. Coarsely quantize
 - Many high frequency components will become zero
 - c. Encode (e.g., with Huffman coding)



Lossless Compression (PNG)

1. Predict that a pixel's value based on its upper-left neighborhood
2. Store difference of predicted and actual value
3. Pkzip it (DEFLATE algorithm)



Three views of Image Filtering

- Image filters in spatial domain
 - Filter is a mathematical operation on values of each patch
 - Smoothing, sharpening, measuring texture
- Image filters in the frequency domain
 - Filtering is a way to modify the frequencies of images
 - Denoising, sampling, image compression
- Templates and Image Pyramids
 - Filtering is a way to match a template to the image
 - Detection, coarse-to-fine registration

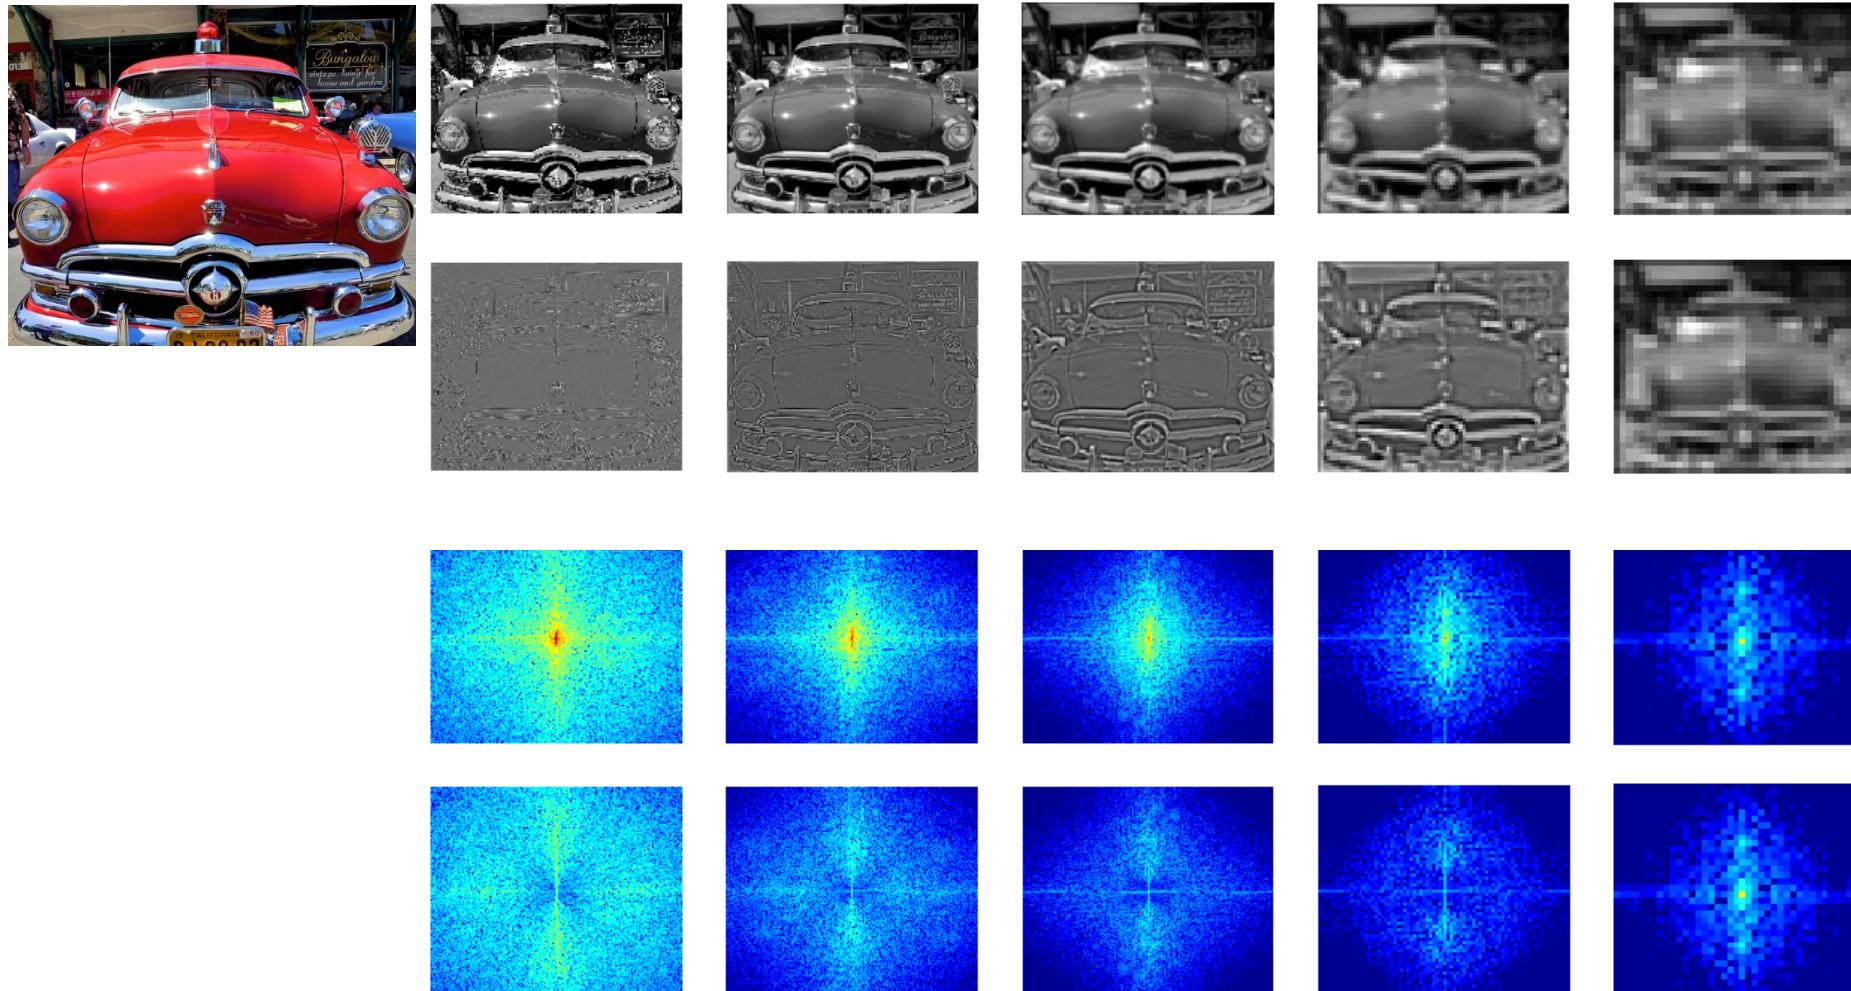


Complex Engineering Task



- Hybrid image =
 $\text{Low-Freq(Image A)} + \text{Hi-Freq(Image B)}$

Complex Engineering Task



Complex Engineering Task



Some Image Denoising Results



Things to Remember

- Template matching (SSD or Normxcorr2)
 - SSD can be done with linear filters, is sensitive to overall intensity
- Gaussian pyramid
 - Coarse-to-fine search, multi-scale detection
- Laplacian pyramid
 - More compact image representation
 - Can be used for compositing in graphics
- Compression
 - In JPEG, coarsely quantize high frequencies

