

# Interest Points

Dr. Sander Ali Khawaja,

Assistant Professor, Department of Telecommunication Engineering  
Faculty of Engineering and Technology, University of Sindh, Pakistan

Senior Member, IEEE – Member, ACM

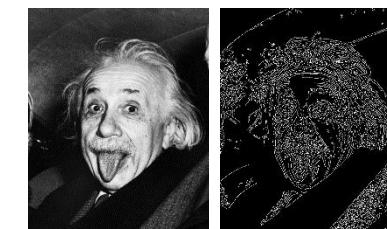
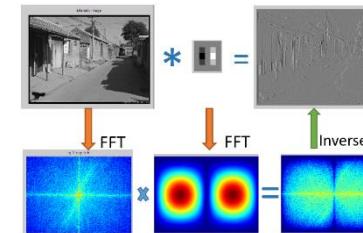
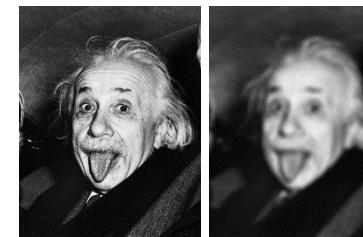
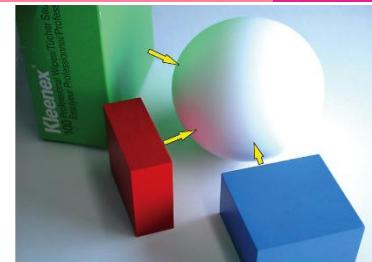
<https://sander-ali.github.io>

# Computer Vision & Image Processing



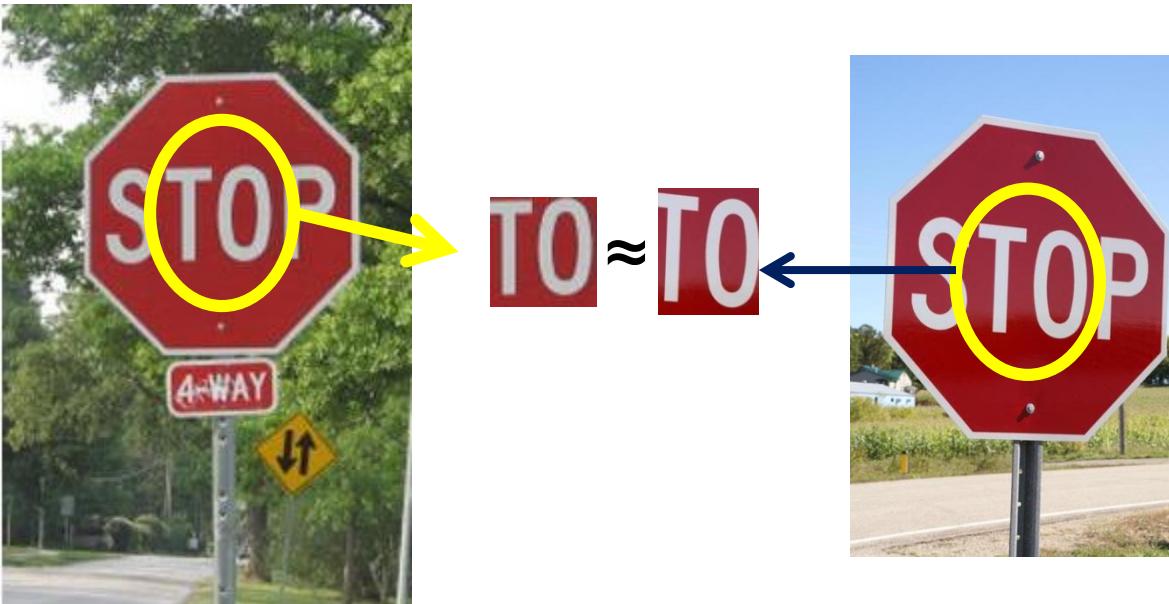
# What we have learned so far

- Light and color
  - What an image records
- Filtering in spatial domain
  - Filtering = weighted sum of neighboring pixels
  - Smoothing, sharpening, measuring texture
- Filtering in frequency domain
  - Filtering = change frequency of the input image
  - Denoising, sampling, image compression
- Image pyramid and template matching
  - Filtering = a way to find a template
  - Image pyramids for coarse-to-fine search and multi-scale detection
- Edge detection
  - Canny edge = smooth -> derivative -> thin -> threshold -> link
  - Finding straight lines, binary image analysis



# This Module: Correspondence and Alignment

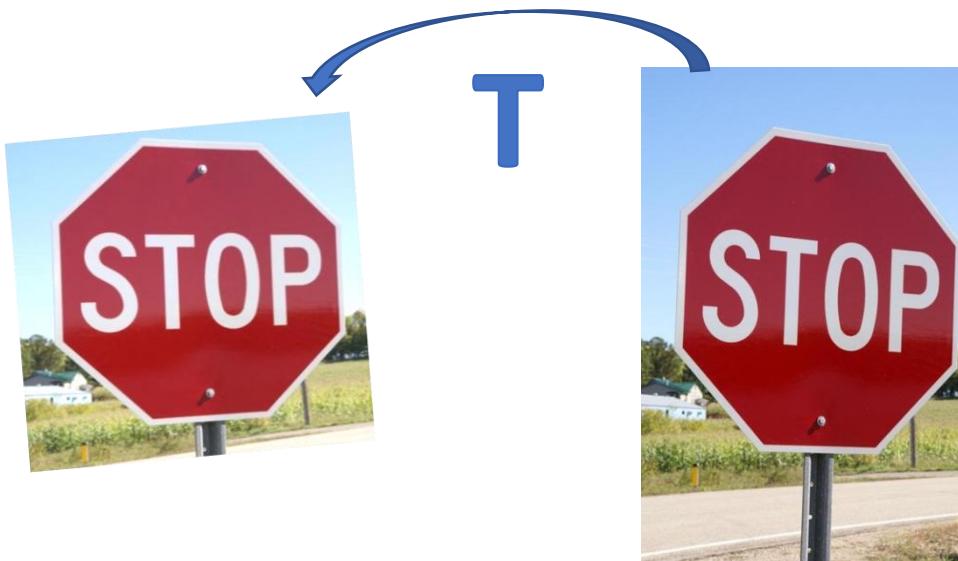
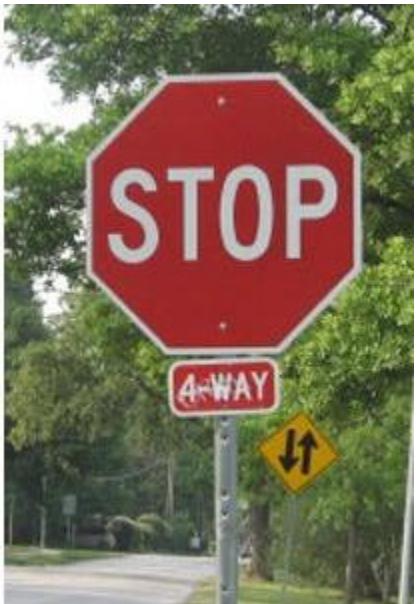
- **Correspondence:**  
matching points, patches, edges, or regions across images



Slide credit: Derek Hoiem

# This module: Correspondence and Alignment

- **Alignment/registration:**  
solving the transformation that makes two things  
match better



Slide credit: Derek Hoiem

## 1. Registration

## 2. Registration

## 3. Registration



Takeo Kanade (CMU)

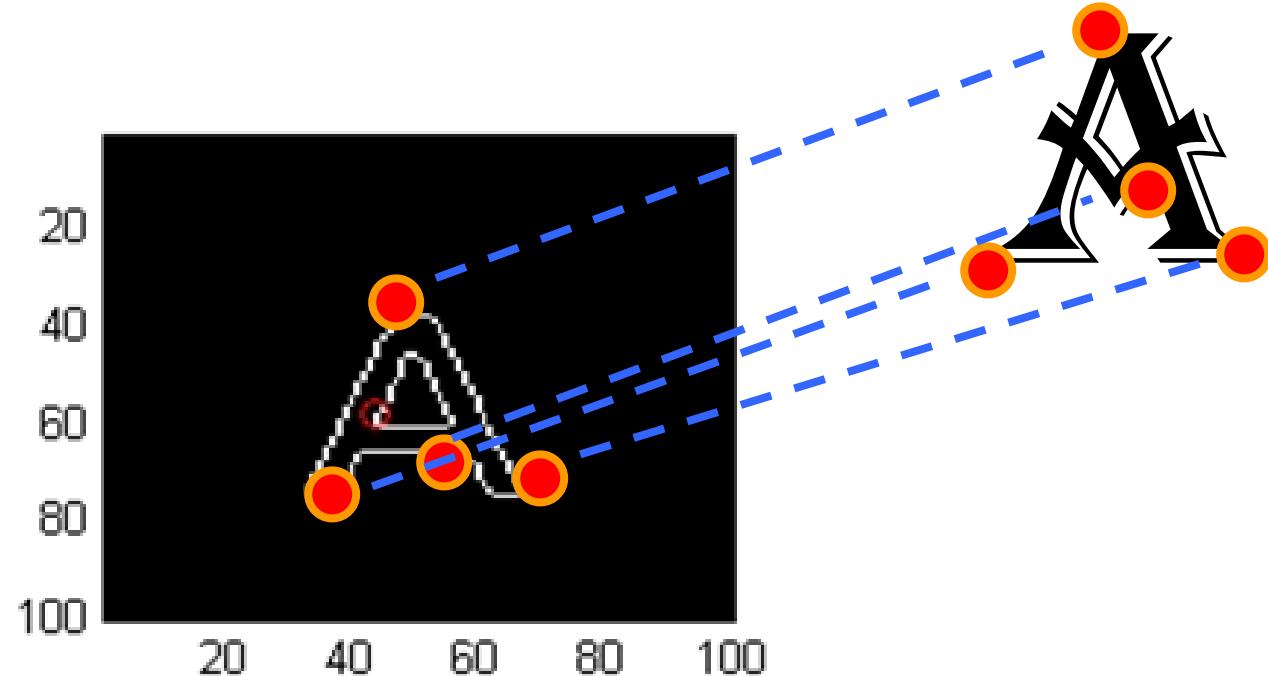
# Automatic Panoramas



Credit: Matt Brown

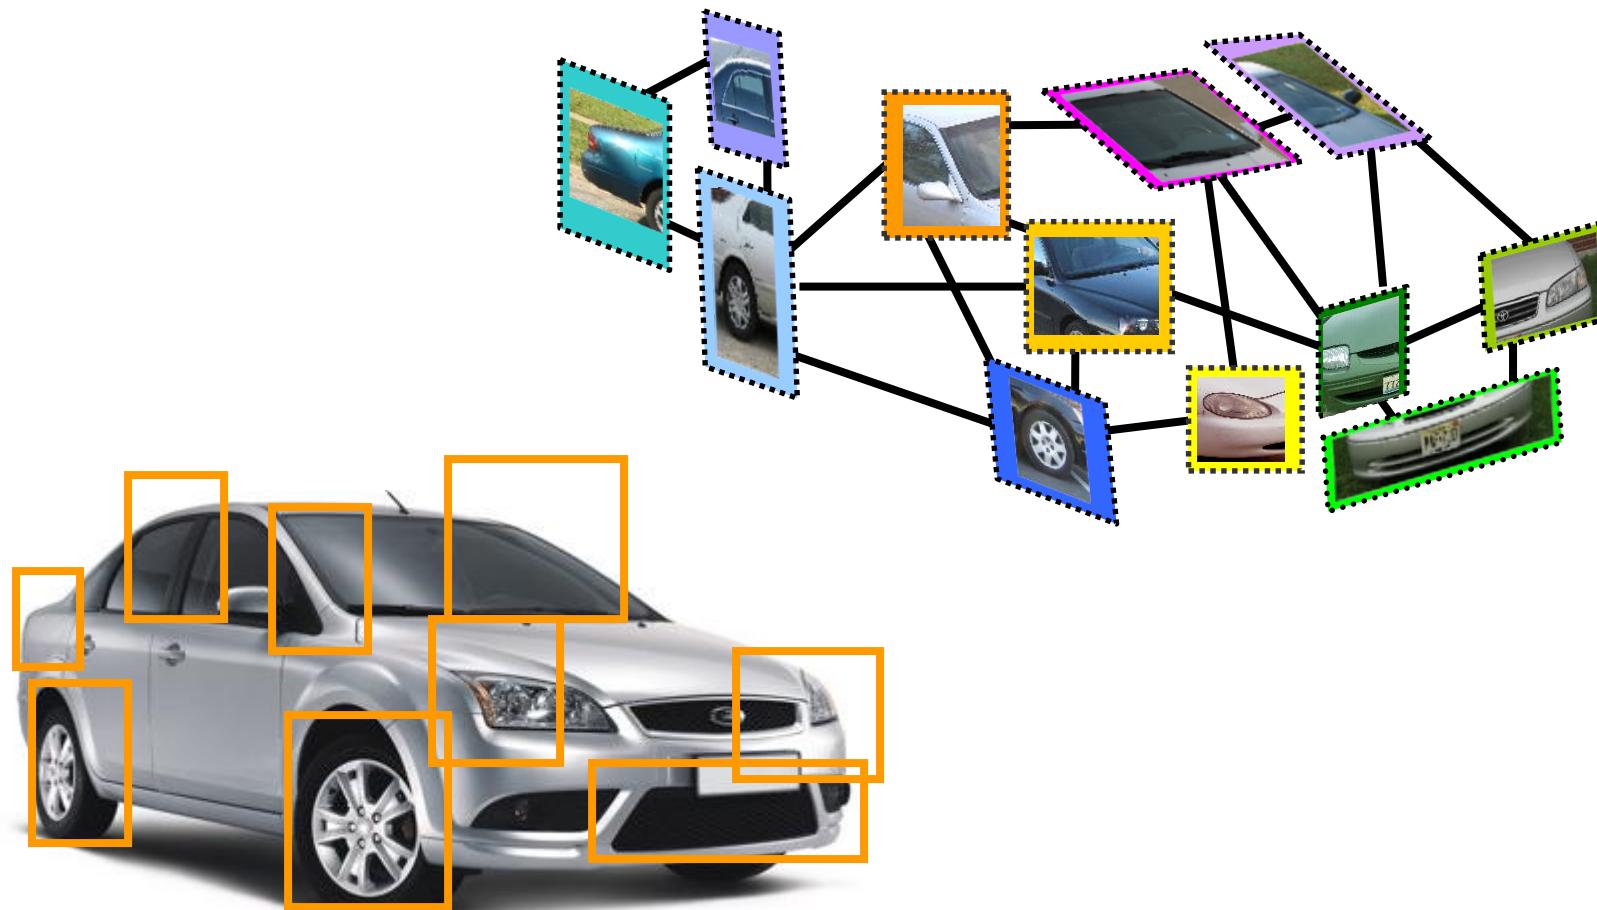
Dr. Sander Ali Khowaja

# Fitting a 2D shape template



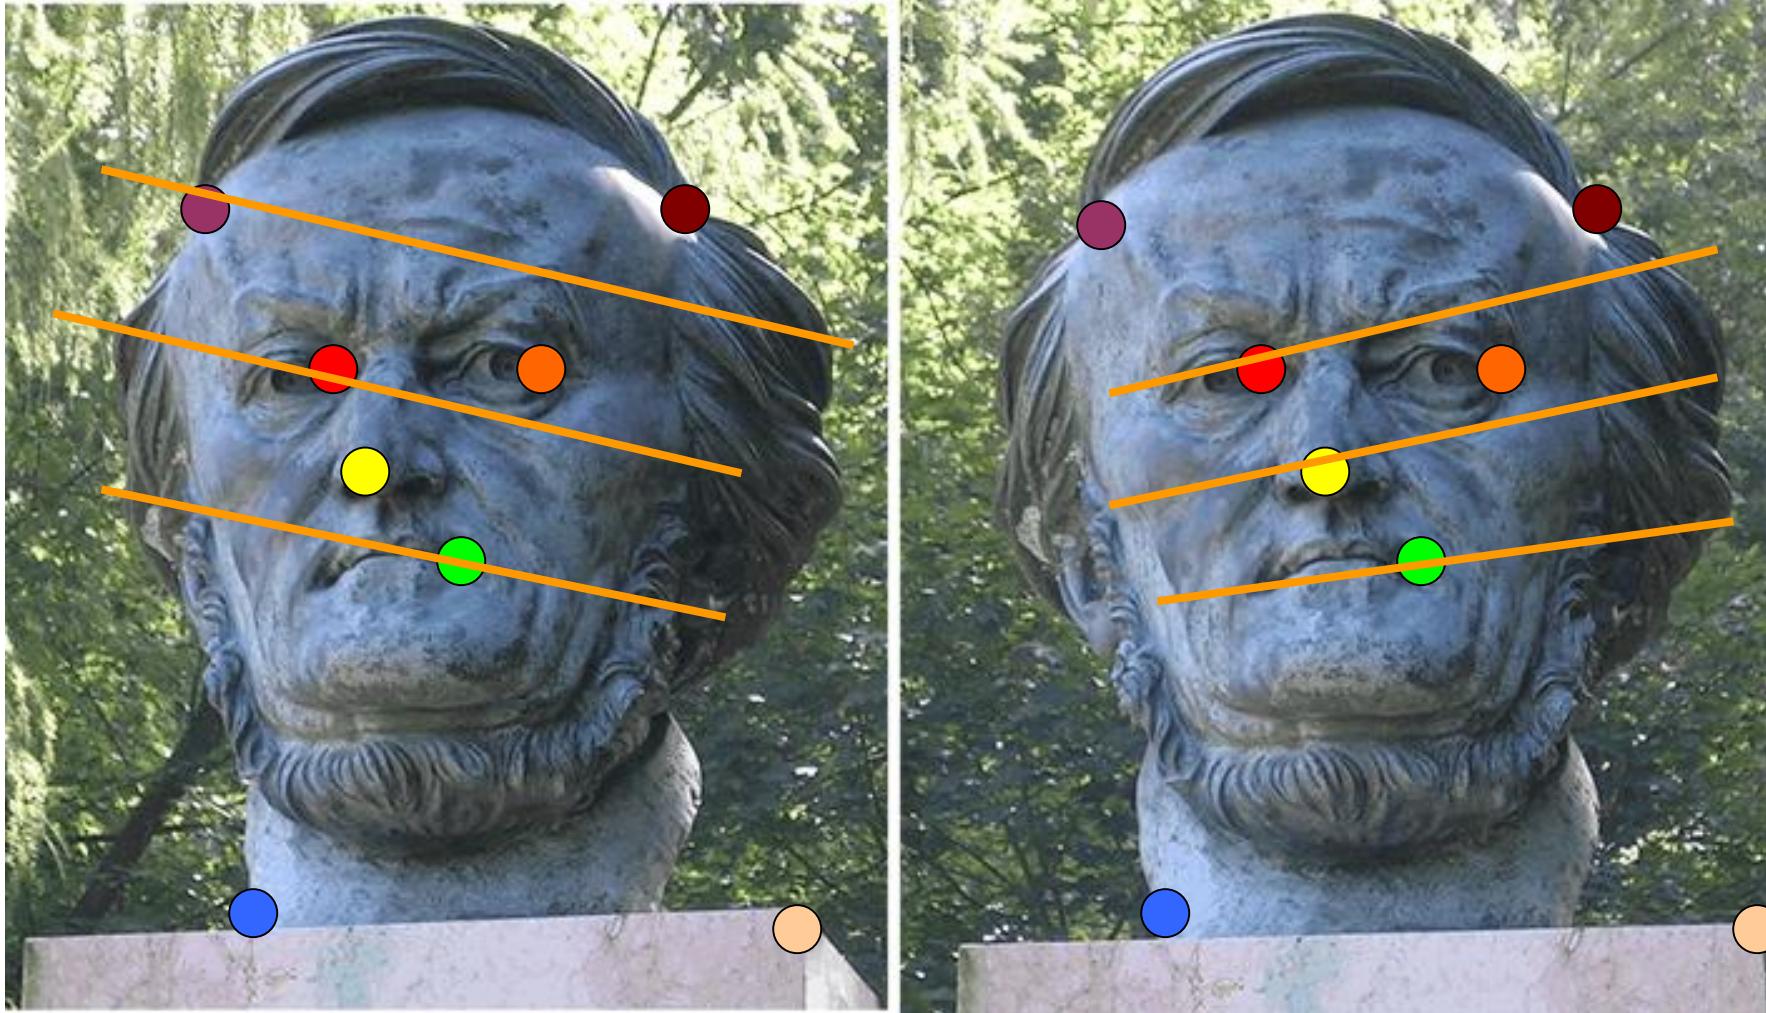
Slide from Silvio Savarese

# Fitting a 3D object model



Slide from Silvio Savarese

# Estimating Fundamental Matrix that Correspond two Views

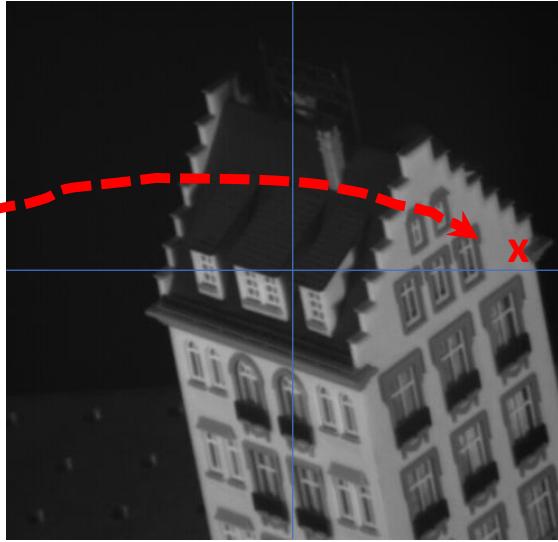


Slide from Silvio Savarese

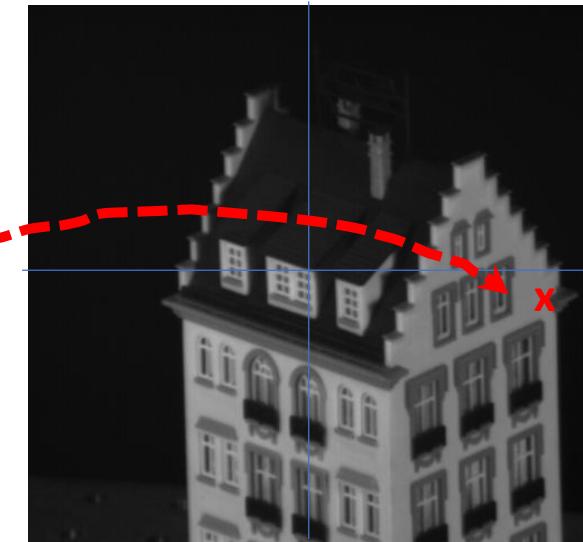
# Tracking Points



frame 0



frame 22



frame 49

# This Lecture

- What is interest point?
- Corner detection
- Handling scale and orientation
- Feature matching



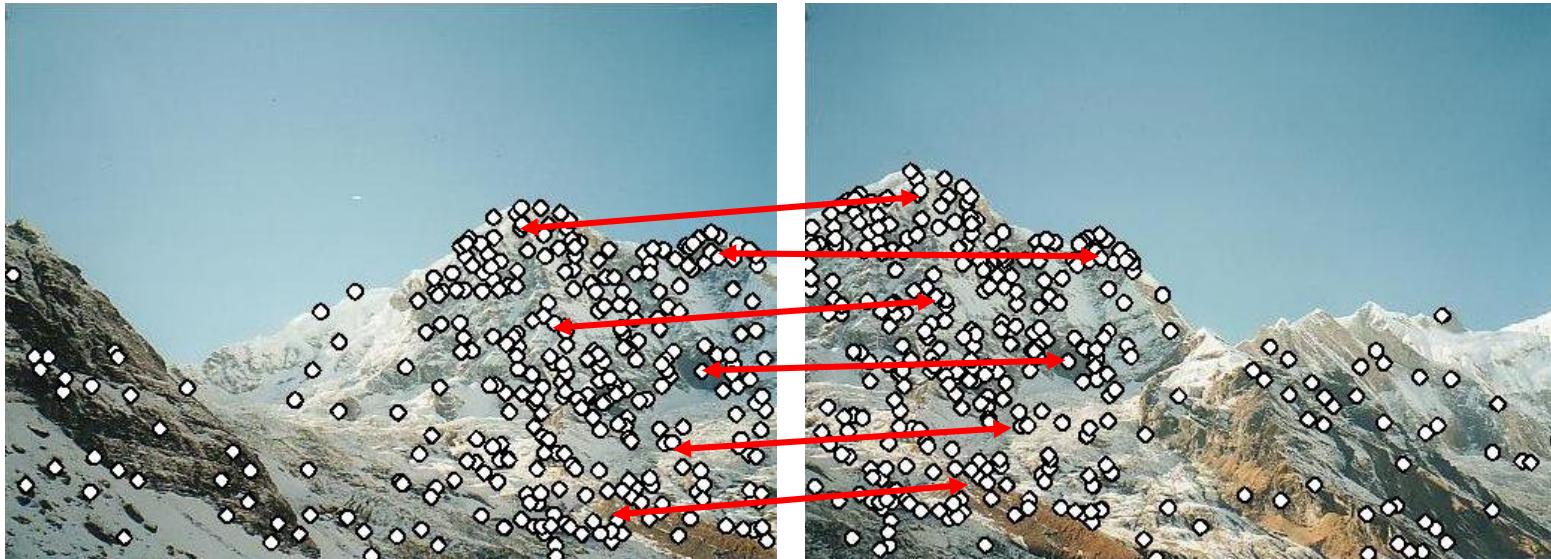
# Why Extract Features

- Motivation: panorama stitching
  - We have two images – how do we combine them?



# Why Extract Features

- Motivation: panorama stitching
  - We have two images – how do we combine them?

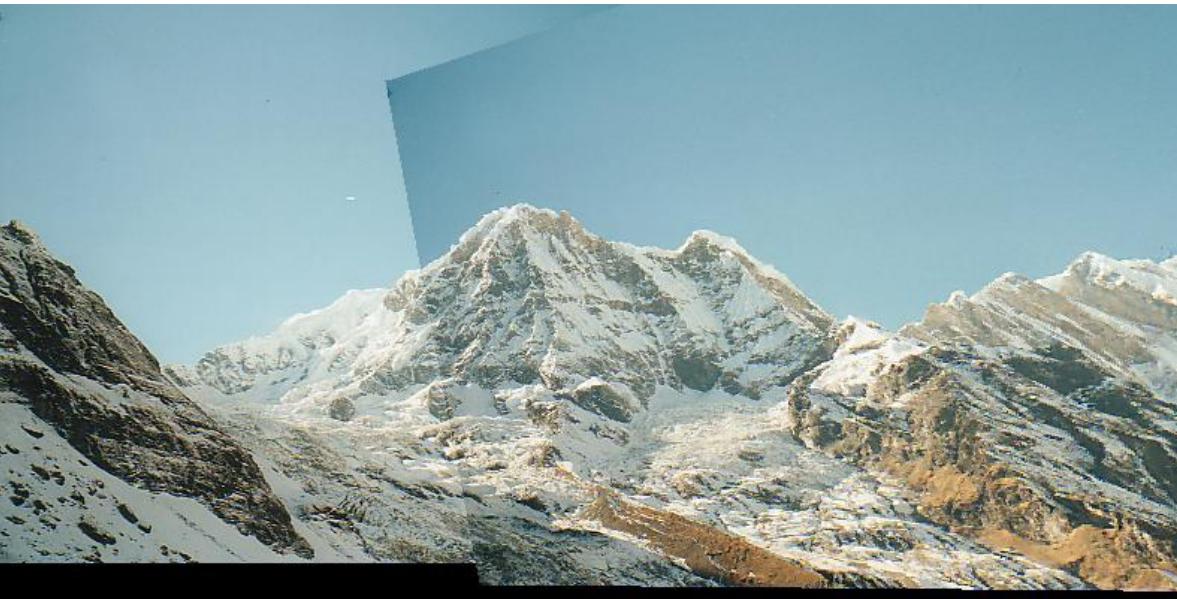


Step 1: extract features

Step 2: match features

# Why Extract Features

- Motivation: panorama stitching
  - We have two images – how do we combine them?



- Step 1: extract features
- Step 2: match features
- Step 3: align images

# Image Matching



by [Diva Sian](#)



by [swashford](#)

# Image Matching: Harder Case



by [Diva Sian](#)

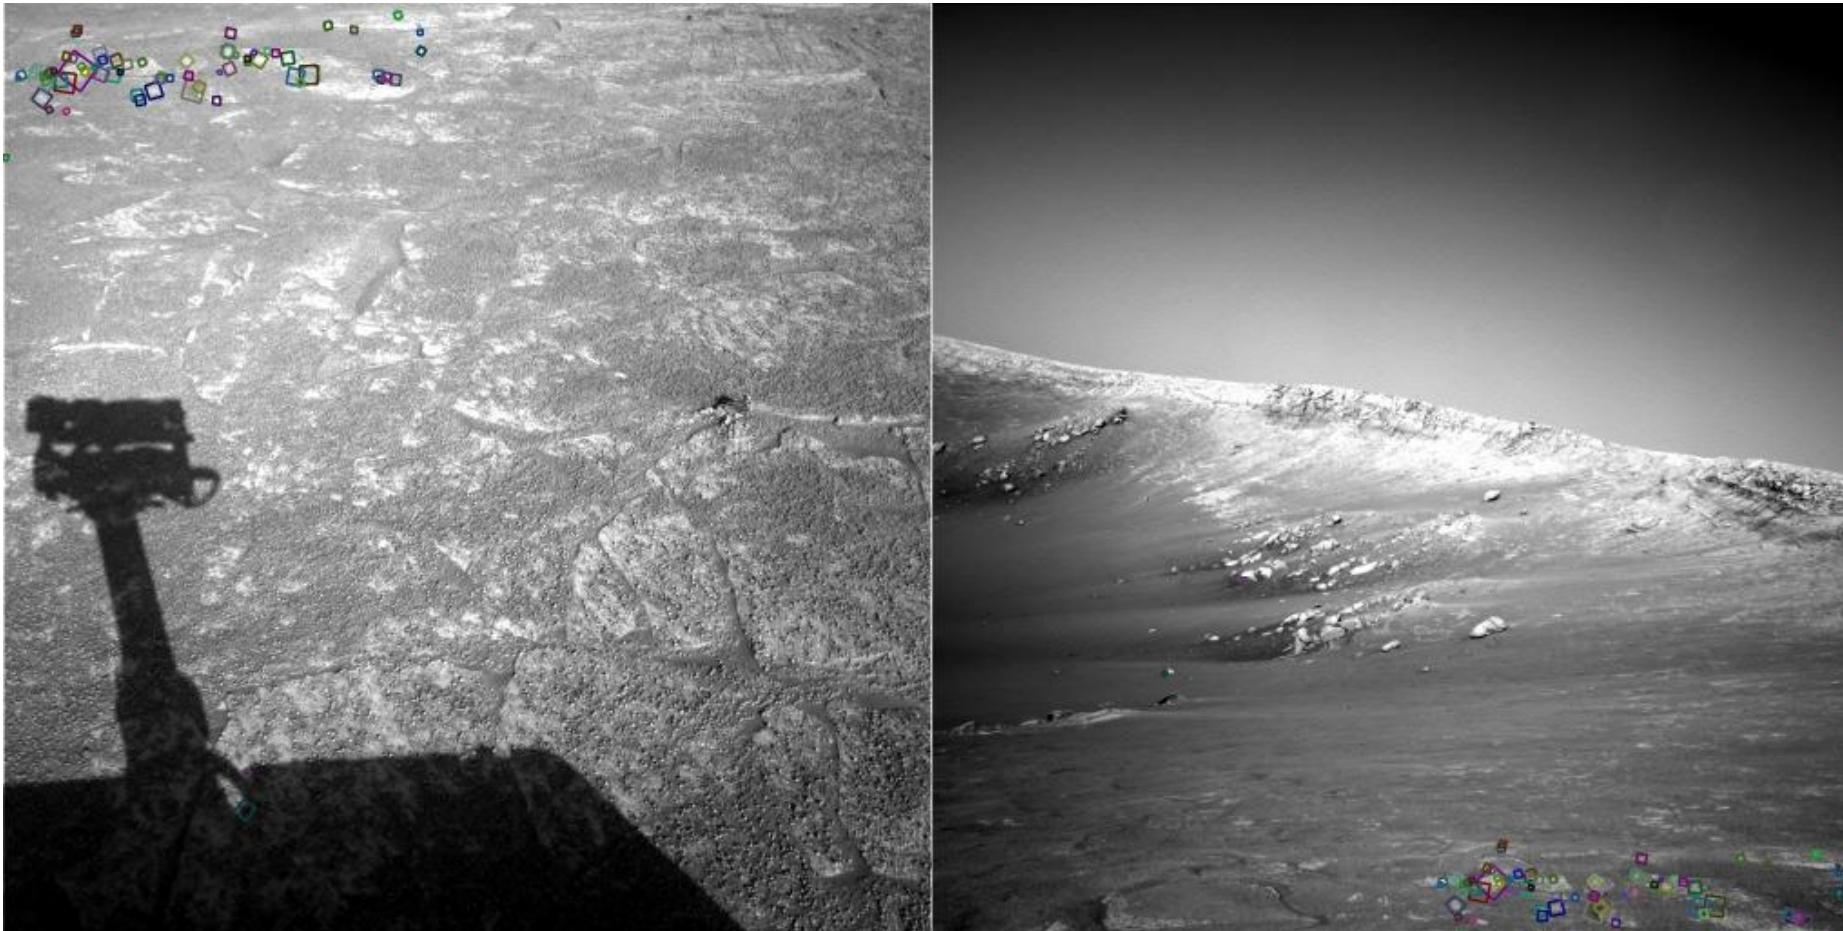


by [scgbt](#)

# What do you think about this



# Answer: Look for Tiny Squares



NASA Mars Rover images  
with SIFT feature matches

# Applications

- Keypoints are used for:
  - Image alignment
  - 3D reconstruction
  - Motion tracking
  - Robot navigation
  - Indexing and database retrieval
  - Object recognition



# Advantages of Local Features

## Locality

- features are local, so robust to occlusion and clutter

## Quantity

- hundreds or thousands in a single image

## Distinctiveness:

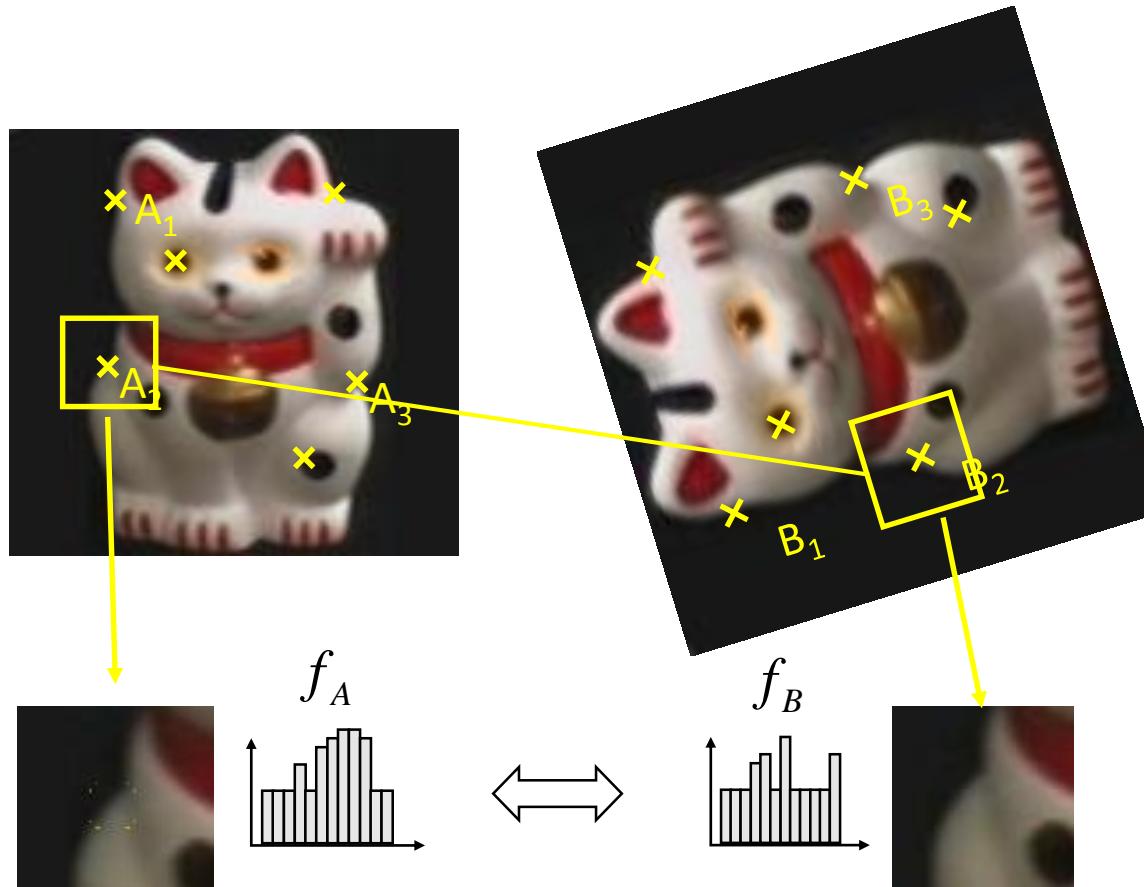
- can differentiate a large database of objects

## Efficiency

- real-time performance achievable



# Overview of Keypoint Matching



$$d(f_A, f_B) < T$$

K. Grauman, B. Leibe

1. Find a set of distinctive keypoints
2. Define a region around each keypoint
3. Extract and normalize the region content
4. Compute a local descriptor from the normalized region
5. Match local descriptors

# Goals for Keypoints



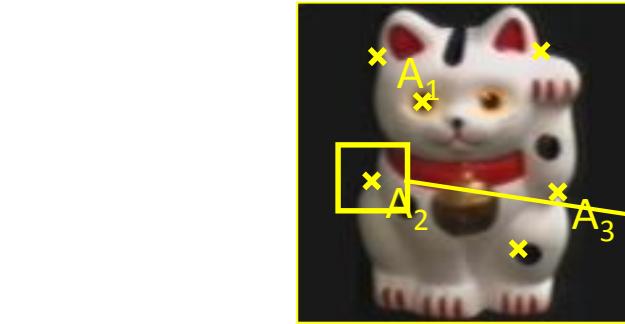
Detect points that are *repeatable* and *distinctive*

# Key Trade-offs

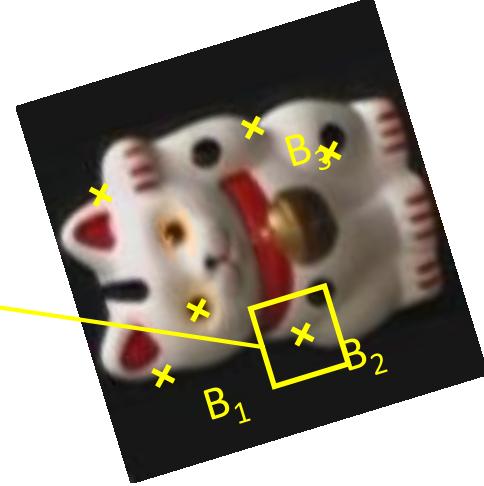
## Detection



More Repeatable  
Robust detection  
Precise localization



More Points  
Robust to occlusion  
Works with less texture



## Description



More Distinctive  
Minimize wrong matches

More Flexible  
Robust to expected variations  
Maximize correct matches

# Choosing Interest Points

Where would you tell  
your friend to meet  
you?



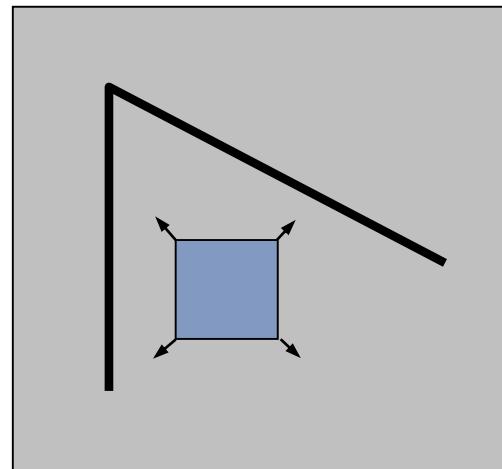
# Choosing Interest Points

Where would you tell  
your friend to meet  
you?

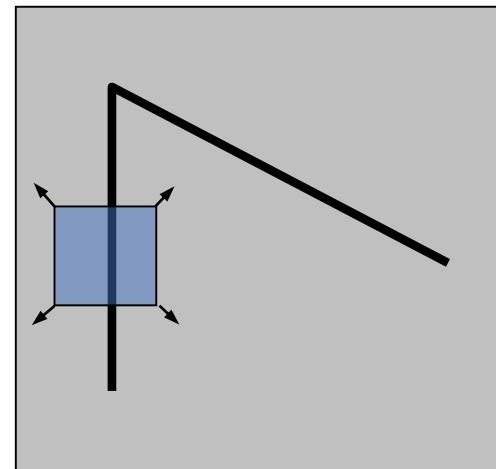


# Corner Detection: Basic Idea

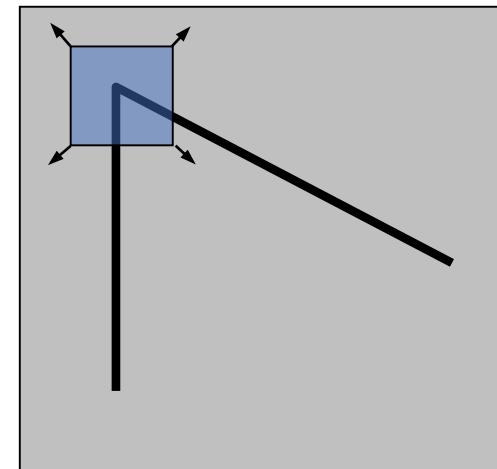
- How does the window change when you shift it?
- Shifting the window in any direction causes a big change



“flat” region:  
no change in all  
directions



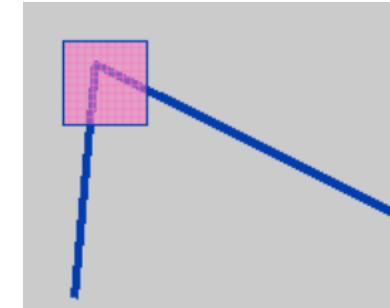
“edge”:  
no change along the  
edge direction



“corner”:  
significant change in  
all directions

Credit: S. Seitz, D. Frolova, D. Simakov

- Point whose local neighborhood stands in two dominant and different edge directions.



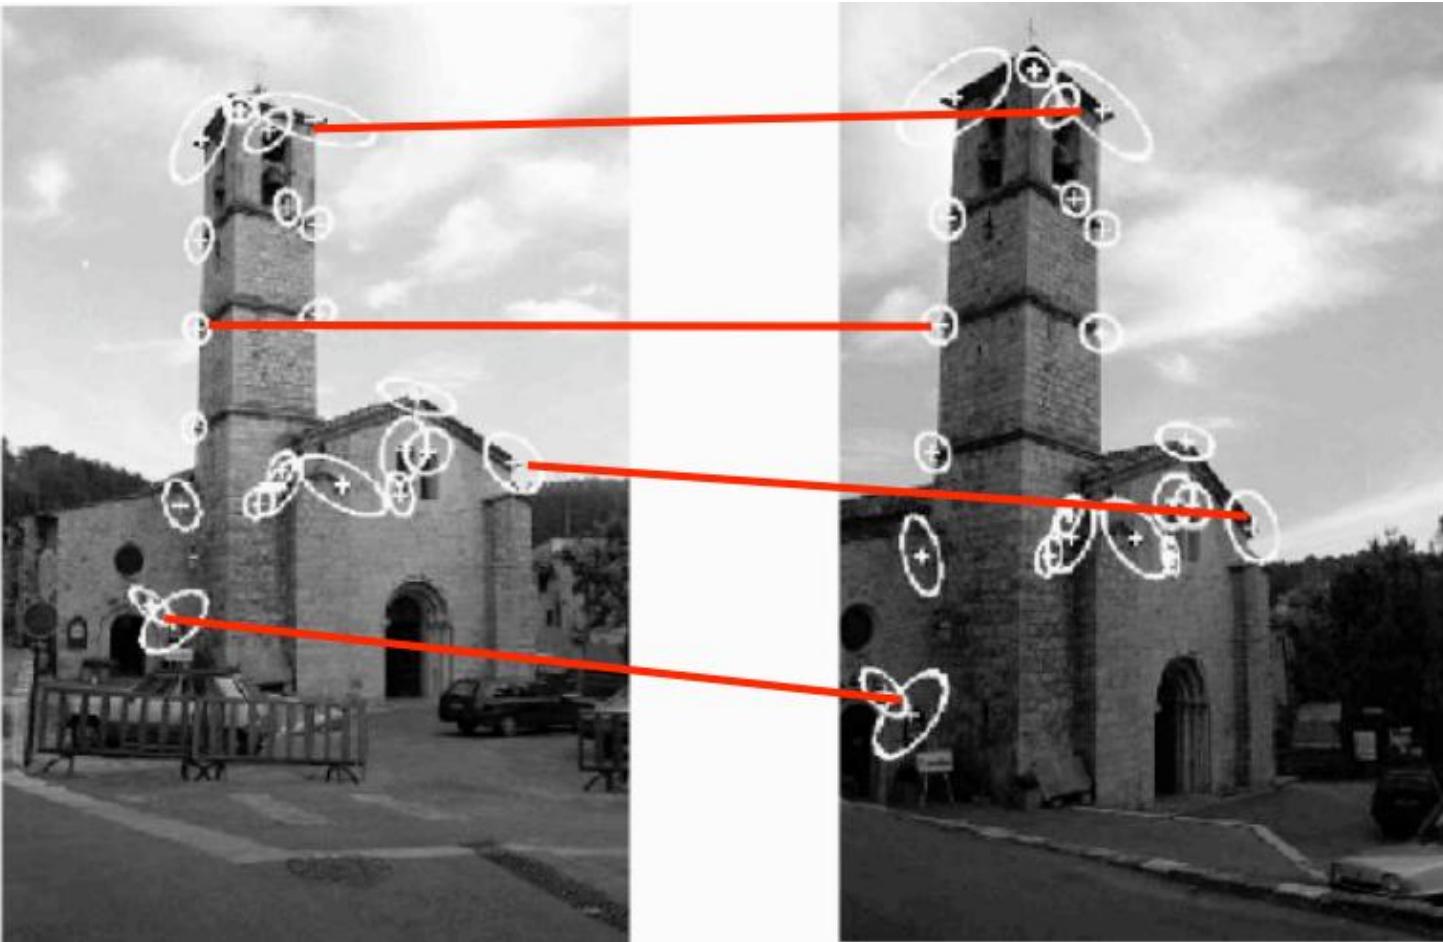
- Junction of two edges.

- Corners are mostly termed as interest points as they are invariant to translation, rotation, and illumination.

- Corners are used as a basic building block for various applications such as Motion tracking, image stitching, 2D mosaics, stereo vision, image representation, and so forth.

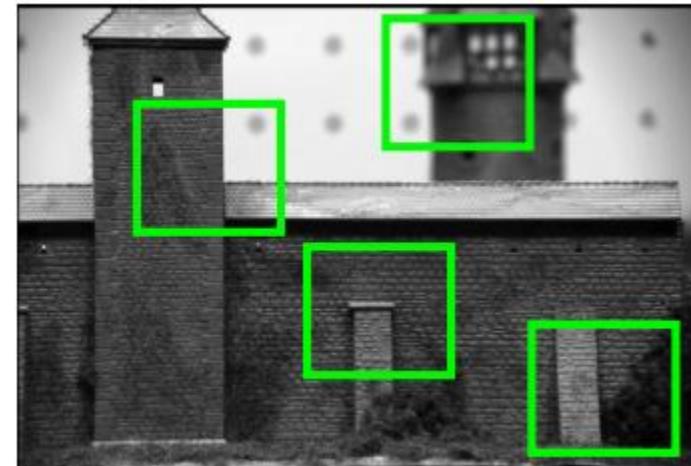
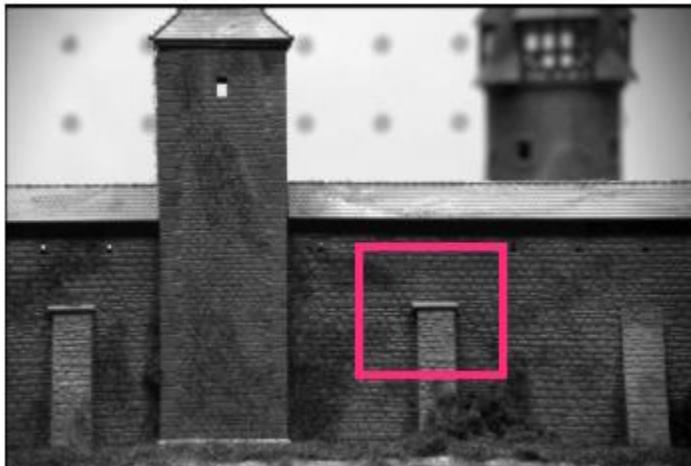
# Motivation – Matching Problem

- Vision tasks such as stereo and motion estimation require finding corresponding features across two or more views.

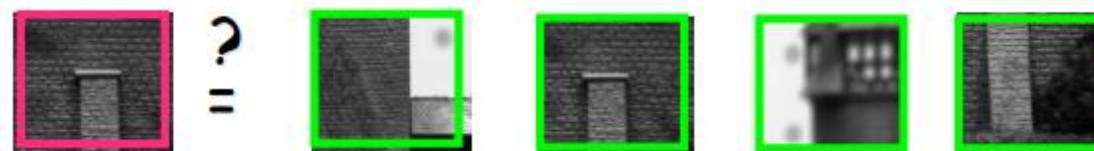


# Motivation – Patch Matching

- Elements to be matched are image patches of fixed size

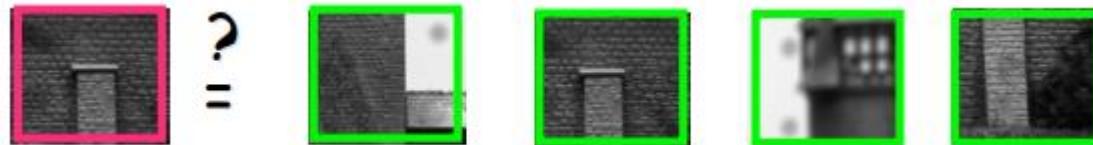
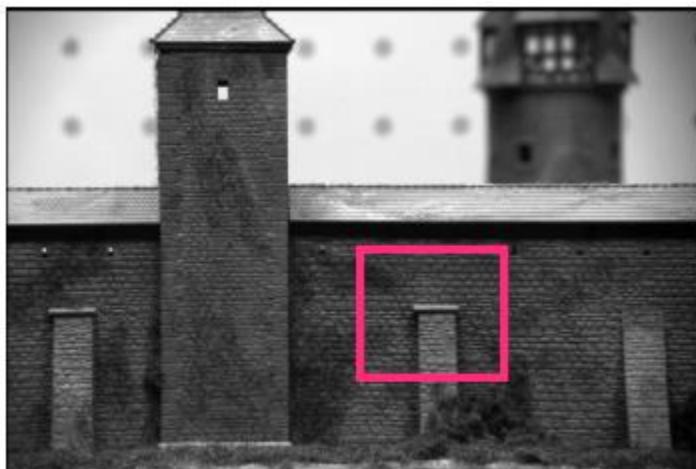


- Task: find the best (most similar) patch in the second image.



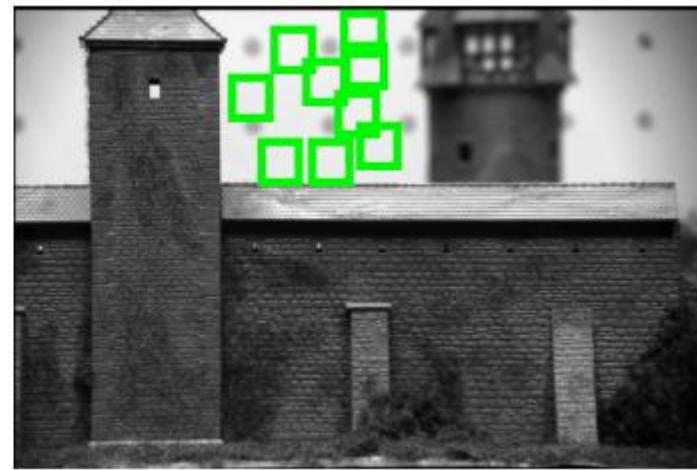
# Not all patches are created equal

- Intuition: This would be a good patch for matching, since it is very distinctive (there is only one patch in the second frame that looks similar).



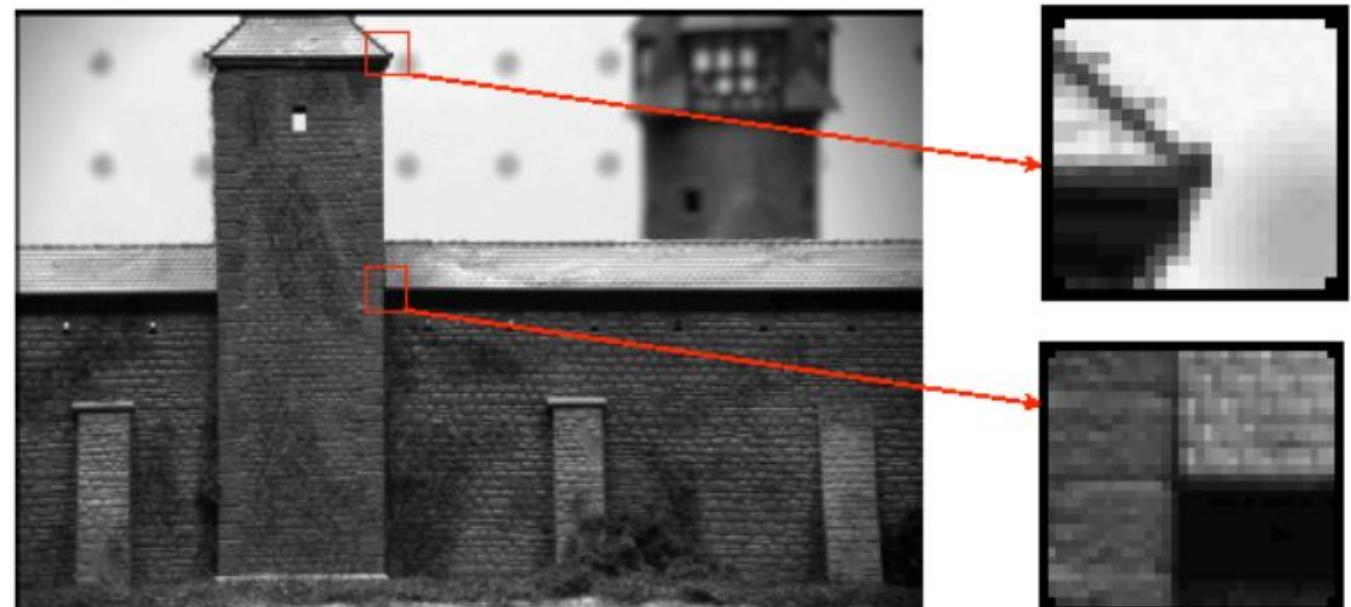
# Not all patches are created equal

- Intuition: This would be a Bad patch for matching, since it is not very distinctive (there are many similar patches in the second frame)



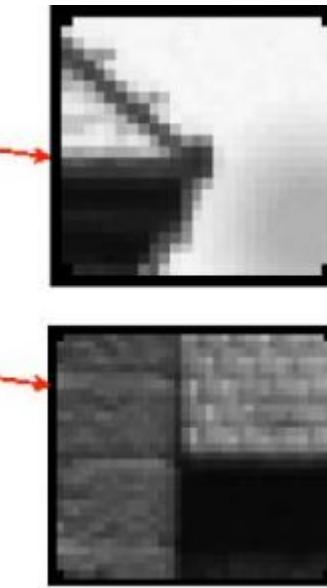
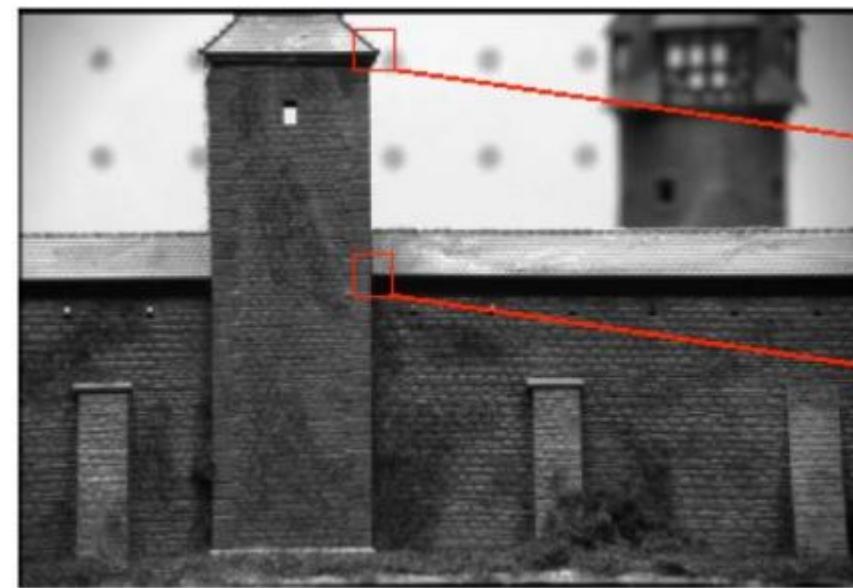
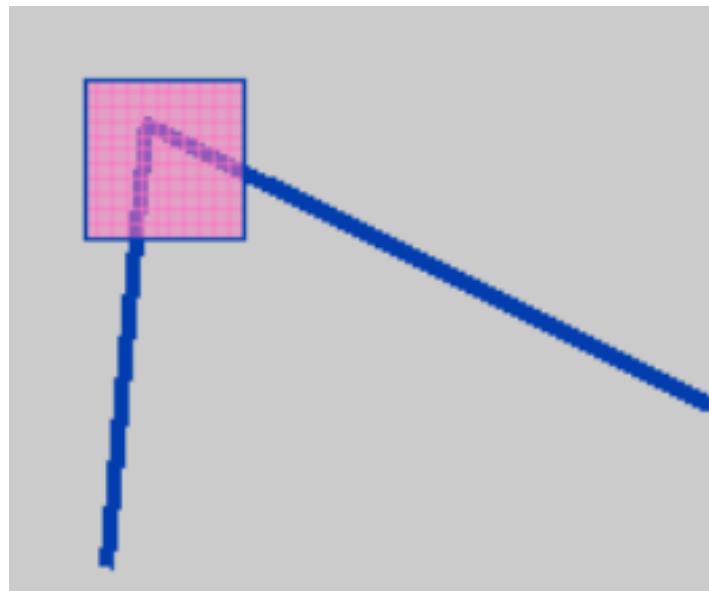
# Why Corners?

- Intuitively, junctions of contours
- Generally more stable features over changes of viewpoint.
- Intuitively, large variations in the neighborhood of the point in all directions.
- Simply, there are good features to match.



# Corner Basic Idea

- We should easily recognize the point by looking at intensity values within a small window.
- Shifting the window in any direction should yield a large change in appearance.



# Sum of Squared Differences

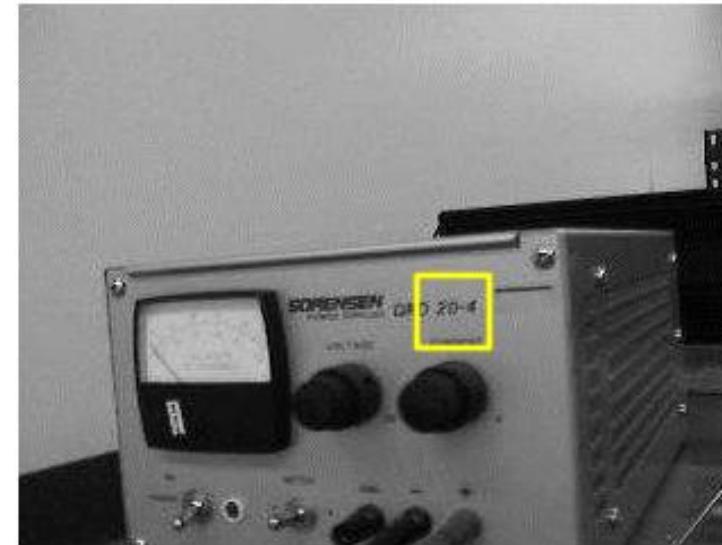
- Say we want to track a patch from image  $I_0$  to  $I_1$

- We'll assume that intensities don't change, and only translational motion
  - So we expect that

$$I_1(\mathbf{x}_i + \mathbf{u}) = I_0(\mathbf{x}_i)$$

- In practice we will have noise, so they won't be exactly equal
- But we can search for the displacement  $\mathbf{u} = (x, y)$  that minimizes the sum of squared differences

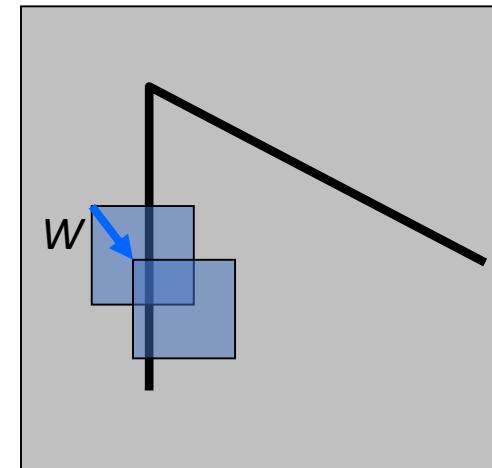
$$E(\mathbf{u}) = \sum_i w(\mathbf{x}_i) [I_1(\mathbf{x}_i + \mathbf{u}) - I_0(\mathbf{x}_i)]^2$$



# Harris Corner Detection

Consider shifting the window  $W$  by  $(u, v)$

- how do the pixels in  $W$  change?
- compare each pixel before and after by summing up the squared differences (SSD)
- this defines an SSD “error”  $E(u, v)$ :



$$E(u, v) = \sum_{(x,y) \in W} (I(x + u, y + v) - I(x, y))^2$$

# Harris Corner Detection Math: Small Motion Assumption

Taylor Series expansion of  $I$ :

$$I(x+u, y+v) = I(x, y) + \frac{\partial I}{\partial x}u + \frac{\partial I}{\partial y}v + \text{higher order terms}$$

If the motion  $(u, v)$  is small, then first order approximation is good

$$\begin{aligned} I(x + u, y + v) &\approx I(x, y) + \frac{\partial I}{\partial x}u + \frac{\partial I}{\partial y}v \\ &\approx I(x, y) + [I_x \ I_y] \begin{bmatrix} u \\ v \end{bmatrix} \end{aligned}$$

shorthand:  $I_x = \frac{\partial I}{\partial x}$

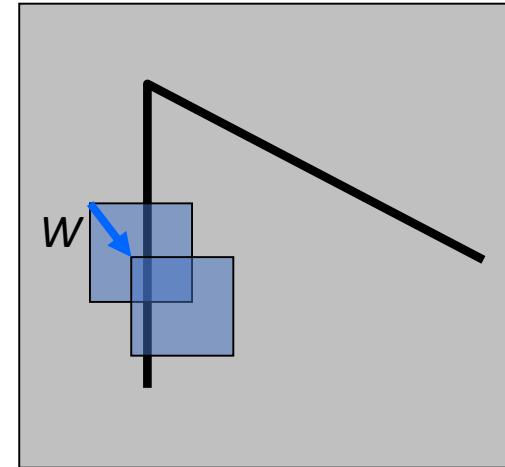
Plugging this into the formula on the previous slide...



# Harris Corner Detection: The Math

Using the small motion assumption,  
replace  $I$  with a linear approximation

(Shorthand:  $I_x = \frac{\partial I}{\partial x}$  )

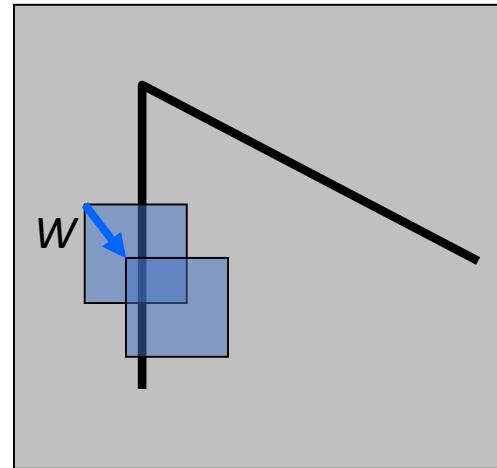


$$\begin{aligned} E(u, v) &= \sum_{(x,y) \in W} (I(x+u, y+v) - I(x, y))^2 \\ &\approx \sum_{(x,y) \in W} (I(x, y) + I_x(x, y)u + I_y(x, y)v - I(x, y))^2 \\ &\approx \sum_{(x,y) \in W} (I_x(x, y)u + I_y(x, y)v)^2 \end{aligned}$$

# Corner Detection: The Math

$$\begin{aligned} E(u, v) &\approx \sum_{(x,y) \in W} (I_x(x, y)u + I_y(x, y)v)^2 \\ &\approx \sum_{(x,y) \in W} (I_x^2 u^2 + 2I_x I_y uv + I_y^2 v^2) \\ &\approx Au^2 + 2Buv + Cv^2 \\ A &= \sum_{(x,y) \in W} I_x^2 \quad B = \sum_{(x,y) \in W} I_x I_y \quad C = \sum_{(x,y) \in W} I_y^2 \end{aligned}$$

- Thus,  $E(u, v)$  is locally approximated as a *quadratic form*



# Second Order Moment

The surface  $E(u,v)$  is locally approximated by a quadratic form.

$$E(u, v) \approx Au^2 + 2Buv + Cv^2$$

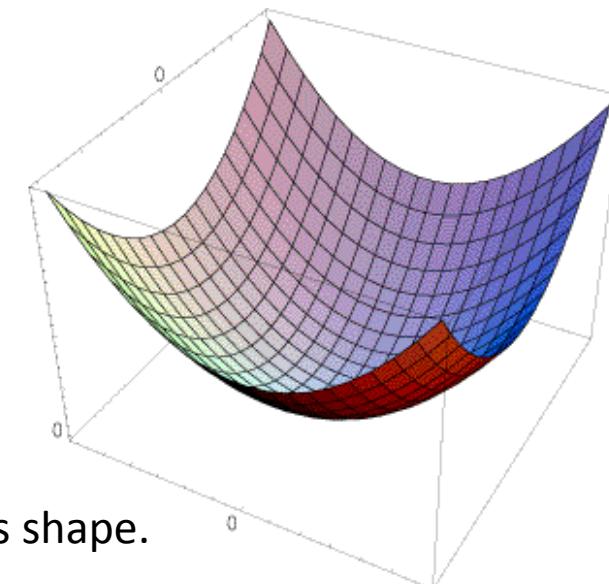
$$\approx \begin{bmatrix} u & v \end{bmatrix} \underbrace{\begin{bmatrix} A & B \\ B & C \end{bmatrix}}_H \begin{bmatrix} u \\ v \end{bmatrix}$$

$$A = \sum_{(x,y) \in W} I_x^2$$

$$B = \sum_{(x,y) \in W} I_x I_y$$

$$C = \sum_{(x,y) \in W} I_y^2$$

Let's try to understand its shape.



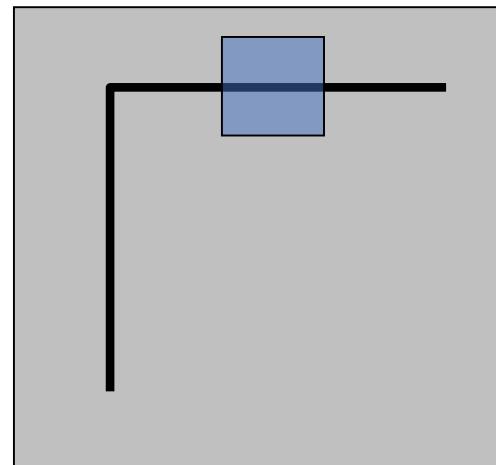
# Corner Detection: The Math

$$E(u, v) \approx \begin{bmatrix} u & v \end{bmatrix} \underbrace{\begin{bmatrix} A & B \\ B & C \end{bmatrix}}_H \begin{bmatrix} u \\ v \end{bmatrix}$$

$$A = \sum_{(x,y) \in W} I_x^2$$

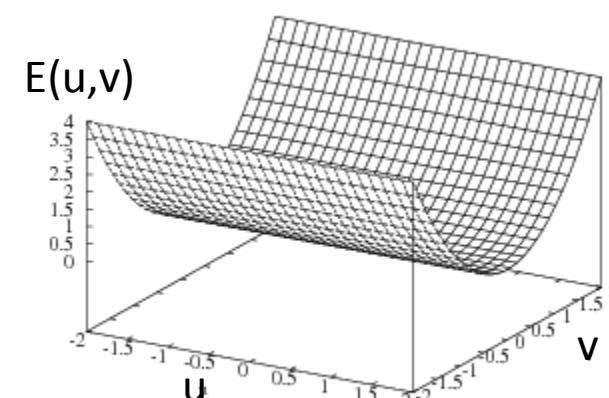
$$B = \sum_{(x,y) \in W} I_x I_y$$

$$C = \sum_{(x,y) \in W} I_y^2$$



Horizontal edge:  $I_x = 0$

$$H = \begin{bmatrix} 0 & 0 \\ 0 & C \end{bmatrix}$$



Dr. Sander Ali Khowaja



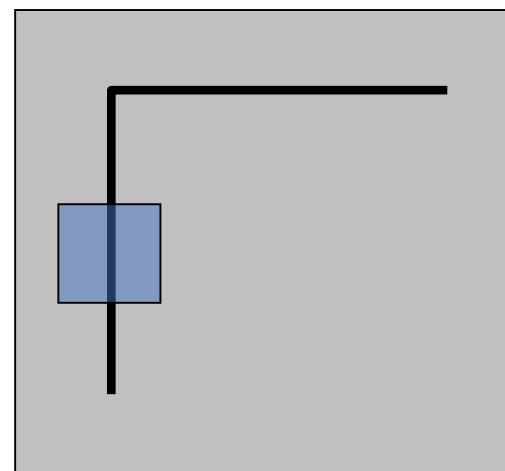
# Corner Detection: The Math

$$E(u, v) \approx \begin{bmatrix} u & v \end{bmatrix} \underbrace{\begin{bmatrix} A & B \\ B & C \end{bmatrix}}_H \begin{bmatrix} u \\ v \end{bmatrix}$$

$$A = \sum_{(x,y) \in W} I_x^2$$

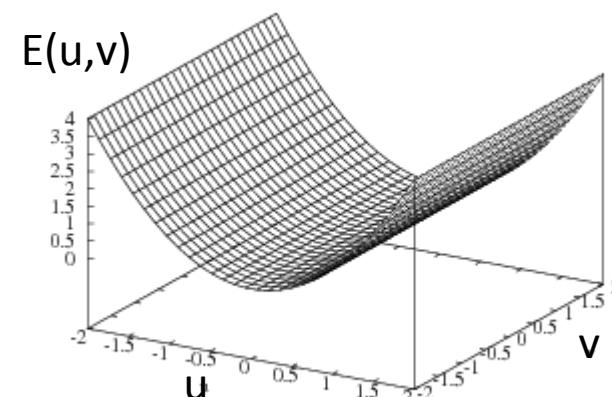
$$B = \sum_{(x,y) \in W} I_x I_y$$

$$C = \sum_{(x,y) \in W} I_y^2$$



Vertical edge:  $I_y = 0$

$$H = \begin{bmatrix} A & 0 \\ 0 & 0 \end{bmatrix}$$



Dr. Sander Ali Khowaja

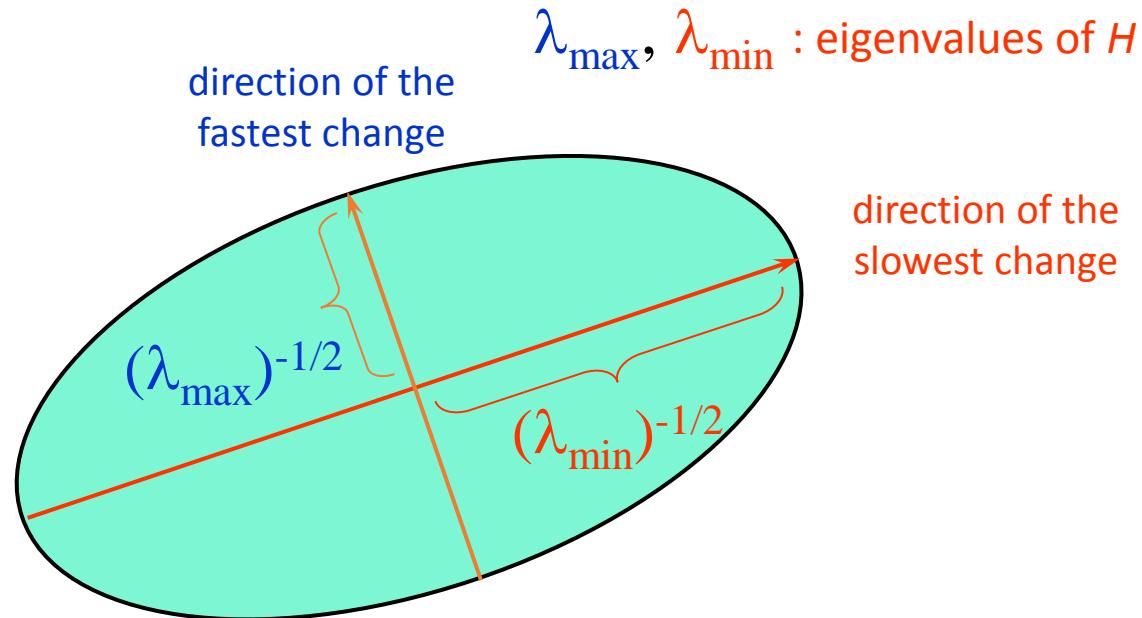
# General Case

The shape of  $H$  tells us something about the *distribution of gradients* around a pixel

We can visualize  $H$  as an ellipse with axis lengths determined by the *eigenvalues* of  $H$  and orientation determined by the *eigenvectors* of  $H$

Ellipse equation:

$$[u \ v] H \begin{bmatrix} u \\ v \end{bmatrix} = \text{const}$$



# Quick Eigenvalue/Eigenvector review

The **eigenvectors** of a matrix  $\mathbf{A}$  are the vectors  $\mathbf{x}$  that satisfy:

$$\mathbf{Ax} = \lambda\mathbf{x}$$

The scalar  $\lambda$  is the **eigenvalue** corresponding to  $\mathbf{x}$

- The eigenvalues are found by solving:

$$\det(\mathbf{A} - \lambda\mathbf{I}) = 0$$

- In our case,  $\mathbf{A} = \mathbf{H}$  is a  $2 \times 2$  matrix, so we have

$$\det \begin{bmatrix} h_{11} - \lambda & h_{12} \\ h_{21} & h_{22} - \lambda \end{bmatrix} = 0$$

- The solution:

$$\lambda_{\pm} = \frac{1}{2} \left[ (h_{11} + h_{22}) \pm \sqrt{4h_{12}h_{21} + (h_{11} - h_{22})^2} \right]$$

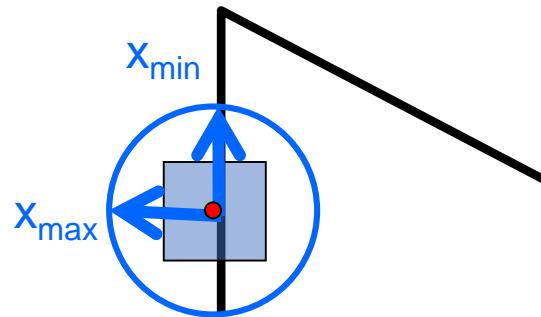
Once you know  $\lambda$ , you find  $\mathbf{x}$  by solving

$$\begin{bmatrix} h_{11} - \lambda & h_{12} \\ h_{21} & h_{22} - \lambda \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} = 0$$



# Corner Detection: The Math

$$E(u, v) \approx \begin{bmatrix} u & v \end{bmatrix} \begin{bmatrix} A & B \\ B & C \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix}$$



$H$

$$Hx_{\max} = \lambda_{\max}x_{\max}$$

$$Hx_{\min} = \lambda_{\min}x_{\min}$$

## Eigenvalues and eigenvectors of $H$

- Define shift directions with the smallest and largest change in error
- $x_{\max}$  = direction of largest increase in  $E$
- $\lambda_{\max}$  = amount of increase in direction  $x_{\max}$
- $x_{\min}$  = direction of smallest increase in  $E$
- $\lambda_{\min}$  = amount of increase in direction  $x_{\min}$

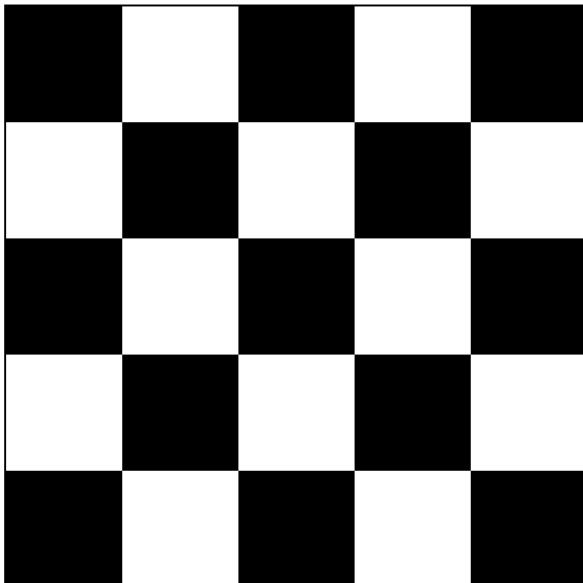
# Corner Detection: The Math

How are  $\lambda_{\max}$ ,  $x_{\max}$ ,  $\lambda_{\min}$ , and  $x_{\min}$  relevant for feature detection?

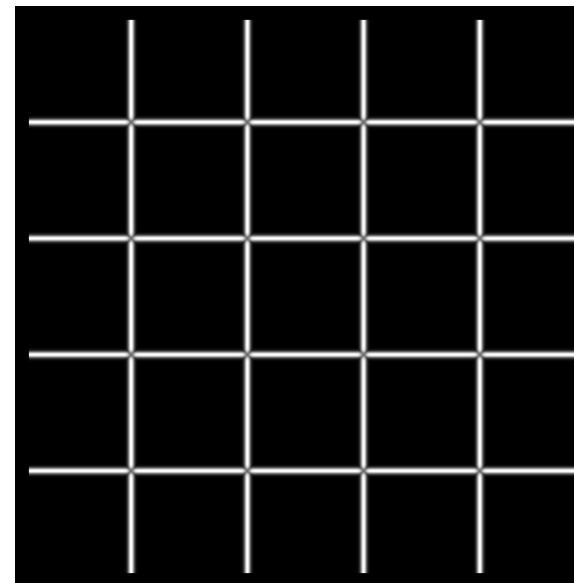
- What's our feature scoring function?

Want  $E(u,v)$  to be large for small shifts in all directions

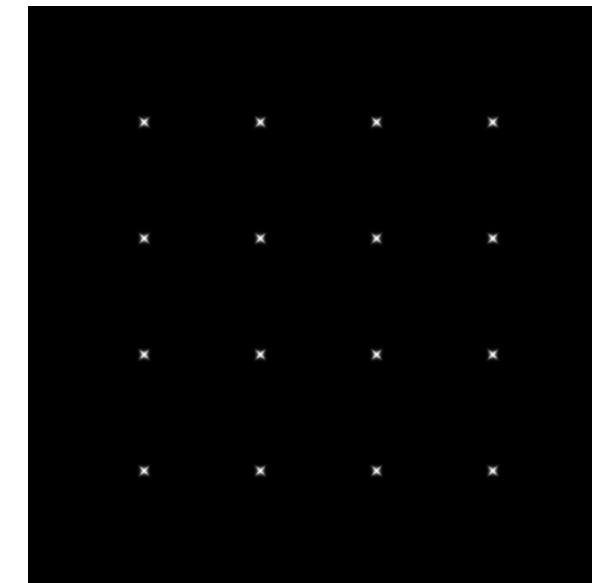
- the minimum of  $E(u,v)$  should be large, over all unit vectors  $[u \ v]$
- this minimum is given by the smaller eigenvalue ( $\lambda_{\min}$ ) of  $H$



$I$



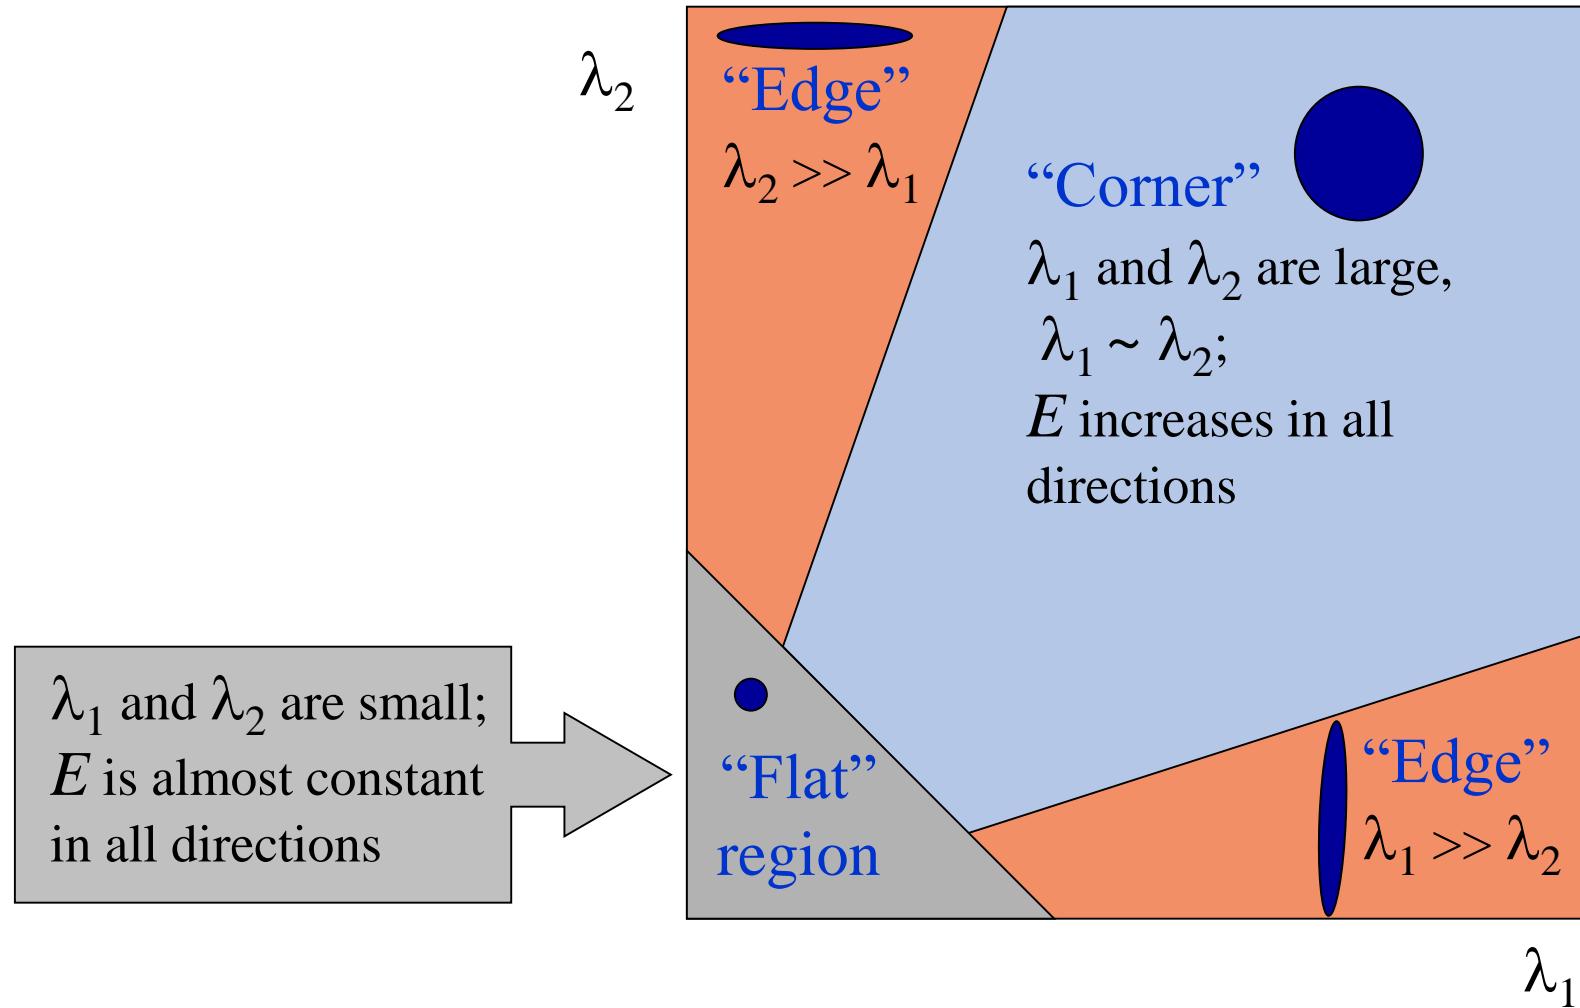
$\lambda_{\max}$



$\lambda_{\min}$

# Interpreting Eigenvalues

Classification of image points using eigenvalues of  $M$ :



# In some books, Corner Detection Math

- We can estimate the stability of the SSD score
- First, take the Taylor series expansion of the image
  - Recall Taylor series expansion of a function  $f(x)$  near  $x_0$

$$f(x) = f(x_0) + \left[ \frac{df}{dx} \right]_{x_0} (x - x_0) + \dots$$

- For vectors  $\mathbf{x}$ ,  $\mathbf{u}$

$$I_0(\mathbf{x}_i + \Delta\mathbf{u}) \approx I_0(\mathbf{x}_i) + \nabla I_0(\mathbf{x}_i) \cdot \Delta\mathbf{u}$$

- Then

$$\begin{aligned} E(\Delta\mathbf{u}) &= \sum_i w(\mathbf{x}_i) [I_0(\mathbf{x}_i + \Delta\mathbf{u}) - I_0(\mathbf{x}_i)]^2 \\ &\approx \sum_i w(\mathbf{x}_i) [I_0(\mathbf{x}_i) + \nabla I_0(\mathbf{x}_i) \cdot \Delta\mathbf{u} - I_0(\mathbf{x}_i)]^2 \\ &= \sum_i w(\mathbf{x}_i) [\nabla I_0(\mathbf{x}_i) \cdot \Delta\mathbf{u}]^2 \end{aligned}$$

where

- $\mathbf{x}_i = (x_i, y_i)^T$  is some image point
- $\Delta\mathbf{u} = (\Delta x, \Delta y)^T$  is the displacement from that point
- $\nabla I_0(\mathbf{x}_i)$  is the gradient of the image at  $\mathbf{x}_i$



# In some books, Corner Detection Math

- Expand the SSD score

$$\begin{aligned} E(\Delta\mathbf{u}) &= \sum_i w(\mathbf{x}_i) [\nabla I_0(\mathbf{x}_i) \cdot \Delta\mathbf{u}]^2 \\ &= \sum_i w(\mathbf{x}_i) \left[ \Delta x \frac{\partial I_0(\mathbf{x}_i)}{\partial x} + \Delta y \frac{\partial I_0(\mathbf{x}_i)}{\partial y} \right]^2 \\ &= \sum_i w(\mathbf{x}_i) \left[ \Delta x \left( \frac{\partial I_0(\mathbf{x}_i)}{\partial x} \right)^2 \Delta x + 2 \Delta x \left( \frac{\partial I_0(\mathbf{x}_i)}{\partial x} \right) \left( \frac{\partial I_0(\mathbf{x}_i)}{\partial y} \right) \Delta y + \Delta y \left( \frac{\partial I_0(\mathbf{x}_i)}{\partial y} \right)^2 \Delta y \right] \\ &= \Delta x \left[ \sum_i w(\mathbf{x}_i) \left( \frac{\partial I_0(\mathbf{x}_i)}{\partial x} \right)^2 \right] \Delta x + 2 \Delta x \left[ \sum_i w(\mathbf{x}_i) \left( \frac{\partial I_0(\mathbf{x}_i)}{\partial x} \right) \left( \frac{\partial I_0(\mathbf{x}_i)}{\partial y} \right) \right] \Delta y + \Delta y \left[ \sum_i w(\mathbf{x}_i) \left( \frac{\partial I_0(\mathbf{x}_i)}{\partial y} \right)^2 \right] \Delta y \\ &= \Delta x (w * I_{xx}) \Delta x + 2 \Delta x (w * I_{xy}) \Delta y + \Delta y (w * I_{yy}) \Delta y \\ &= (\Delta x \quad \Delta y) \begin{pmatrix} w * I_{xx} & w * I_{xy} \\ w * I_{xy} & w * I_{yy} \end{pmatrix} \begin{pmatrix} \Delta x \\ \Delta y \end{pmatrix} \end{aligned}$$

- where

$$I_x^2 = \left( \frac{\partial I}{\partial x} \right)^2, I_{xy} = \left( \frac{\partial I}{\partial x} \right) \left( \frac{\partial I}{\partial y} \right), I_y^2 = \left( \frac{\partial I}{\partial y} \right)^2$$

# In some books, Corner Detection Math

- The SSD score is

$$E(\Delta\mathbf{u}) = (\Delta x \quad \Delta y) \begin{pmatrix} w * I_{xx} & w * I_{xy} \\ w * I_{xy} & w * I_{yy} \end{pmatrix} \begin{pmatrix} \Delta x \\ \Delta y \end{pmatrix}$$
$$= \Delta\mathbf{u}^T \mathbf{A} \Delta\mathbf{u}$$

$$I_x^2 = \left( \frac{\partial I}{\partial x} \right)^2, I_{xy} = \left( \frac{\partial I}{\partial x} \right) \left( \frac{\partial I}{\partial y} \right),$$
$$I_y^2 = \left( \frac{\partial I}{\partial y} \right)^2$$

- where

$$\mathbf{A} = w * \begin{pmatrix} I_x^2 & I_{xy} \\ I_{xy} & I_y^2 \end{pmatrix}$$

Note – the “\*” indicates correlation of  $w$  with each element of  $\mathbf{A}$

- $w$  is a small  $N \times N$  averaging filter such as a box filter
- it essentially averages the values over a small local neighborhood

- $\mathbf{A}$  can tell us how stable the SSD score is, at a given point



# Possible Interest Point Measures

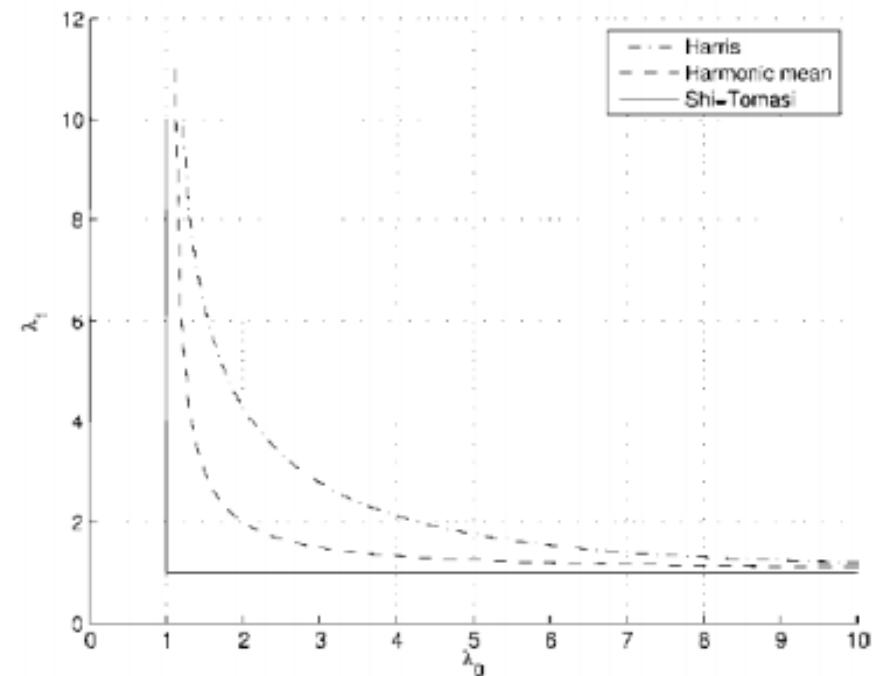
- Look at the smallest eigenvalue – if it is sufficiently large, then we have a candidate point (**Shi-Tomasi**)
- Alternative measures that don't require square roots:

**(Harris)**

$$\det(\mathbf{A}) - \alpha \operatorname{tr}(\mathbf{A})^2 = \lambda_0 \lambda_1 - \alpha(\lambda_0 + \lambda_1)^2$$

**(Harmonic mean)**

$$\frac{\det(\mathbf{A})}{\operatorname{tr}(\mathbf{A})} = \frac{\lambda_0 \lambda_1}{\lambda_0 + \lambda_1}$$



From *Computer Vision:  
Algorithms and Applications*,  
Richard Szeliski, 2010, Springer

# MATLAB Code

```
% Find corner points
clear all
close all

I = double(imread('test000.jpg'));

% Apply Gaussian blur
sd = 1.0;
I = imfilter(I, fspecial('gaussian', round(6*sd), sd));

imshow(I, []);

% Compute the gradient components
Gx = imfilter(I, [-1 1]);
Gy = imfilter(I, [-1; 1]);

% Compute the outer product g'g at each pixel
Gxx = Gx .* Gx;
Gxy = Gx .* Gy;
Gyy = Gy .* Gy;

% Sum the G's over the window size.
A11 = imfilter(Gxx, w);
A12 = imfilter(Gxy, w);
A22 = imfilter(Gyy, w);

% At each pixel (x,y), we have the 2x2 matrix
% [A11(x,y) A12(x,y);
%  A21(x,y) A22(x,y)]
% Of course, A21 = A12.

% Compute the interest score: det(A)/trace(A)
detA = A11.*A22 - A12.^2; % Computes det(A) at each pixel
traceA = A11 + A22; % Computer trace(A) at each pixel
s = detA./traceA; % The interest score

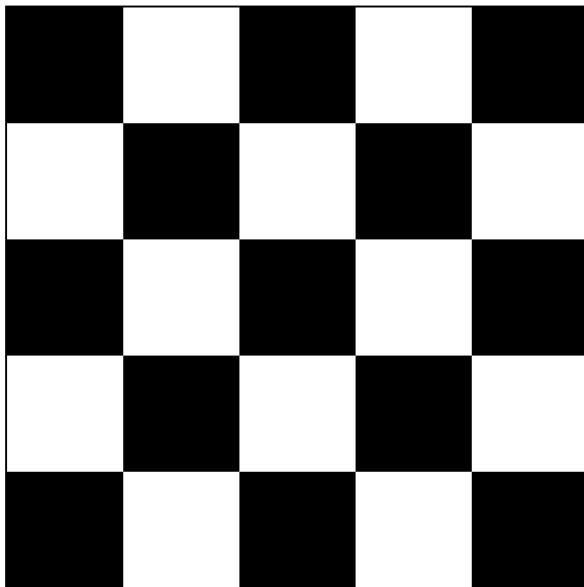
% If trace(A)=0 anywhere, we get "not a number".
s(isnan(s)) = 0; % If any NaN's, just set them to zero
```



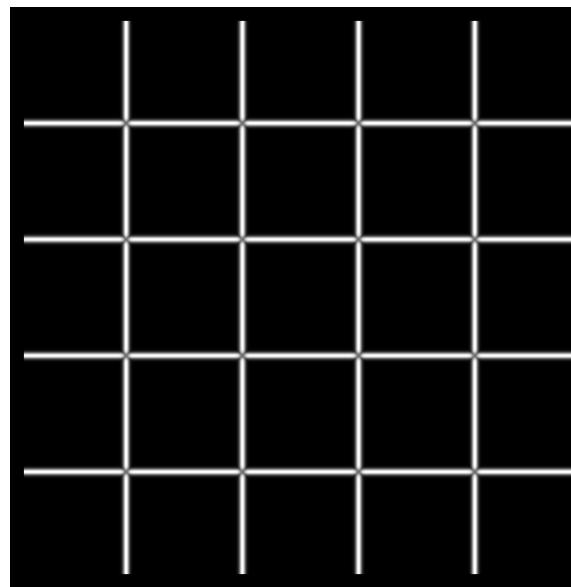
# Corner Detection Summary

Here's what you do

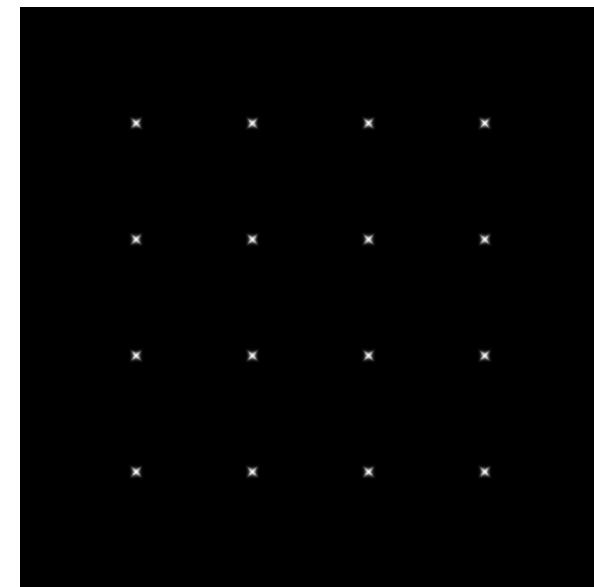
- Compute the gradient at each point in the image
- Create the  $H$  matrix from the entries in the gradient
- Compute the eigenvalues.
- Find points with large response ( $\lambda_{\min} > \text{threshold}$ )
- Choose those points where  $\lambda_{\min}$  is a local maximum as features



$I$



$\lambda_{\max}$

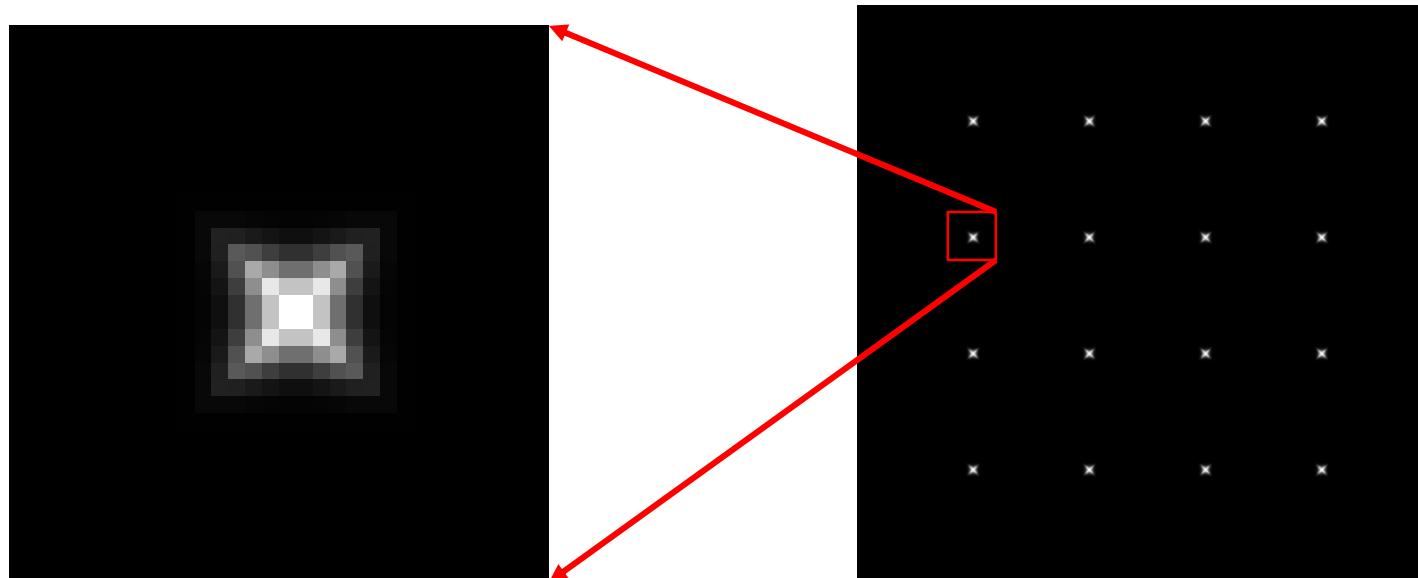


$\lambda_{\min}$

# Corner Detection Summary

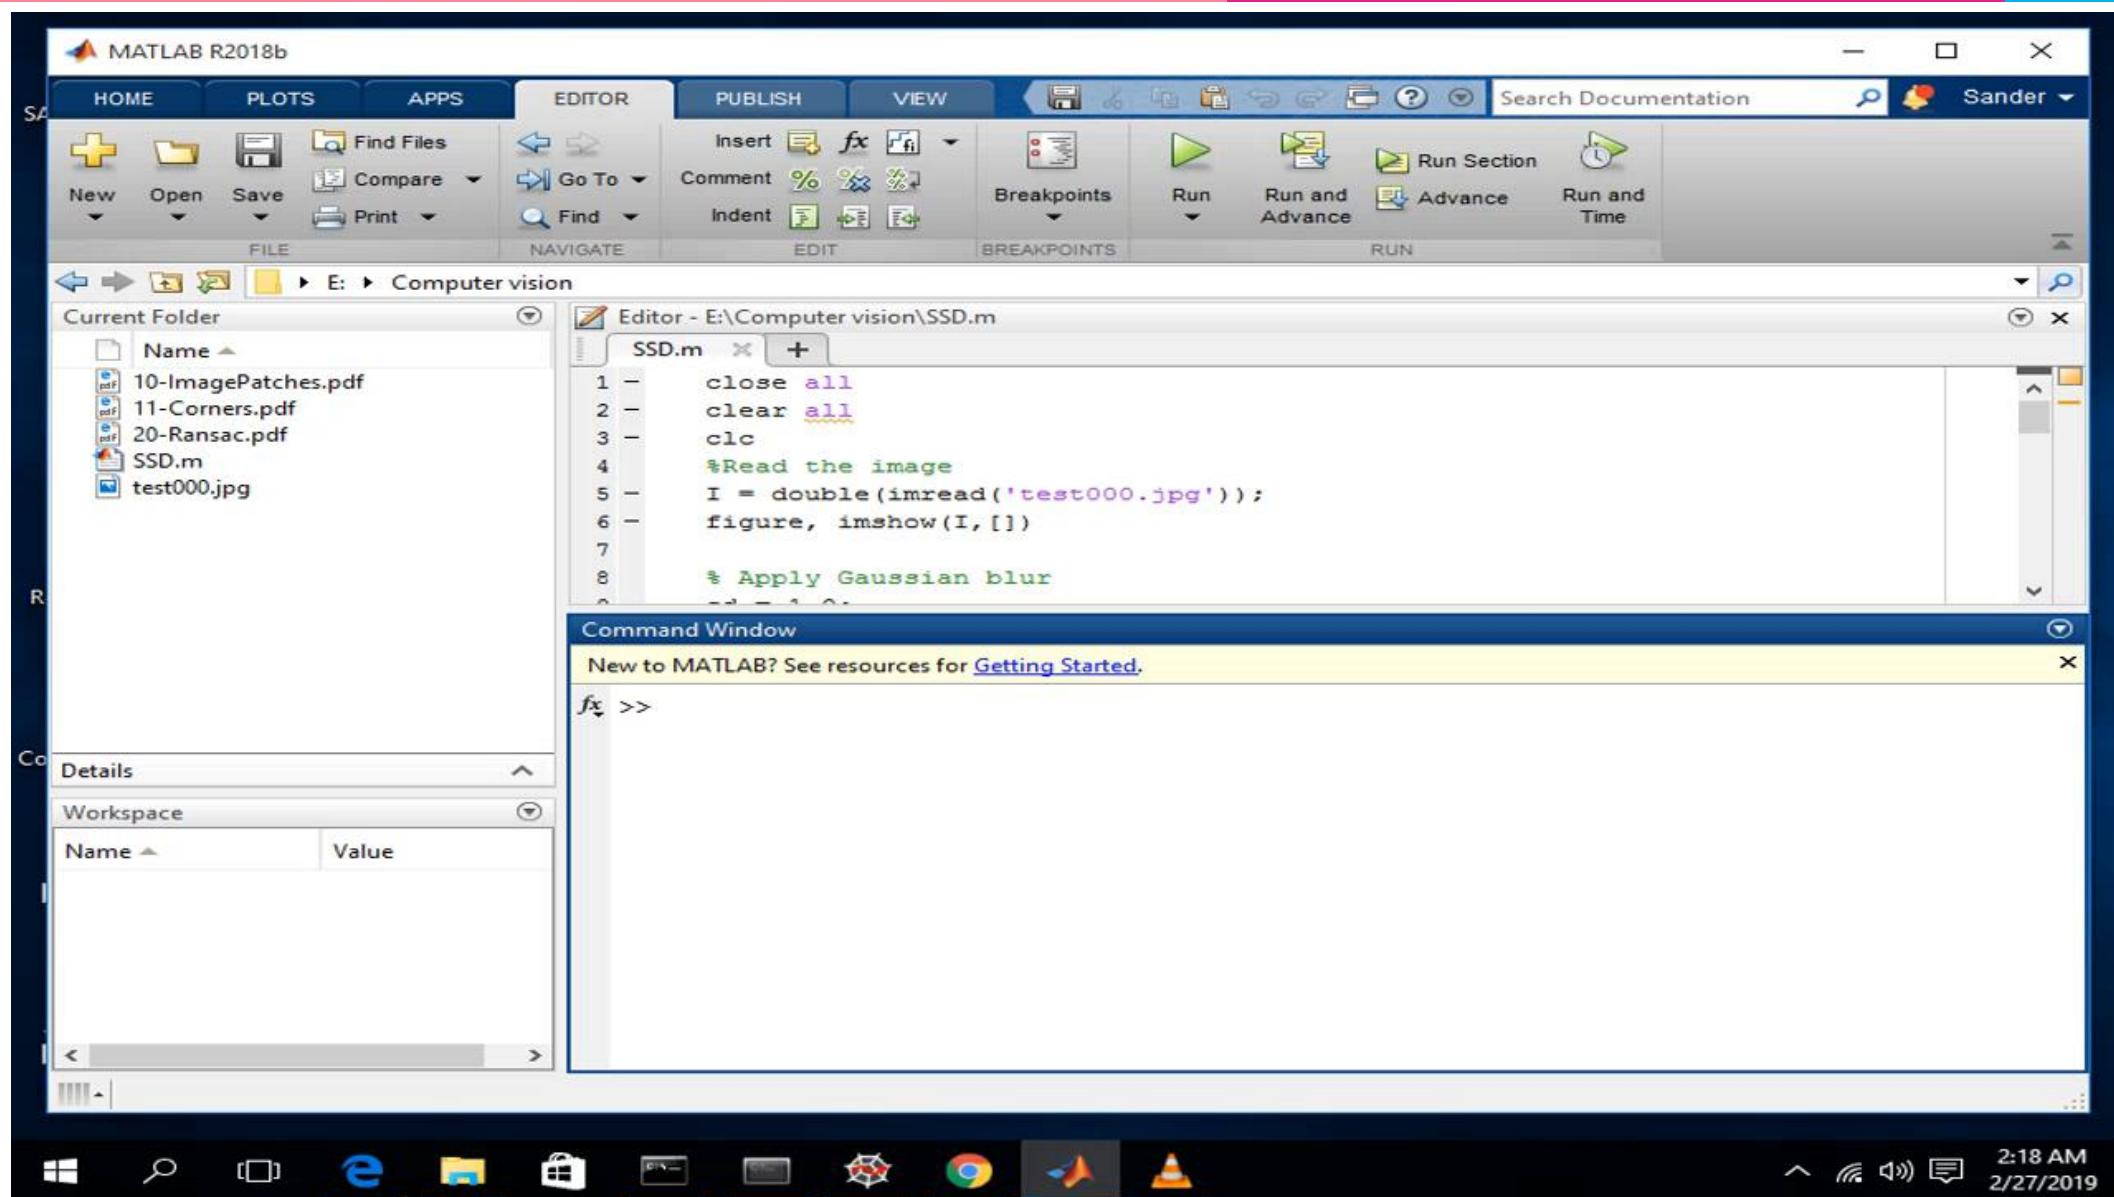
Here's what you do

- Compute the gradient at each point in the image
- Create the  $H$  matrix from the entries in the gradient
- Compute the eigenvalues.
- Find points with large response ( $\lambda_{\min} > \text{threshold}$ )
- Choose those points where  $\lambda_{\min}$  is a local maximum as features



$$\lambda_{\min}$$

# Demo



# The Harris Operator

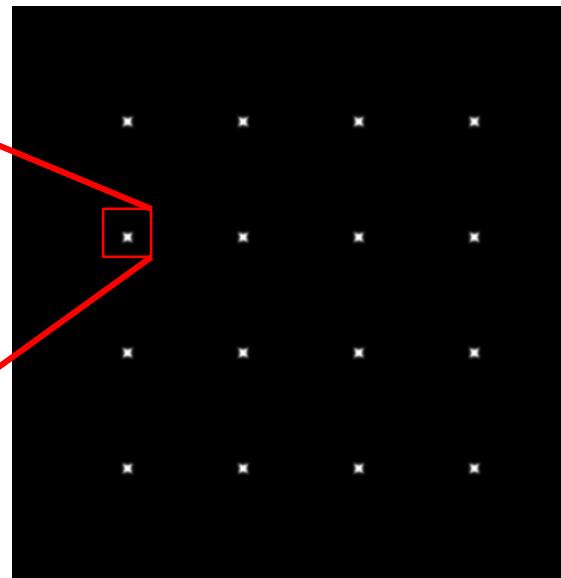
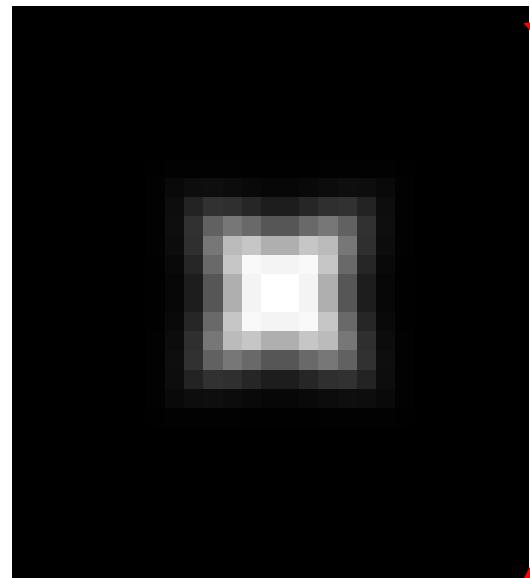
$\lambda_{\min}$  is a variant of the “Harris operator” for feature detection

$$f = \frac{\lambda_1 \lambda_2}{\lambda_1 + \lambda_2}$$
$$= \frac{\text{determinant}(H)}{\text{trace}(H)}$$

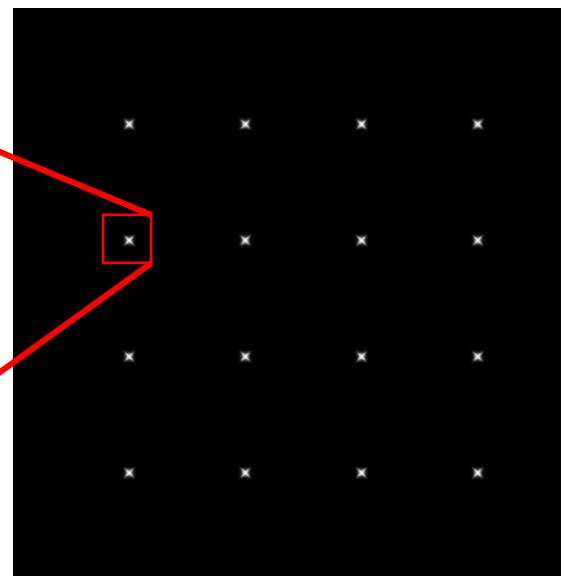
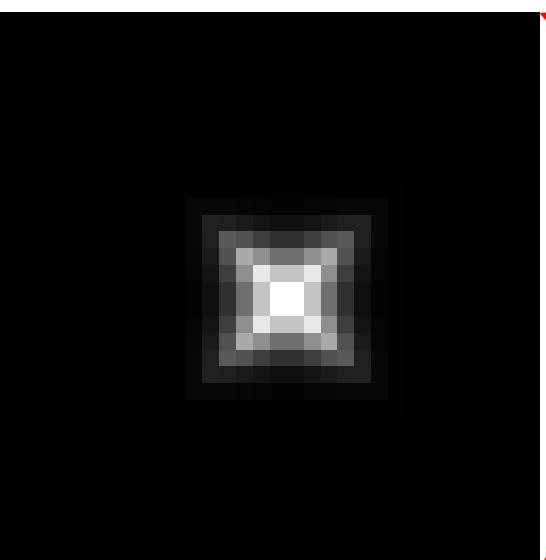
- The *trace* is the sum of the diagonals, i.e.,  $\text{trace}(H) = h_{11} + h_{22}$
- Very similar to  $\lambda_{\min}$  but less expensive (no square root)
- Called the “Harris Corner Detector” or “Harris Operator”
- Lots of other detectors, this is one of the most popular



# The Harris Operator

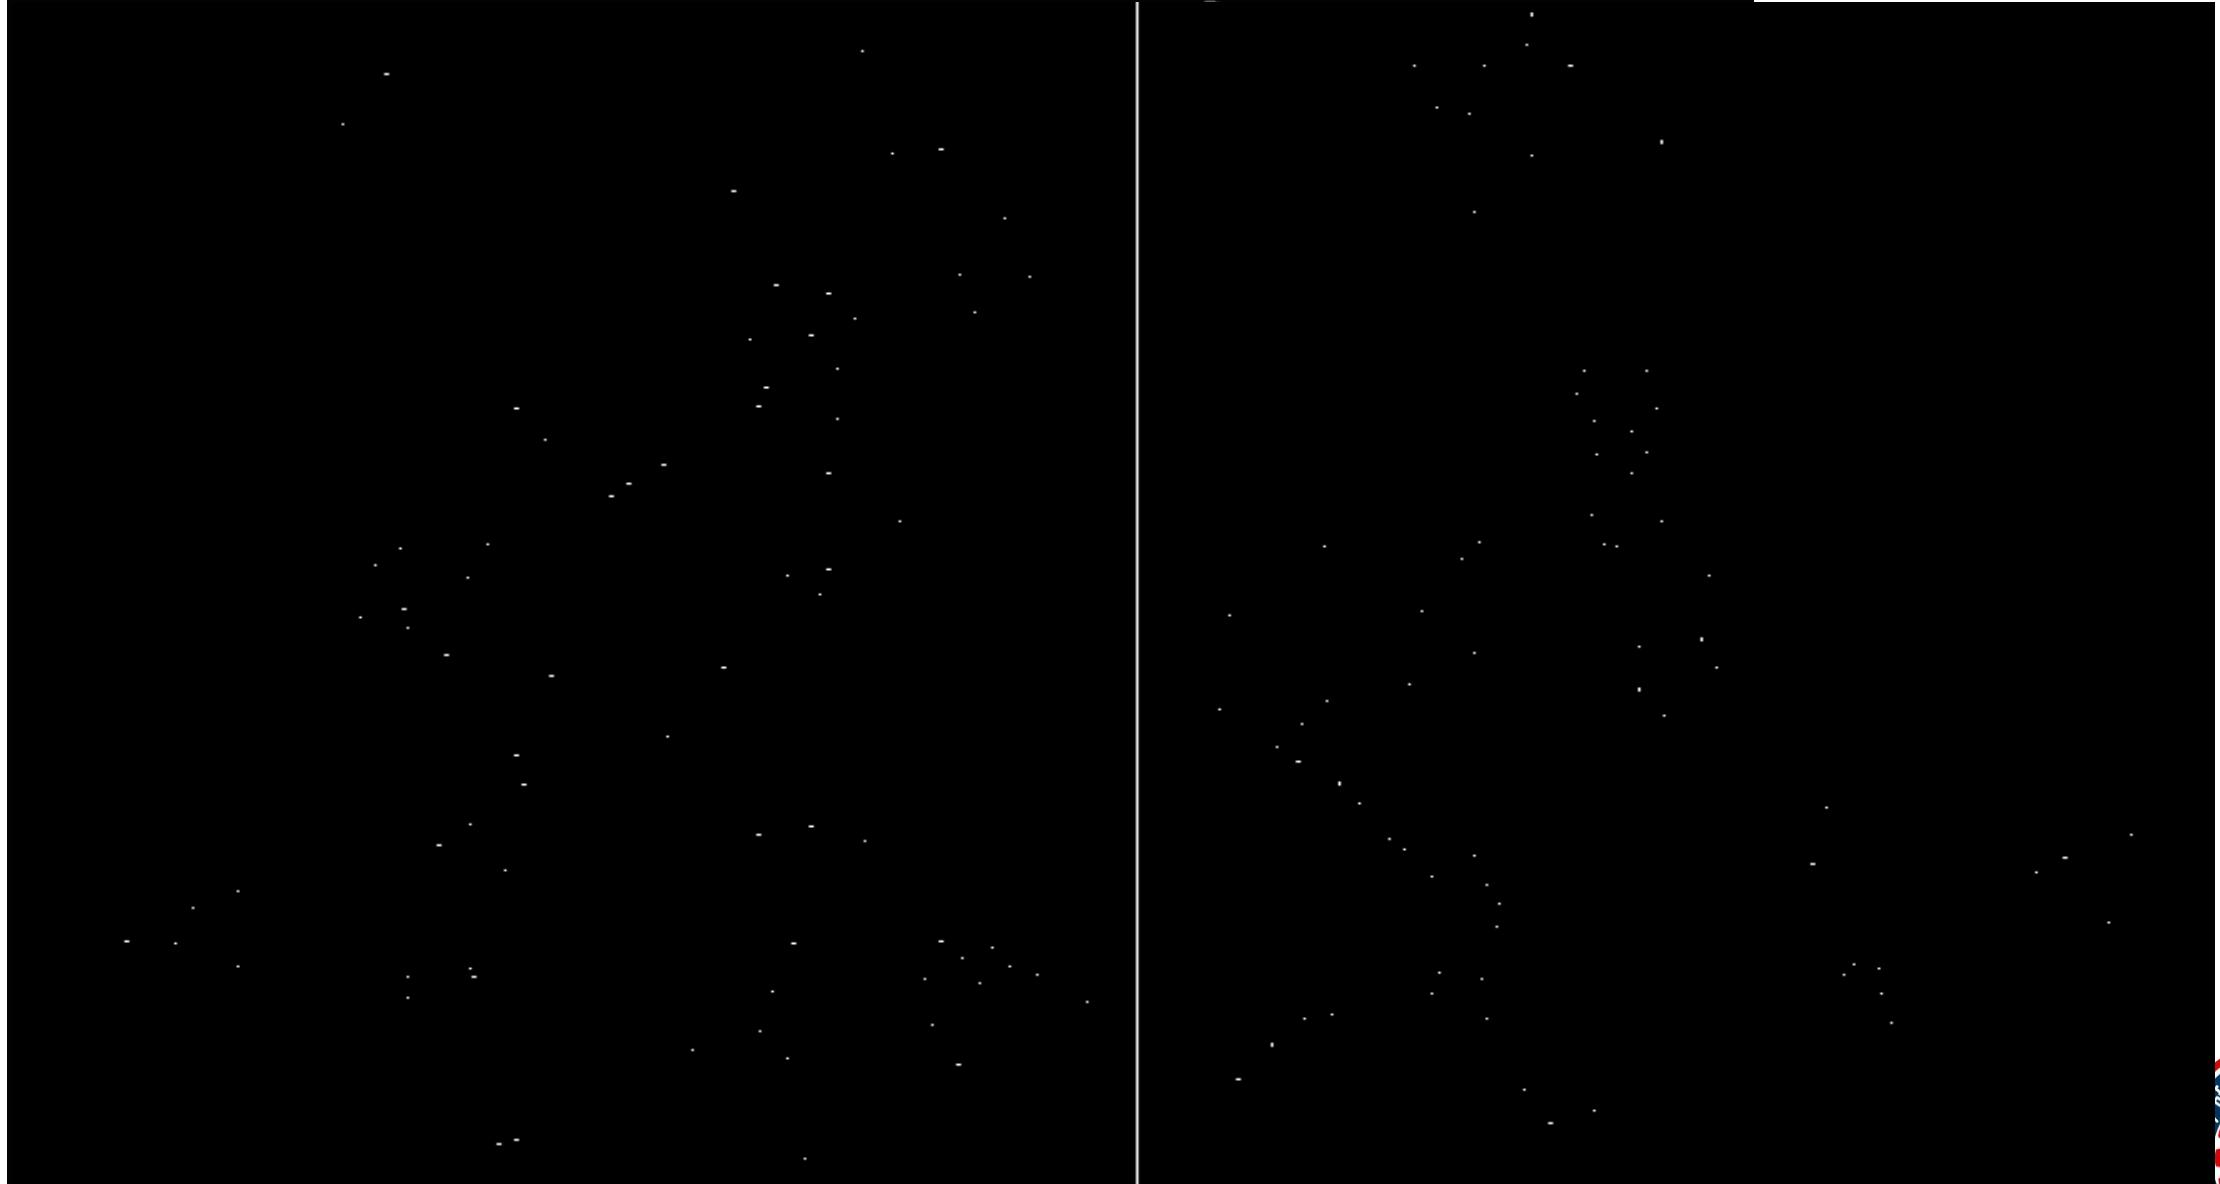


Harris operator



$\lambda_{\min}$

# Harris Detector Example



# Harris Features in Red



# Code for Harris Corner Detection

```
1 # organizing imports
2 import cv2
3 import numpy as np
4
5 # path to input image specified and
6 # image is loaded with imread command
7 image = cv2.imread('8.jpg')
8
9 # convert the input image into
10 # grayscale color space
11 operatedImage = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)
12
13 # modify the data type
14 # setting to 32-bit floating point
15 operatedImage = np.float32(operatedImage)
16
17 # apply the cv2.cornerHarris method
18 # to detect the corners with appropriate
19 # values as input parameters
20 dest = cv2.cornerHarris(operatedImage, 2, 5, 0.07)
21
22 # Results are marked through the dilated corners
23 dest = cv2.dilate(dest, None)
24
25 # Reverting back to the original image,
26 # with optimal threshold value
27 image[dest > 0.01 * dest.max()]=[0, 0, 255]
28
29 # the window showing output image with corners
30 cv2.imshow('Image with Borders', image)
31
32 # De-allocate any associated memory usage
33 if cv2.waitKey(0) & 0xff == 27:
34     cv2.destroyAllWindows()
35
```



# Results for Harris Corner Detection



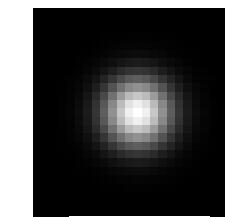
# Weighting the derivatives

- In practice, using a simple window  $W$  doesn't work too well

$$H = \sum_{(x,y) \in W} \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

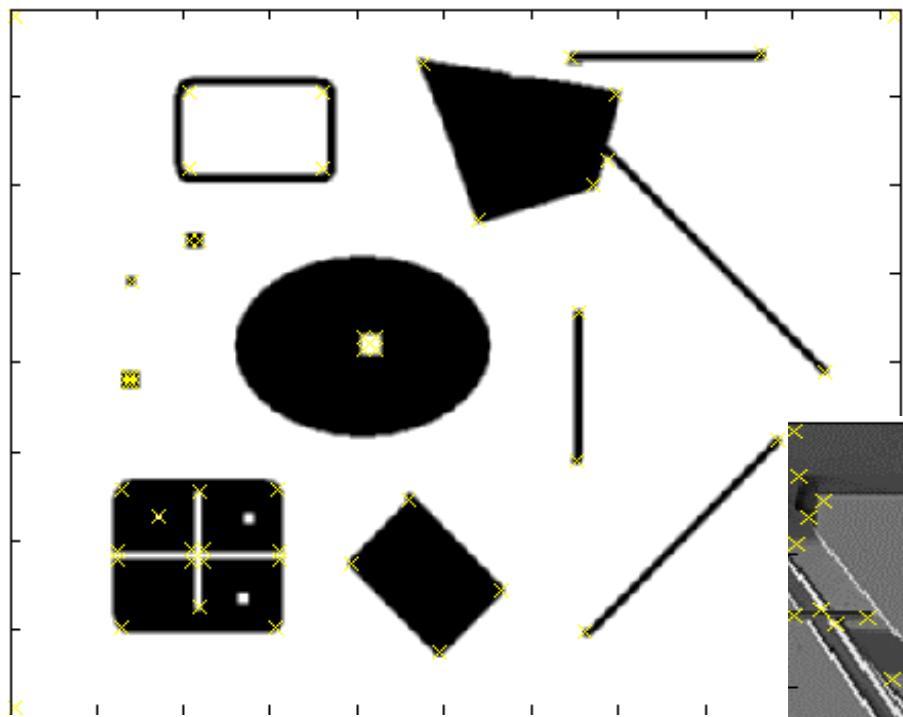
- Instead, we'll *weight* each derivative value based on its distance from the center pixel

$$H = \sum_{(x,y) \in W} w_{x,y} \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

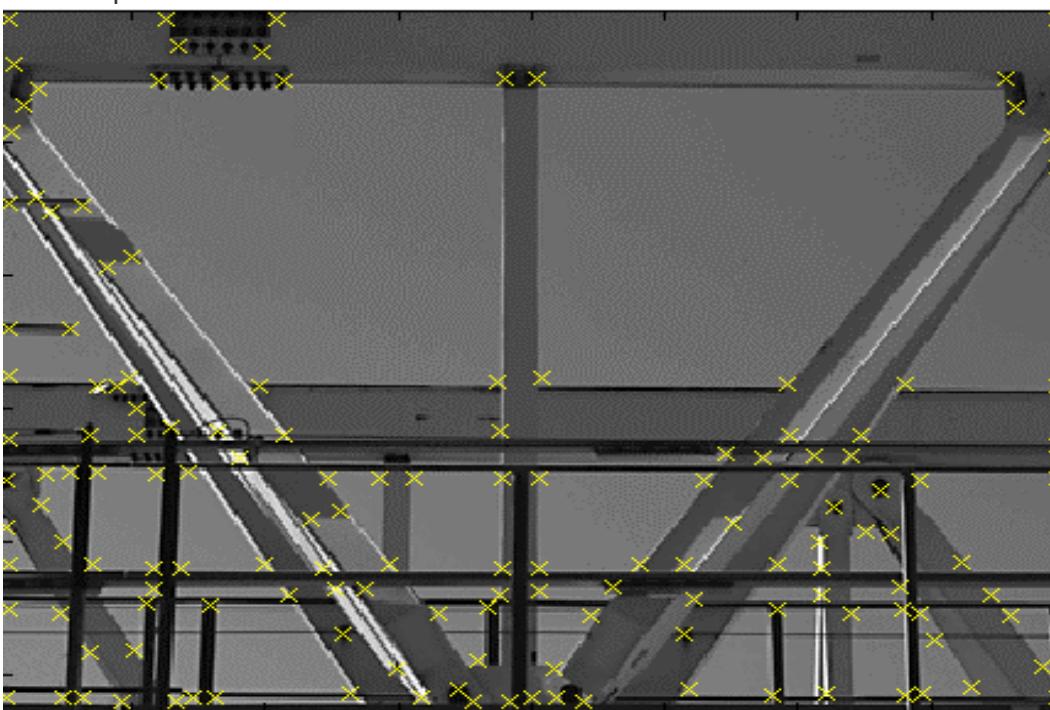


$w_{x,y}$

# Harris Detector – Responses [Harris88]



**Effect:** A very precise corner detector.

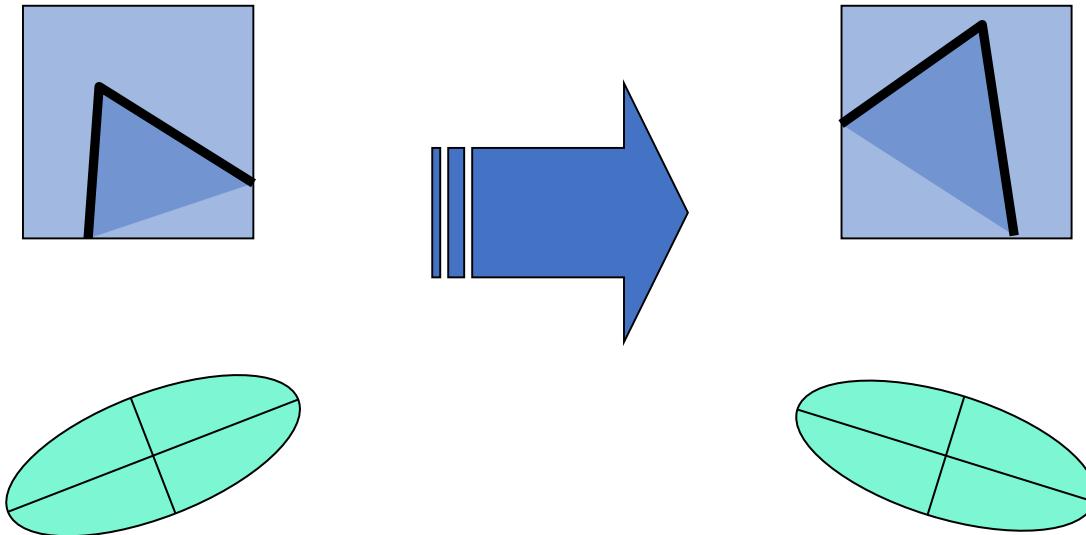


# Harris Detector – Responses [Harris88]



# Harris Detector: Invariance Properties

- Rotation



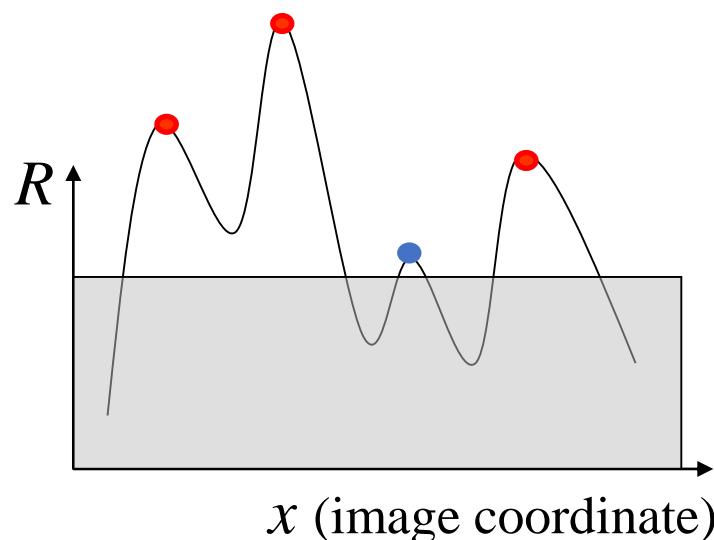
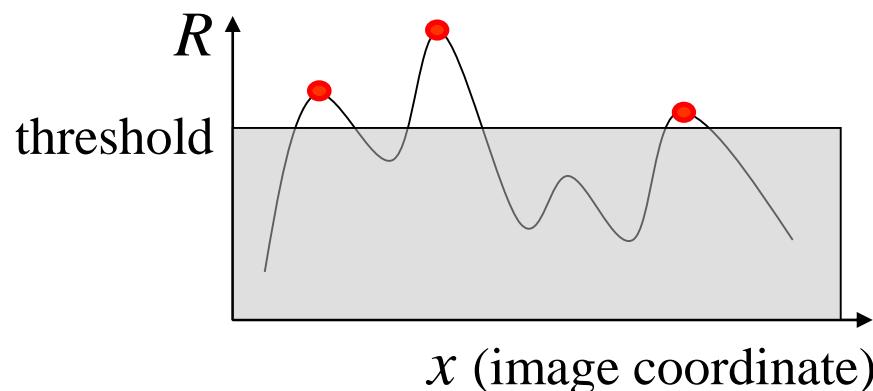
Ellipse rotates but its shape (i.e. eigenvalues)  
remains the same

Corner response is invariant to image rotation

# Harris Detector: Invariance Properties

- Affine intensity change:  $I \rightarrow aI + b$ 
  - ✓ Only derivatives are used => invariance to intensity shift  $I \rightarrow I + b$

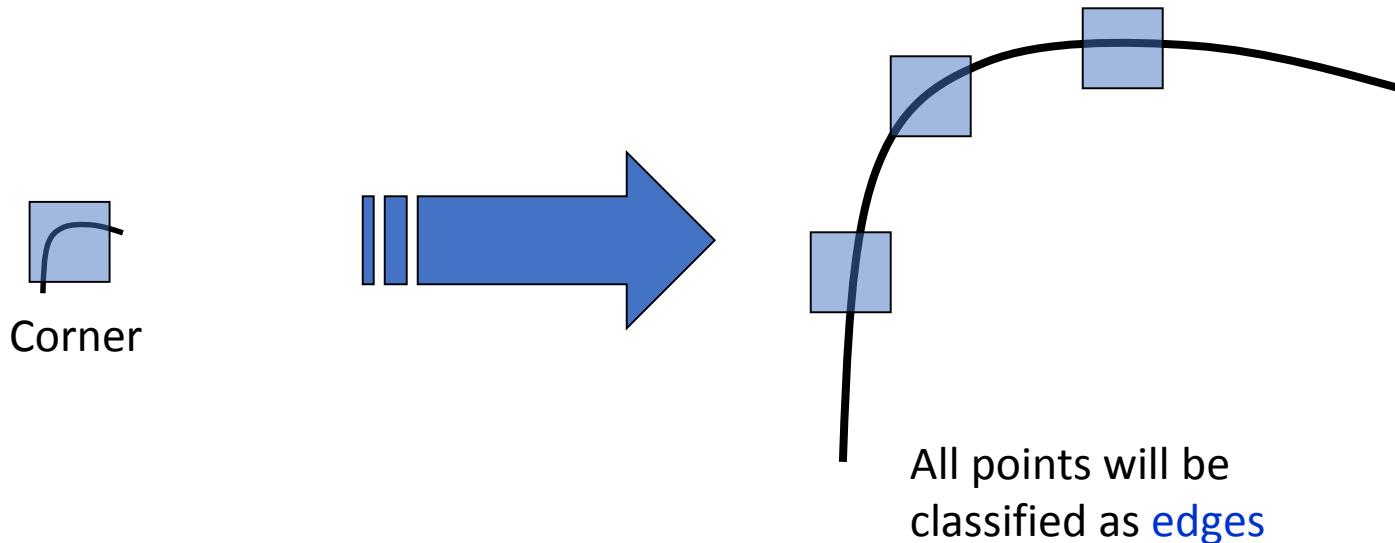
- ✓ Intensity scale:  $I \rightarrow a I$



*Partially invariant to affine intensity change*

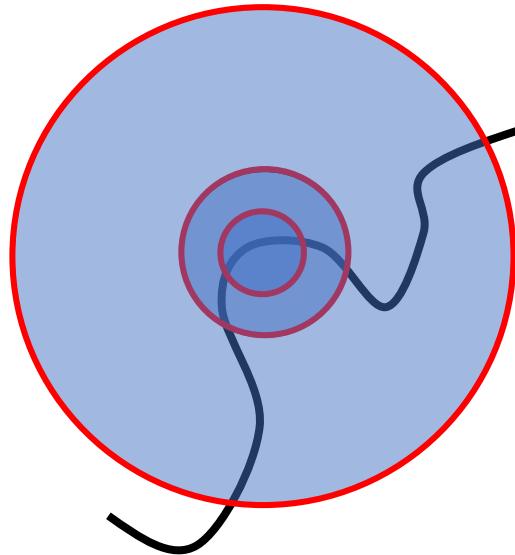
# Harris Detector: Invariance Properties

- Scaling



*Not invariant to scaling*

Suppose you're looking for corners



Key idea: find scale that gives local maximum of  $f$

- in both position and scale
- One definition of  $f$ : the Harris operator

# Automatic Scale Selection



$$f(I_{i_1 \dots i_m}(x, \sigma)) = f(I_{i_1 \dots i_m}(x', \sigma'))$$

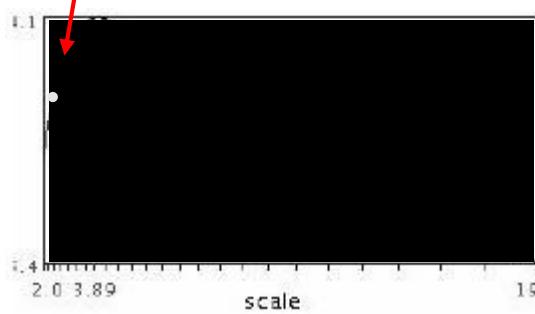
How to find corresponding patch sizes?

K. Grauman, B. Leibe

Dr. Sander Ali Khowaja

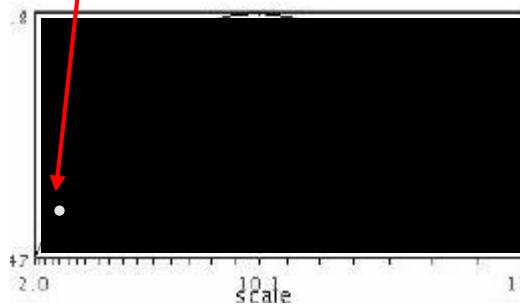
# Automatic Scale Selection

- Function responses for increasing scale (scale signature)



$$f(I_{i_1...i_m}(x, \sigma))$$

K. Grauman, B. Leibe



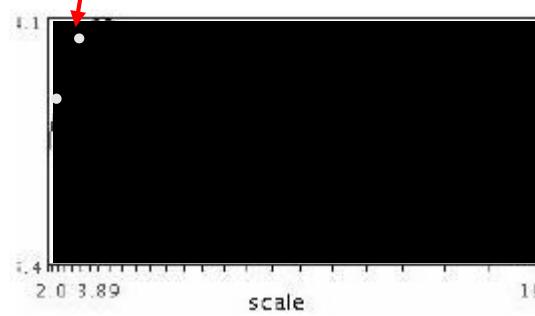
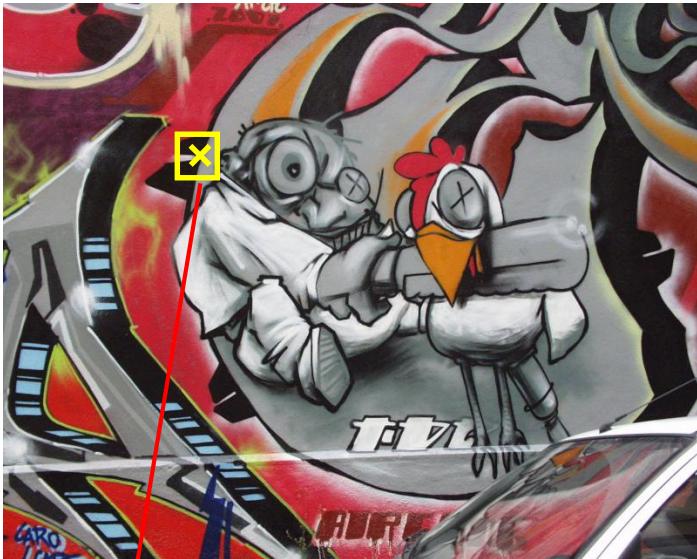
$$f(I_{i_1...i_m}(x', \sigma))$$

Dr. Sander Ali Khowaja



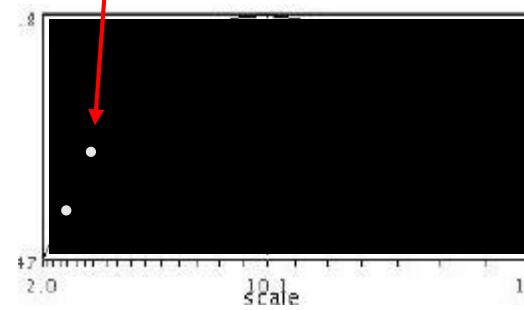
# Automatic Scale Selection

- Function responses for increasing scale (scale signature)



$$f(I_{i_1...i_m}(x, \sigma))$$

K. Grauman, B. Leibe

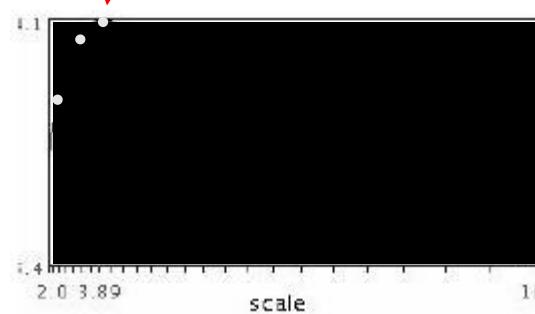
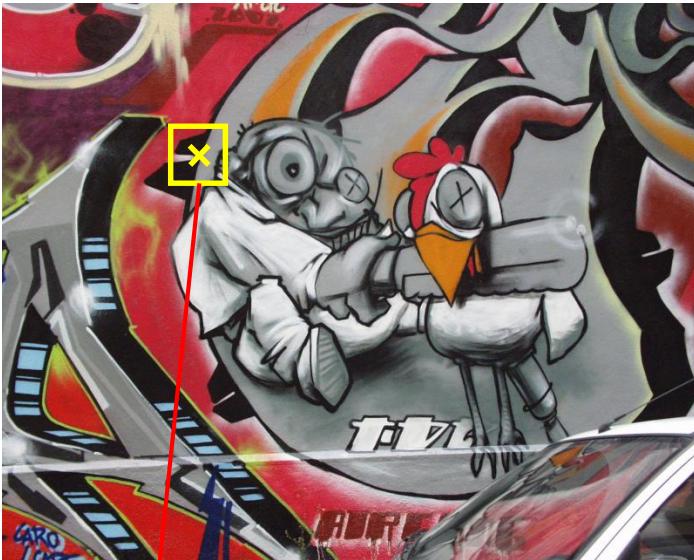


$$f(I_{i_1...i_m}(x', \sigma))$$

Dr. Sander Ali Khowaja

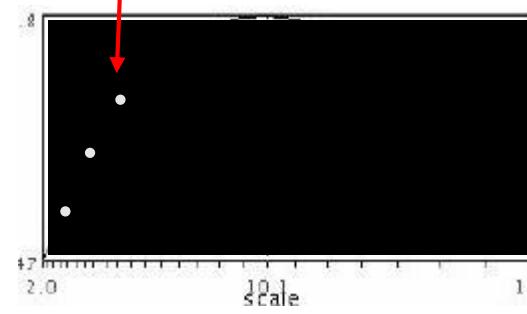
# Automatic Scale Selection

- Function responses for increasing scale (scale signature)



$$f(I_{i_1 \dots i_m}(x, \sigma))$$

K. Grauman, B. Leibe

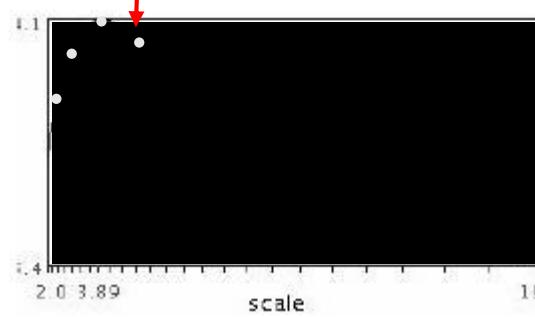
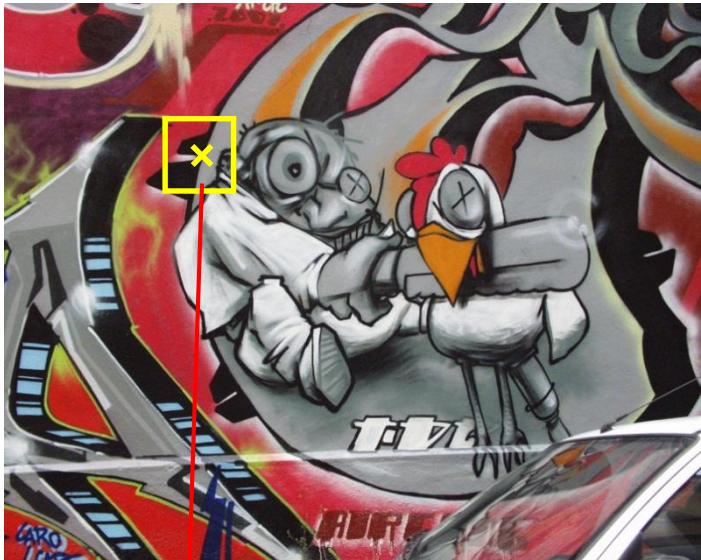


$$f(I_{i_1 \dots i_m}(x', \sigma))$$

Dr. Sander Ali Khowaja

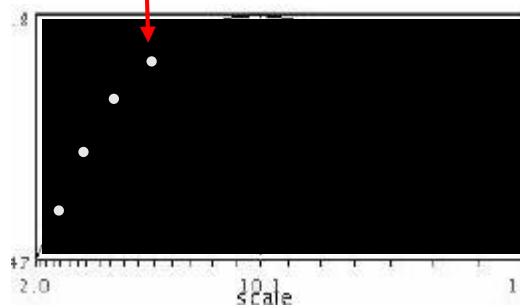
# Automatic Scale Selection

- Function responses for increasing scale (scale signature)



$$f(I_{i_1 \dots i_m}(x, \sigma))$$

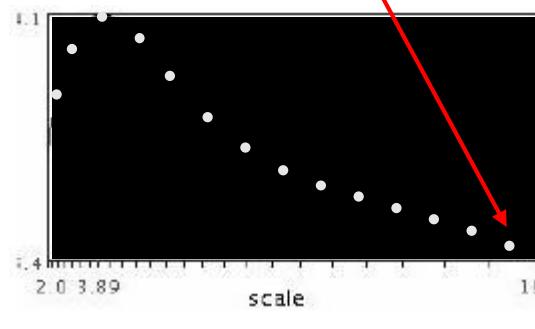
K. Grauman, B. Leibe



$$f(I_{i_1 \dots i_m}(x', \sigma))$$

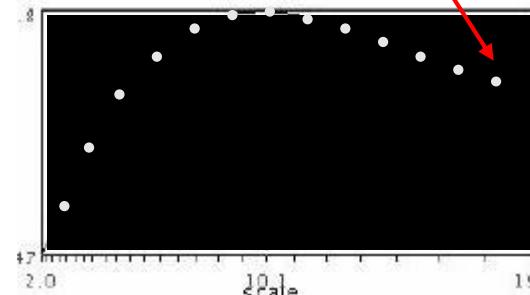
# Automatic Scale Selection

- Function responses for increasing scale (scale signature)



$$f(I_{i_1 \dots i_m}(x, \sigma))$$

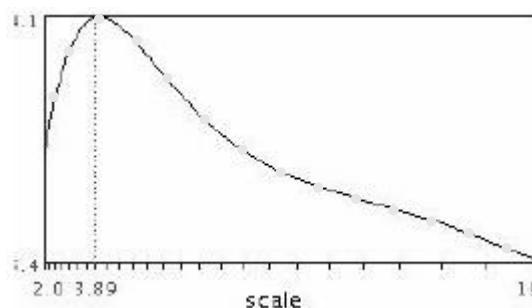
K. Grauman, B. Leibe



$$f(I_{i_1 \dots i_m}(x', \sigma))$$

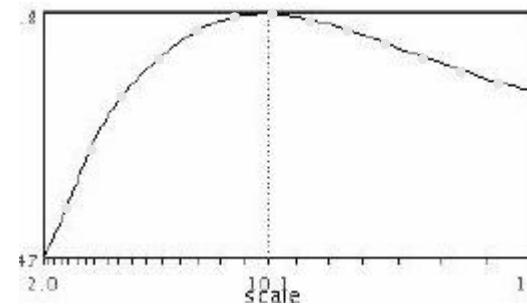
# Automatic Scale Selection

- Function responses for increasing scale (scale signature)



$$f(I_{i_1 \dots i_m}(x, \sigma))$$

K. Grauman, B. Leibe

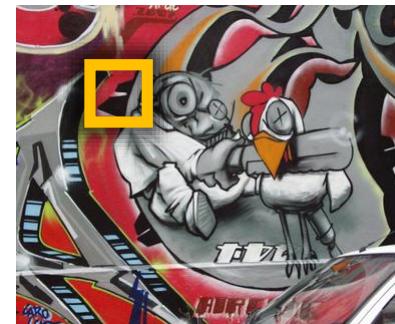


$$f(I_{i_1 \dots i_m}(x', \sigma'))$$

Dr. Sander Ali Khowaja

# Implementation

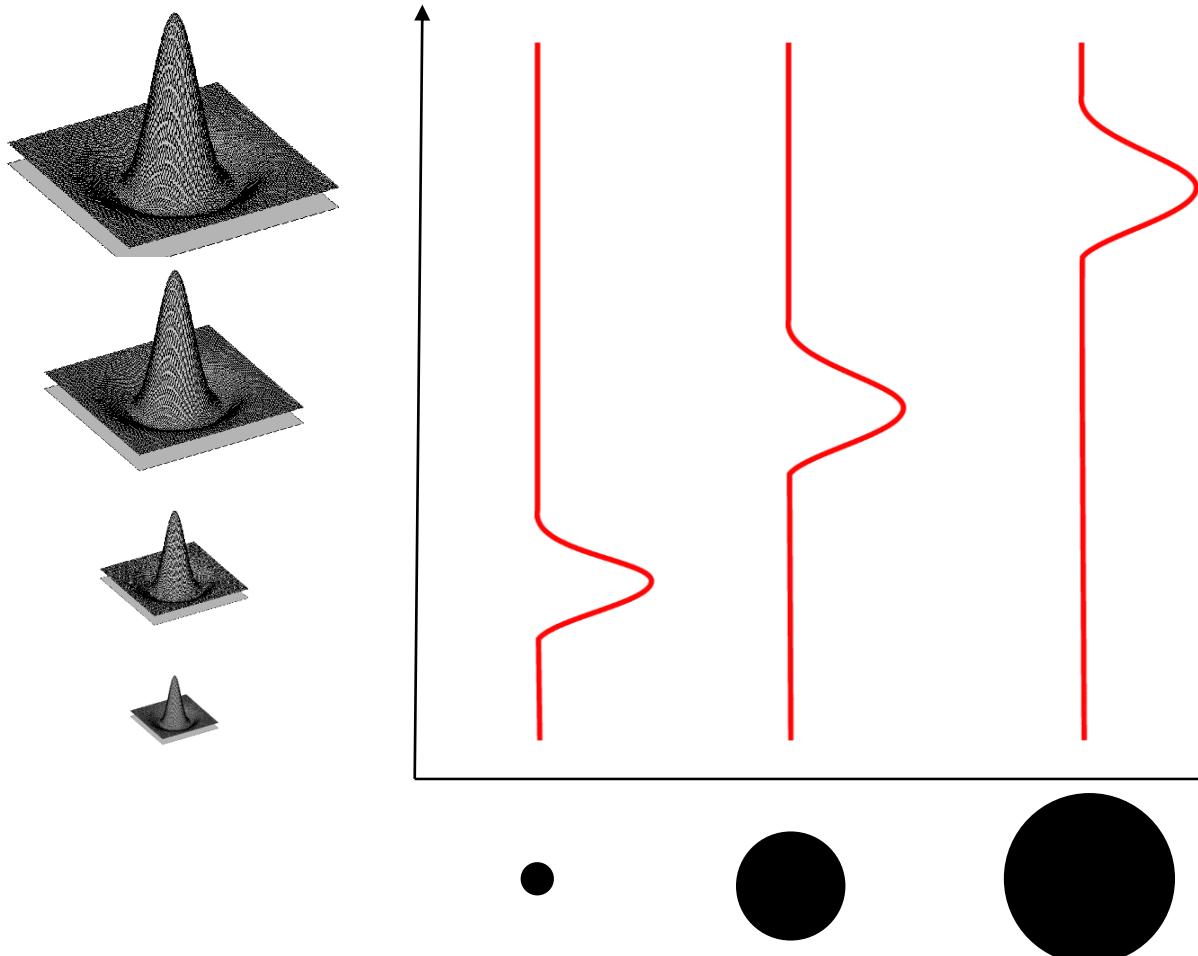
- Instead of computing  $f$  for larger and larger windows, we can implement using a fixed window size with a Gaussian pyramid



(sometimes need to create in-between levels, e.g. a  $\frac{3}{4}$ -size image)

# What is a useful signature function

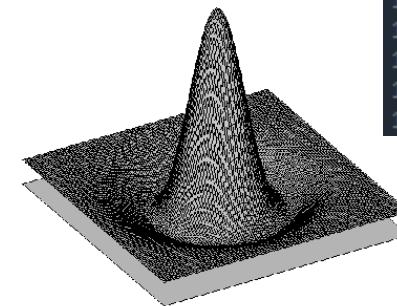
- Difference-of-Gaussian = “blob” detector



K. Grauman, B. Leibe

Dr. Sander Ali Khowaja

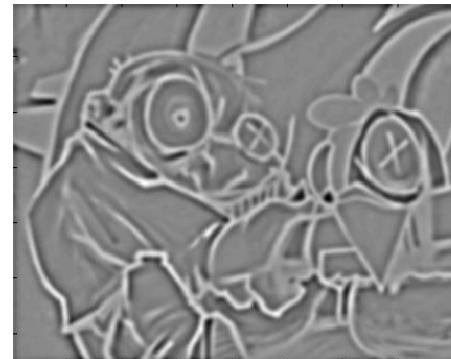
# Difference of Gaussian



```
1 import cv2
2
3 img = cv2.imread('54.jpg')
4
5 # Apply 3x3 and 7x7 Gaussian blur
6 low_sigma = cv2.GaussianBlur(img,(3,3),0)
7 high_sigma = cv2.GaussianBlur(img,(5,5),0)
8
9 # Calculate the DoG by subtracting
10 dog = low_sigma - high_sigma
11
12 # the window showing output image with corners
13 cv2.imshow('Difference of Gaussian', dog)
14
15 # De-allocate any associated memory usage
16 if cv2.waitKey(0) & 0xff == 27:
17     cv2.destroyAllWindows()
```



=

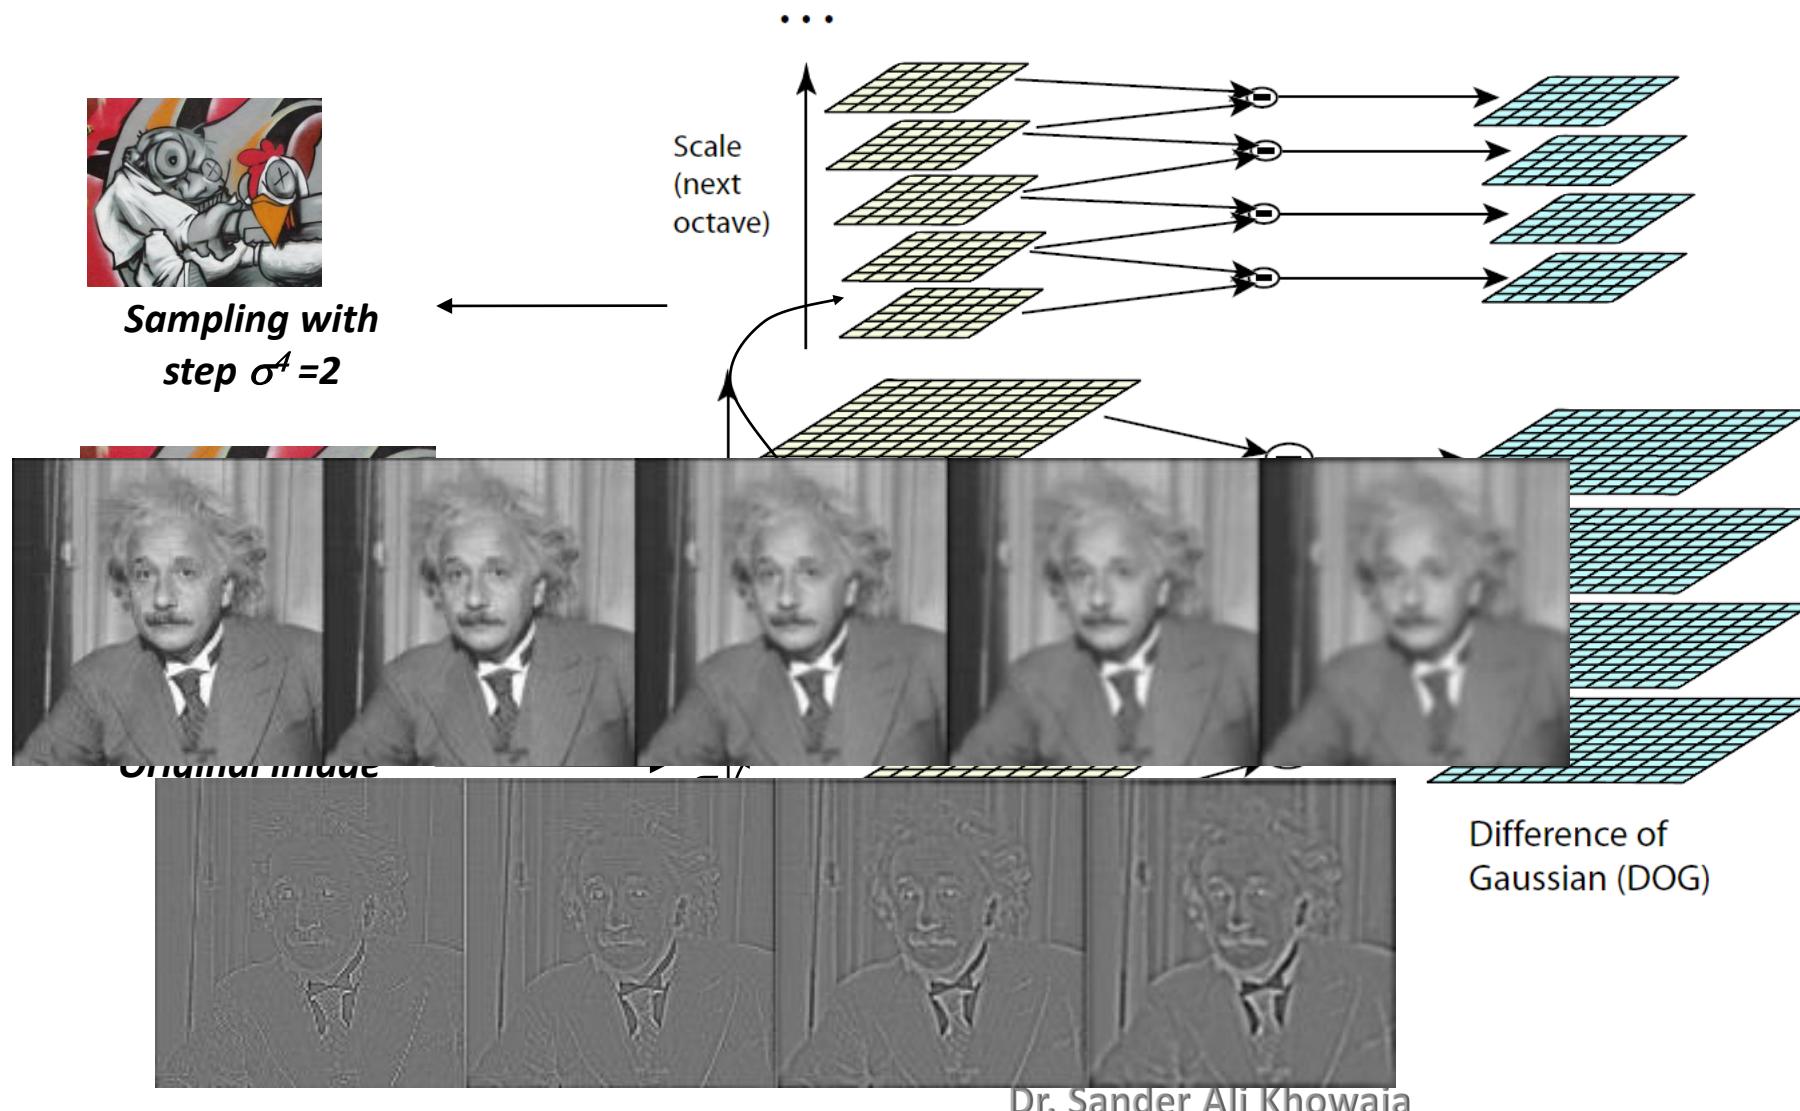


K. Grauman, B. Leibe

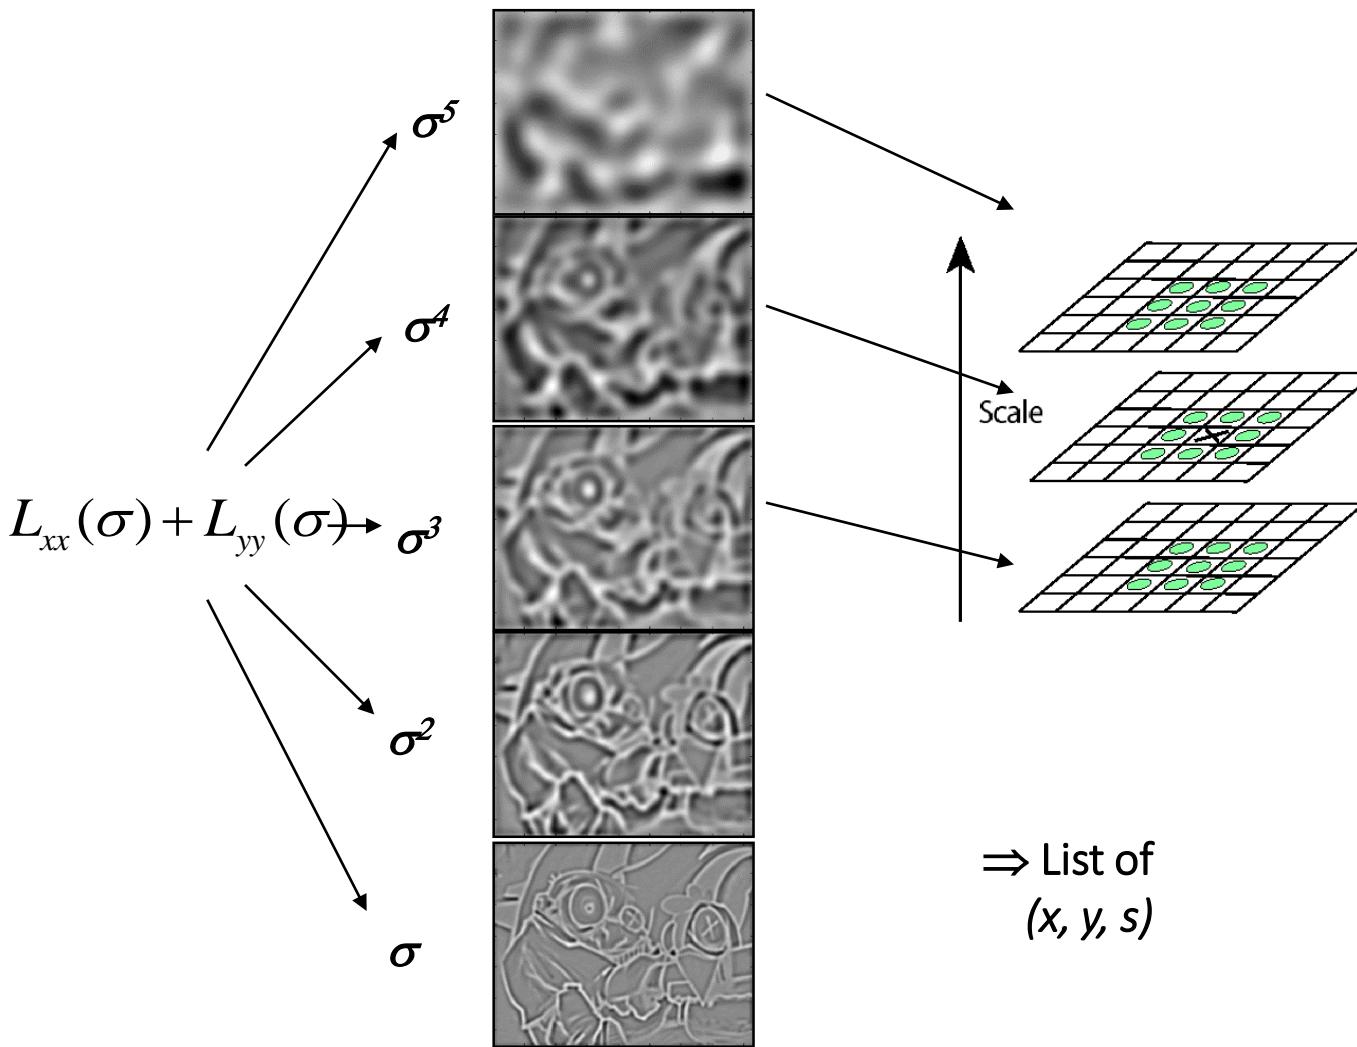
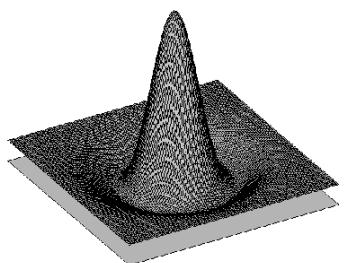
Dr. Sander Ali Khowaja

# Difference of Gaussian – Efficient Computation

- Computation in Gaussian scale pyramid



# Find Local Maxima in Position-scale space of Difference-of-Gaussian



K. Grauman, B. Leibe

Dr. Sander Ali Khowaja

# Example of Gaussian Kernel



Original video



Blurred with a  
gaussian kernel



Blurred with a different  
gaussian kernel

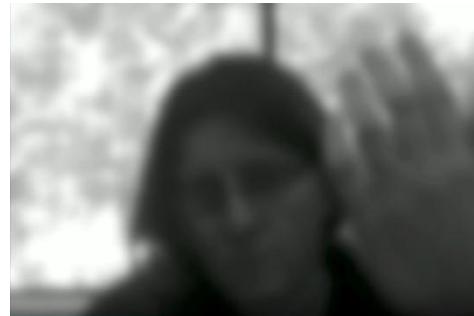
What happens if you subtract one blurred image from another?

[source](#)

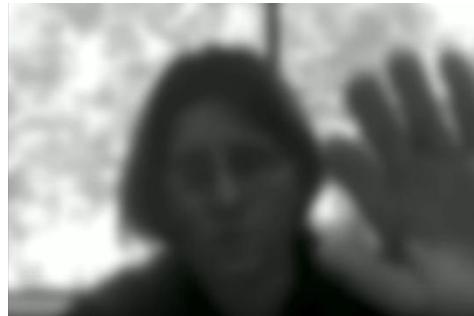
# Difference of Gaussians



Original video



Blurred with a  
gaussian kernel:  $k_1$



Blurred with a different  
gaussian kernel:  $k_2$



DoG:  $k_1 - k_2$



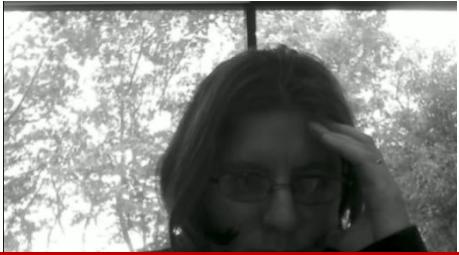
DoG:  $k_1 - k_3$



DoG:  $k_1 - k_4$

Source: [https://www.youtube.com/watch?v=oTud1De\\_W4s](https://www.youtube.com/watch?v=oTud1De_W4s)

# Difference of Gaussians



At different resolutions of kernel size, we see different fine details of the image. In other words, we can capture keypoints at varying scales.



DoG:  $k_1 - k_2$

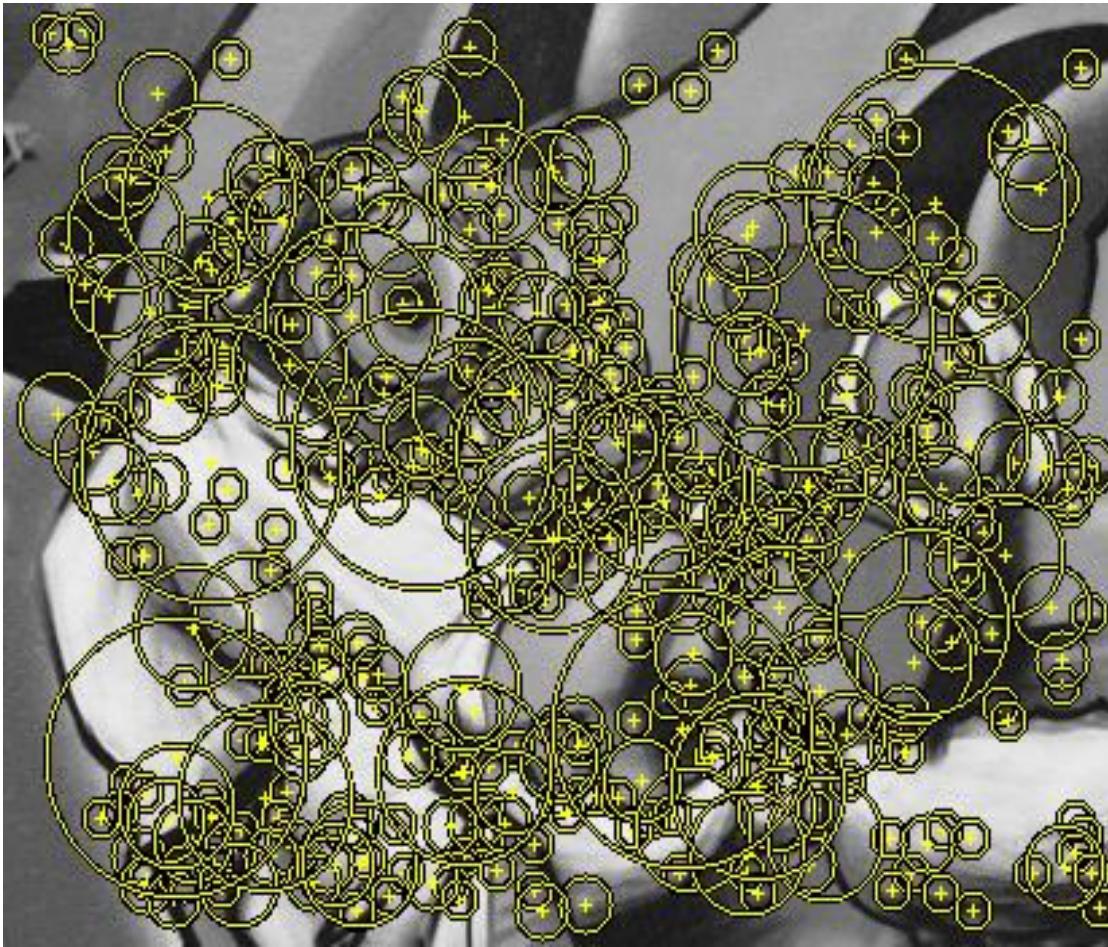


DoG:  $k_1 - k_3$



DoG:  $k_1 - k_4$

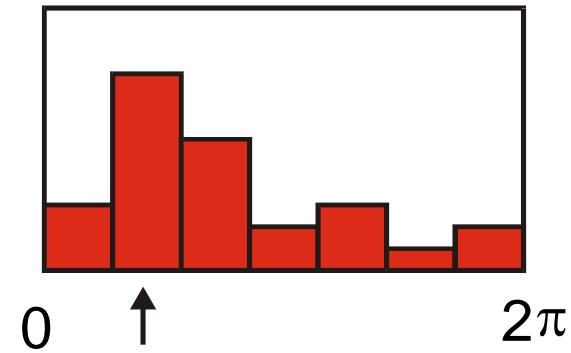
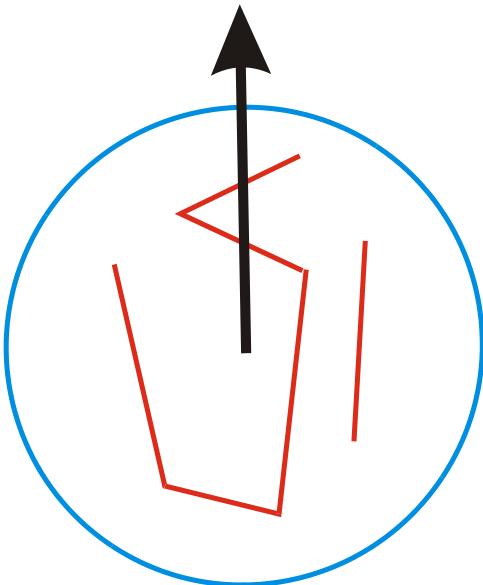
# Results of Interest points using Difference of Gaussian



# Orientation Normalization

[Lowe, SIFT, 1999]

- Compute orientation histogram
  - Select dominant orientation
  - Normalize: rotate to fixed orientation



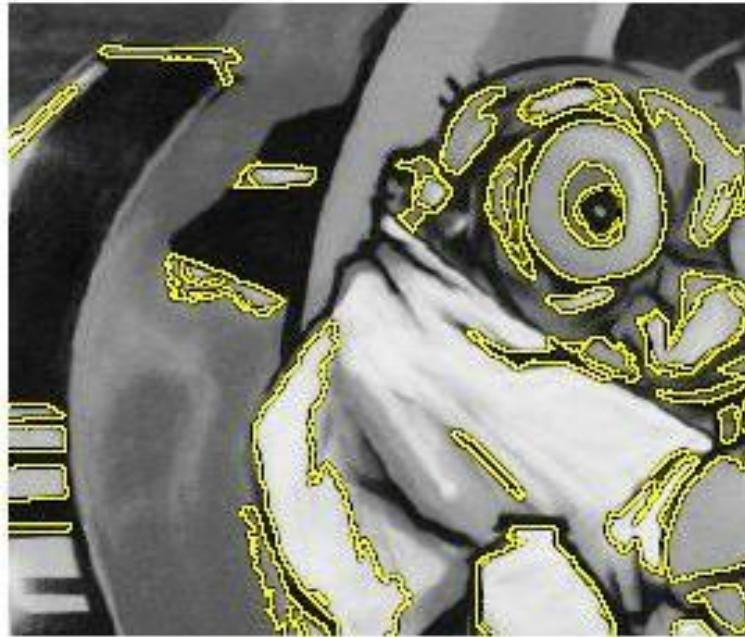
# Maximally Stable Extremal Regions [Matas '02]

- Based on Watershed segmentation algorithm
- Select regions that stay stable over a large parameter range



K. Grauman, B. Leibe

# Example Results: MSER



K. Grauman, B. Leibe

Dr. Sander Ali Khowaja

# Feature Detector codes

- For most local feature detectors, executables are available online:
  - <http://www.robots.ox.ac.uk/~vgg/research/affine>
  - <http://www.cs.ubc.ca/~lowe/keypoints/>
  - <http://www.vision.ee.ethz.ch/~surf>



# Local Descriptors

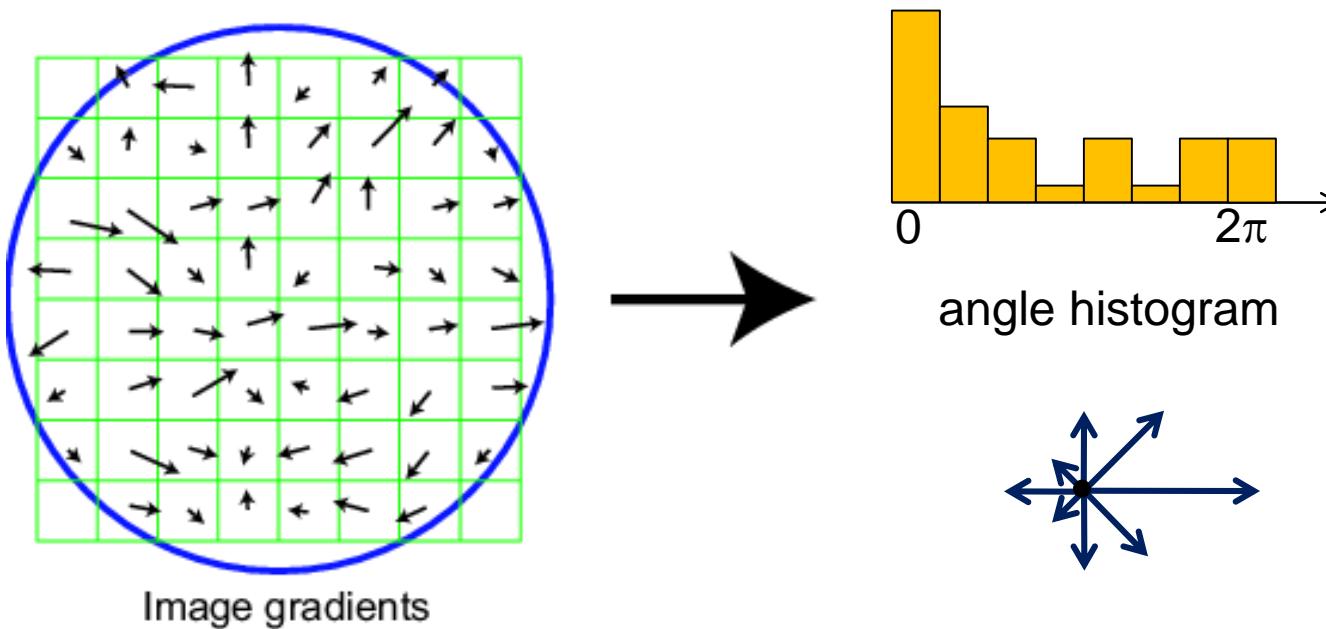
- The ideal descriptor should be
  - Robust
  - Distinctive
  - Compact
  - Efficient
- Most available descriptors focus on edge/gradient information
  - Capture texture information
  - Color rarely used



# Scale Invariant Feature Transform (SIFT)

Basic idea:

- Take 16x16 square window around detected feature
- Compute edge orientation (angle of the gradient - 90°) for each pixel
- Throw out weak edges (threshold gradient magnitude)
- Create histogram of surviving edge orientations

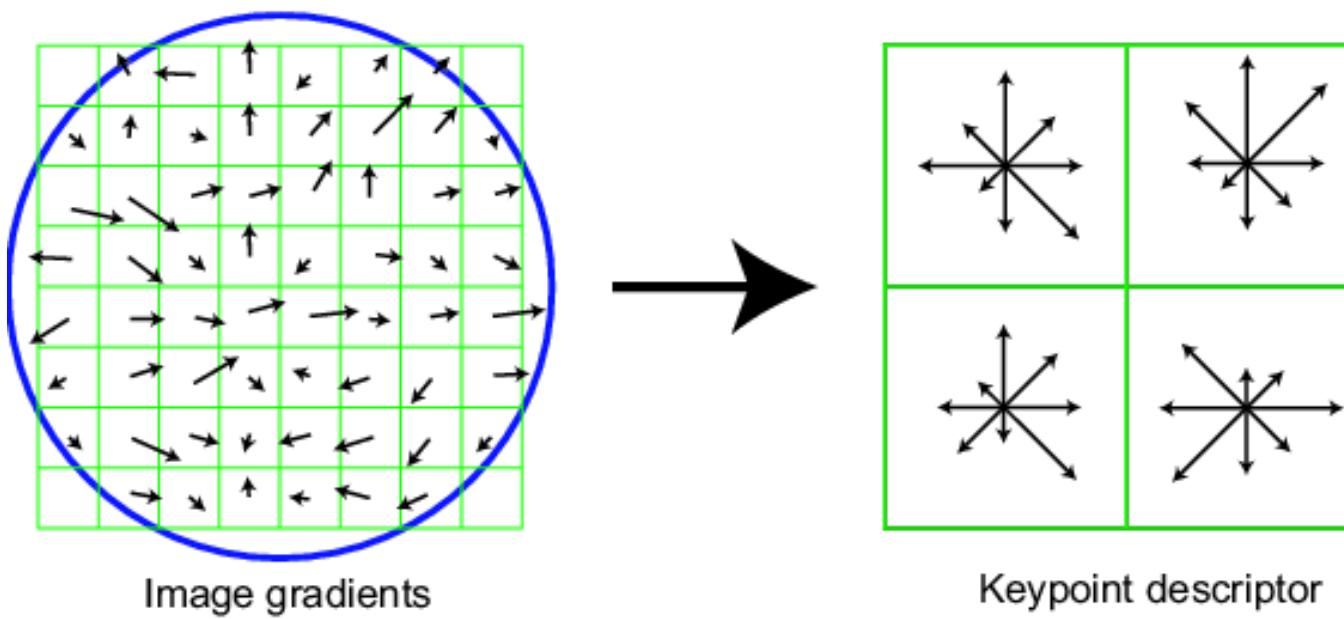


Adapted from slide by David Lowe

# SIFT Descriptor

Full version

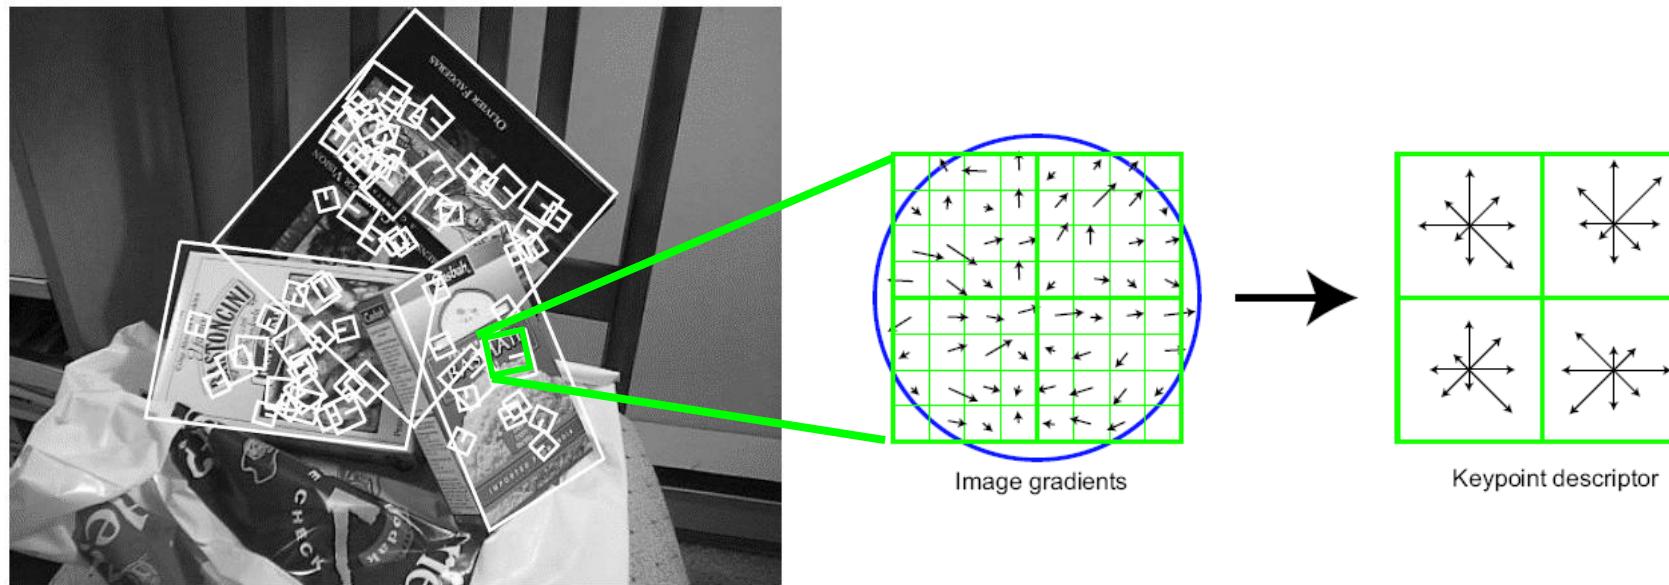
- Divide the  $16 \times 16$  window into a  $4 \times 4$  grid of cells (2x2 case shown below)
- Compute an orientation histogram for each cell
- $16 \text{ cells} * 8 \text{ orientations} = 128 \text{ dimensional descriptor}$



Adapted from slide by David Lowe

Dr. Sander Ali Khowaja

# Local Descriptors: SIFT Descriptor



## Histogram of oriented gradients

- Captures important texture information
- Robust to small translations / affine deformations

[Lowe, ICCV 1999]

K. Grauman, B. Leibe

# Details of Lowe's SIFT Algorithm

- Run DoG detector
  - Find maxima in location/scale space
  - Remove edge points
- Find all major orientations
  - Bin orientations into 36 bin histogram
    - Weight by gradient magnitude
    - Weight by distance to center (Gaussian-weighted mean)
  - Return orientations within 0.8 of peak
    - Use parabola for better orientation fit
- For each (x,y,scale,orientation), create descriptor:
  - Sample 16x16 gradient mag. and rel. orientation
  - Bin 4x4 samples into 4x4 histograms
  - Threshold values to max of 0.2, divide by L2 norm
  - Final descriptor: 4x4x8 normalized histograms

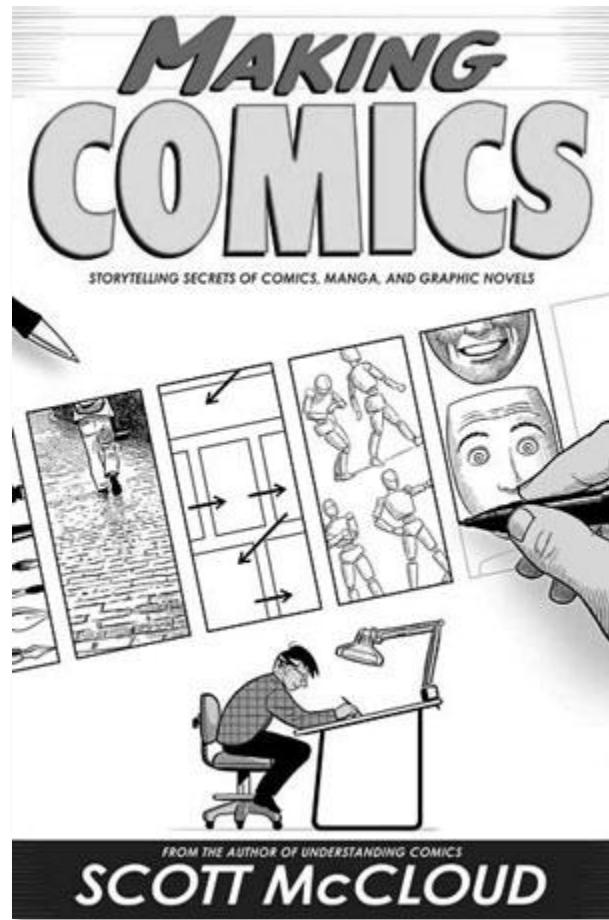
$$\mathbf{H} = \begin{bmatrix} D_{xx} & D_{xy} \\ D_{xy} & D_{yy} \end{bmatrix}$$

$$\frac{\text{Tr}(\mathbf{H})^2}{\text{Det}(\mathbf{H})} < \frac{(r+1)^2}{r}$$

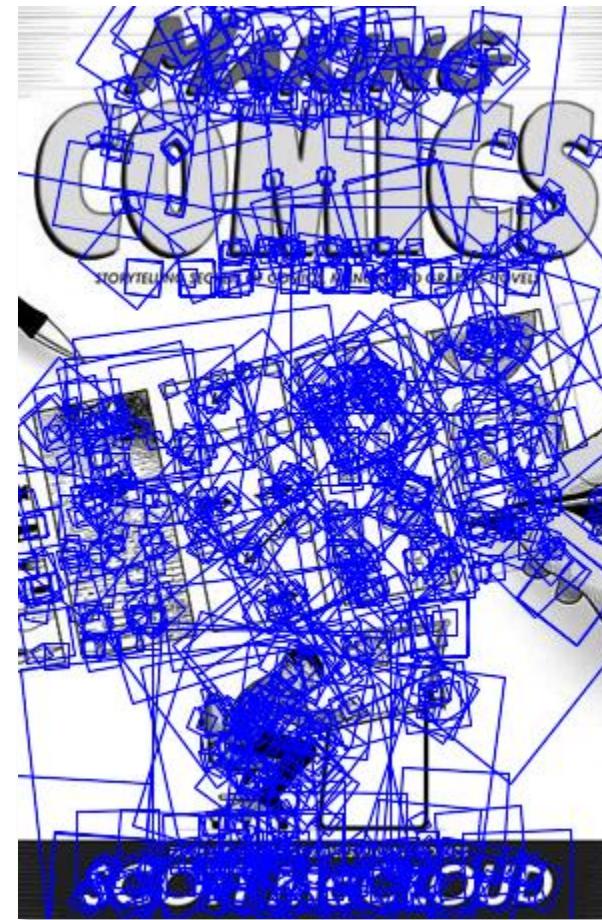
Lowe IJCV 2004



# SIFT Example



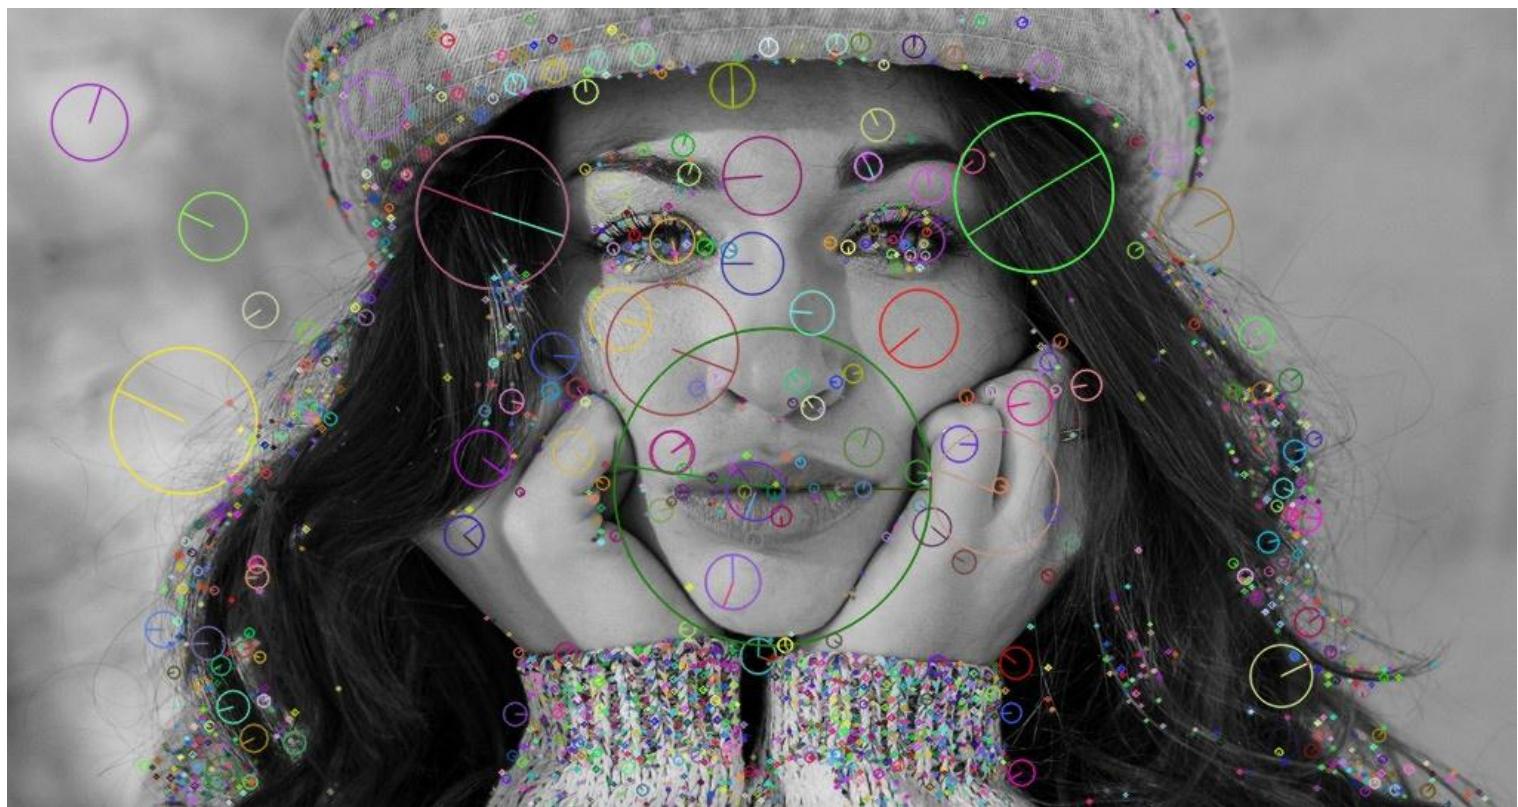
sift



868 SIFT features

# SIFT Python

```
1 # Important NOTE: Use opencv <= 3.4.2.16 as
2 # SIFT is no longer available in
3 # opencv > 3.4.2.16
4 import cv2
5
6 # Loading the image
7 img = cv2.imread('8.jpg')
8
9 # Converting image to grayscale
10 gray= cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
11
12 # Applying SIFT detector
13 sift = cv2.xfeatures2d.SIFT_create()
14 kp = sift.detect(gray, None)
15
16 # Marking the keypoint on the image using circles
17 img=cv2.drawKeypoints(gray ,
18 kp ,
19 img ,
20 flags=cv2.DRAW_MATCHES_FLAGS_DRAW_RICH_KEYPOINTS)
21
22 cv2.imwrite('image-with-keypoints.jpg', img)
```



# Feature Matching

Given a feature in  $I_1$ , how to find the best match in  $I_2$ ?

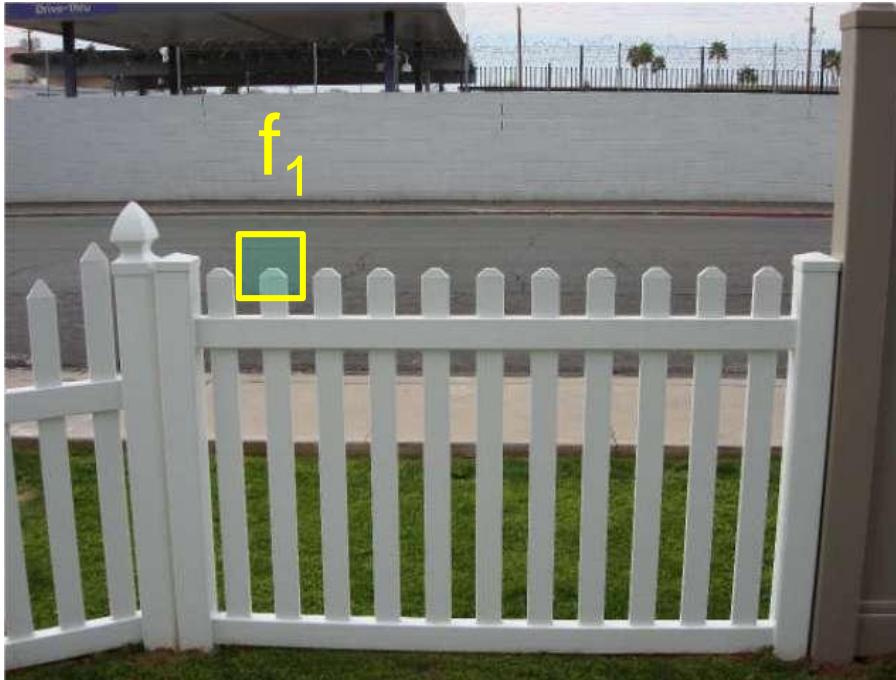
1. Define distance function that compares two descriptors
2. Test all the features in  $I_2$ , find the one with min distance



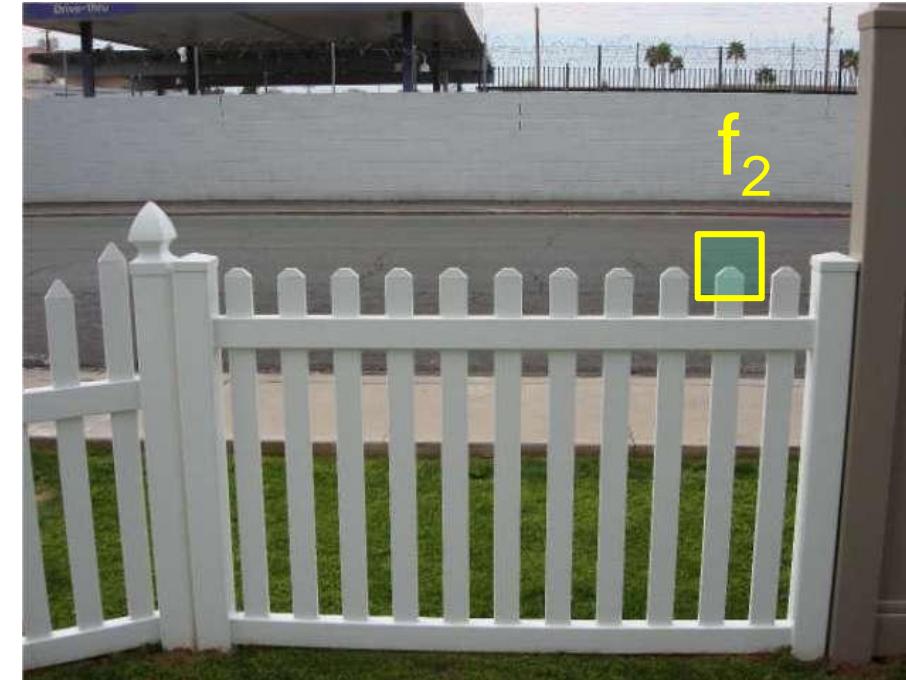
# Feature Distance

How to define the difference between two features  $f_1, f_2$ ?

- Simple approach: L<sub>2</sub> distance,  $\|f_1 - f_2\|$
- can give good scores to ambiguous (incorrect) matches



$I_1$

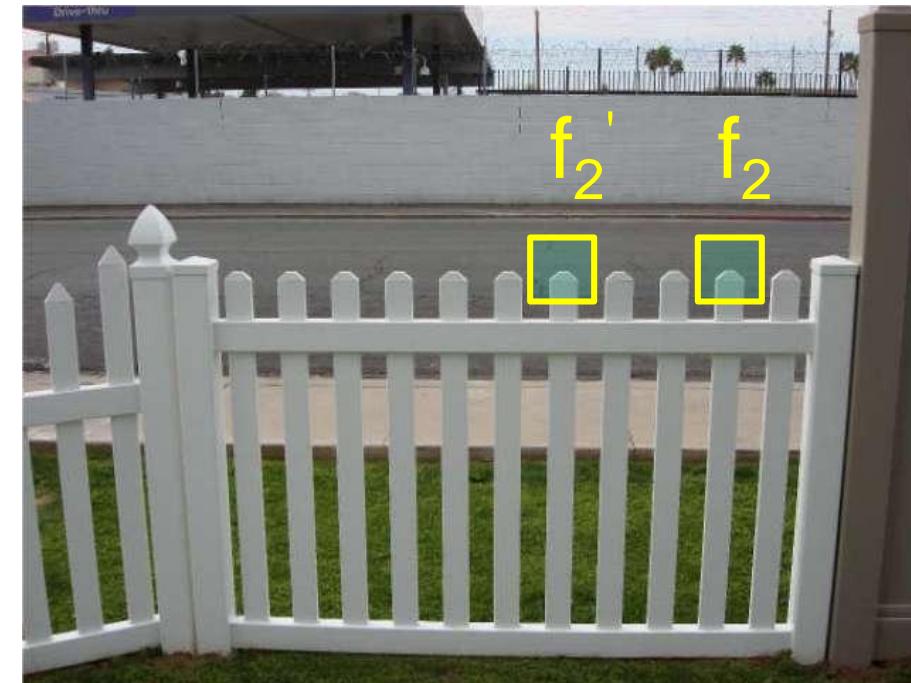
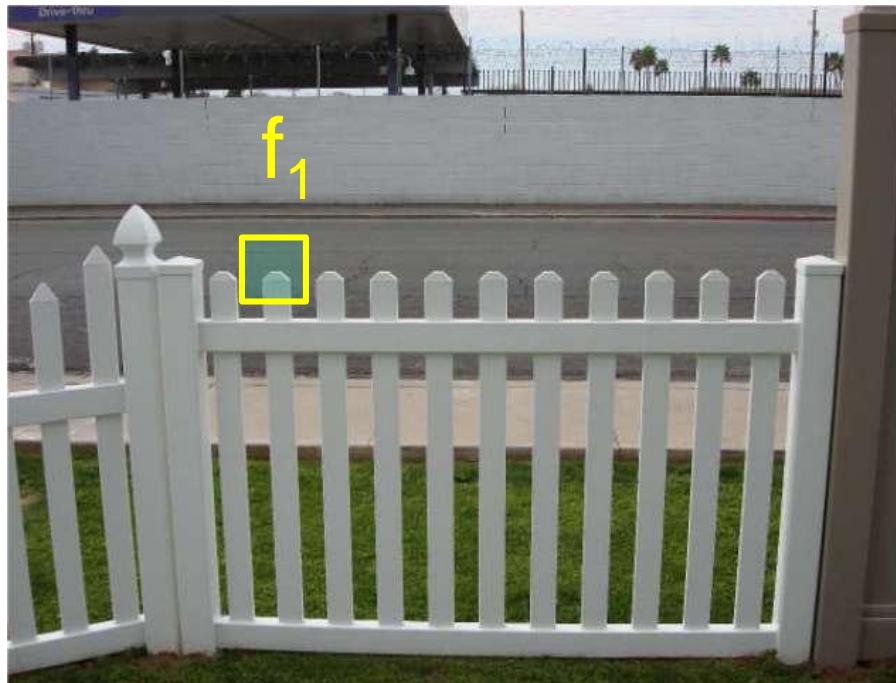


$I_2$

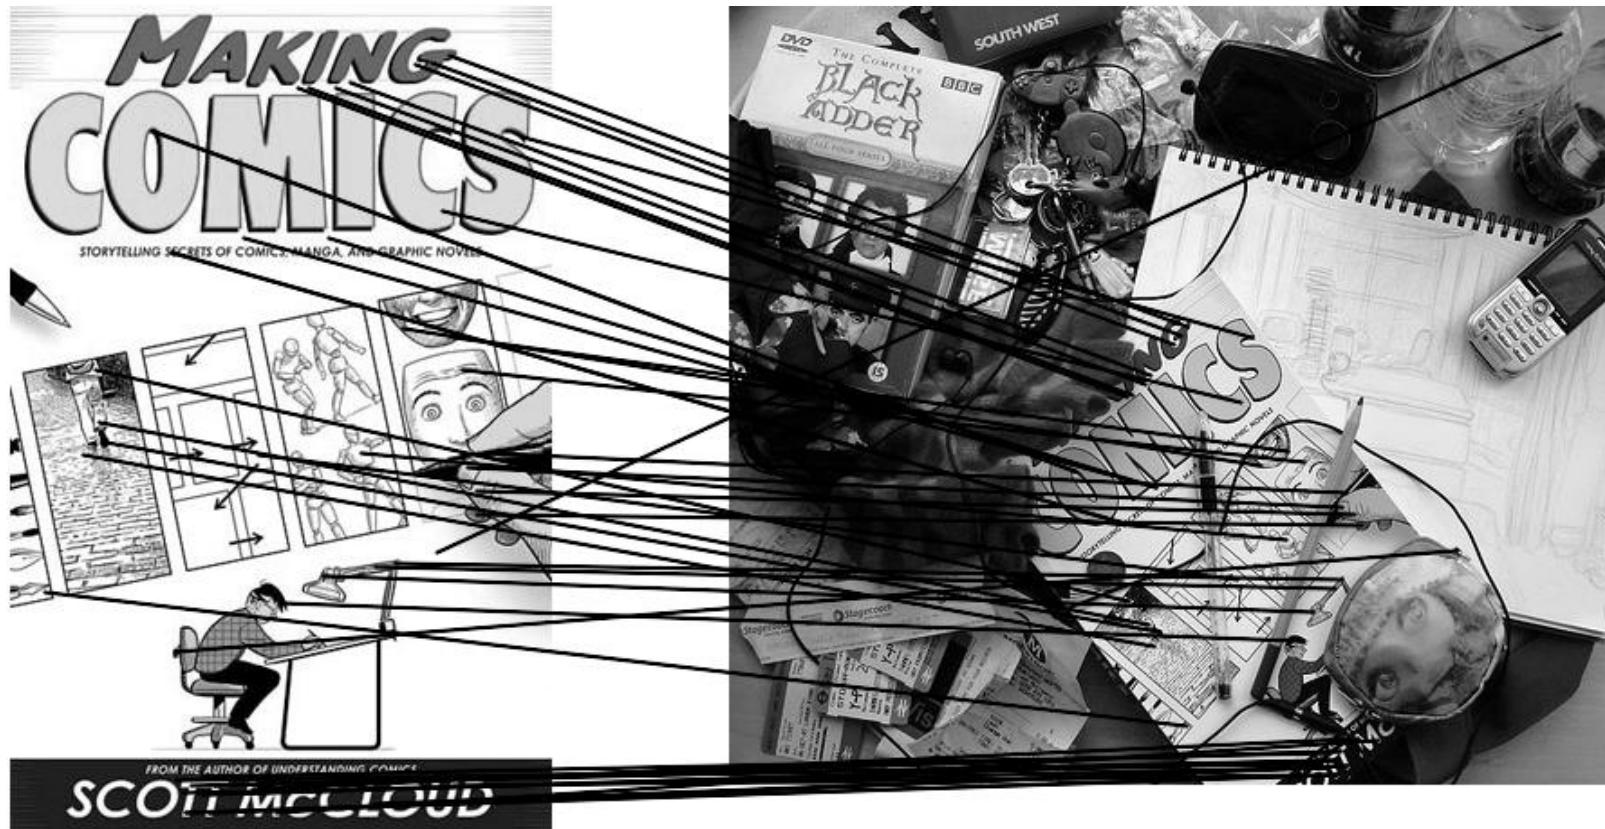
# Feature Distance

How to define the difference between two features  $f_1, f_2$ ?

- Better approach: ratio distance =  $\|f_1 - f_2\| / \|f_1 - f_2'\|$ 
  - $f_2$  is best SSD match to  $f_1$  in  $I_2$
  - $f_2'$  is 2<sup>nd</sup> best SSD match to  $f_1$  in  $I_2$
  - gives large values for ambiguous matches

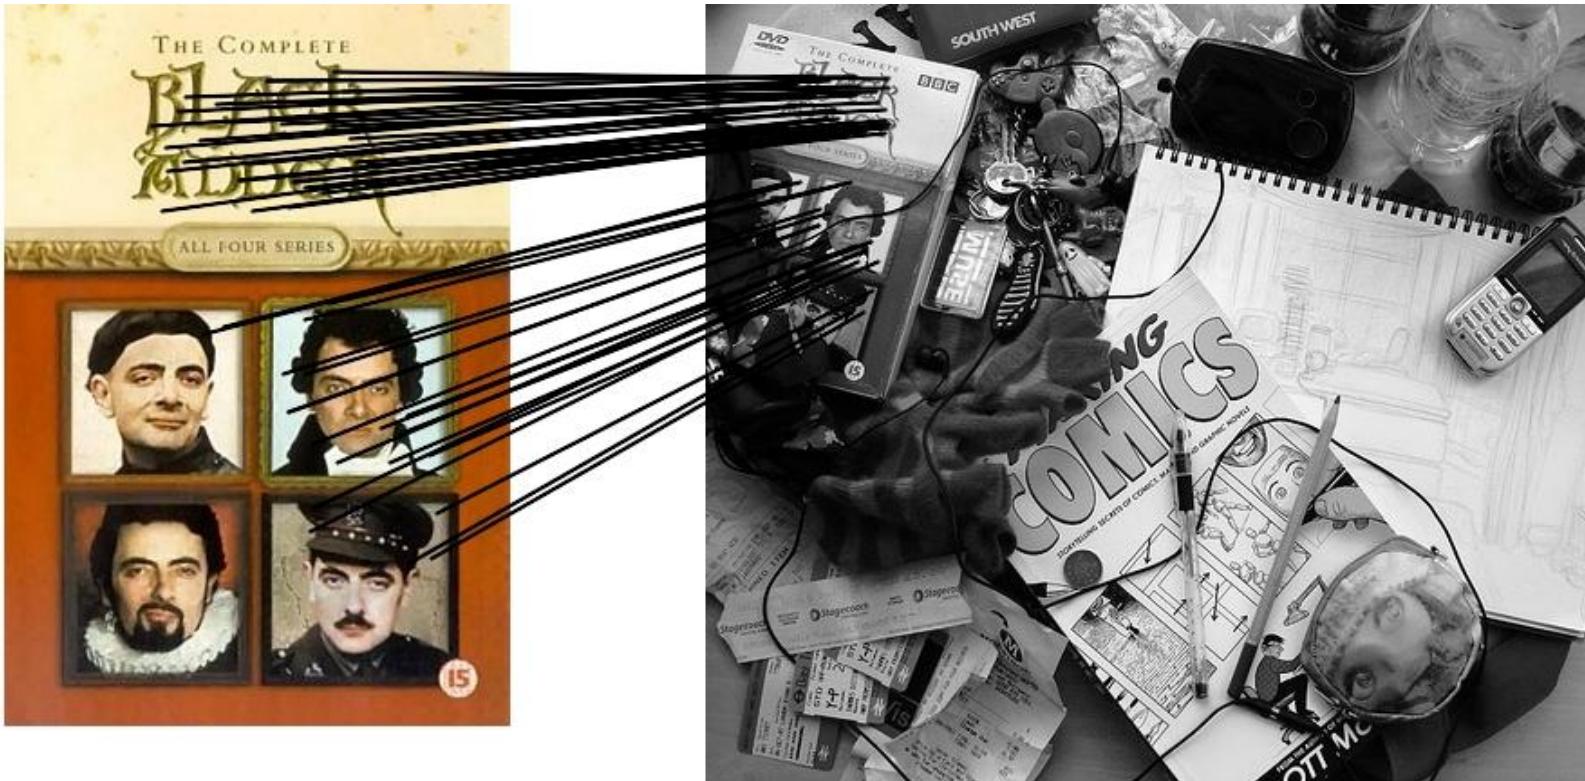


# Feature Matching Example



51 matches

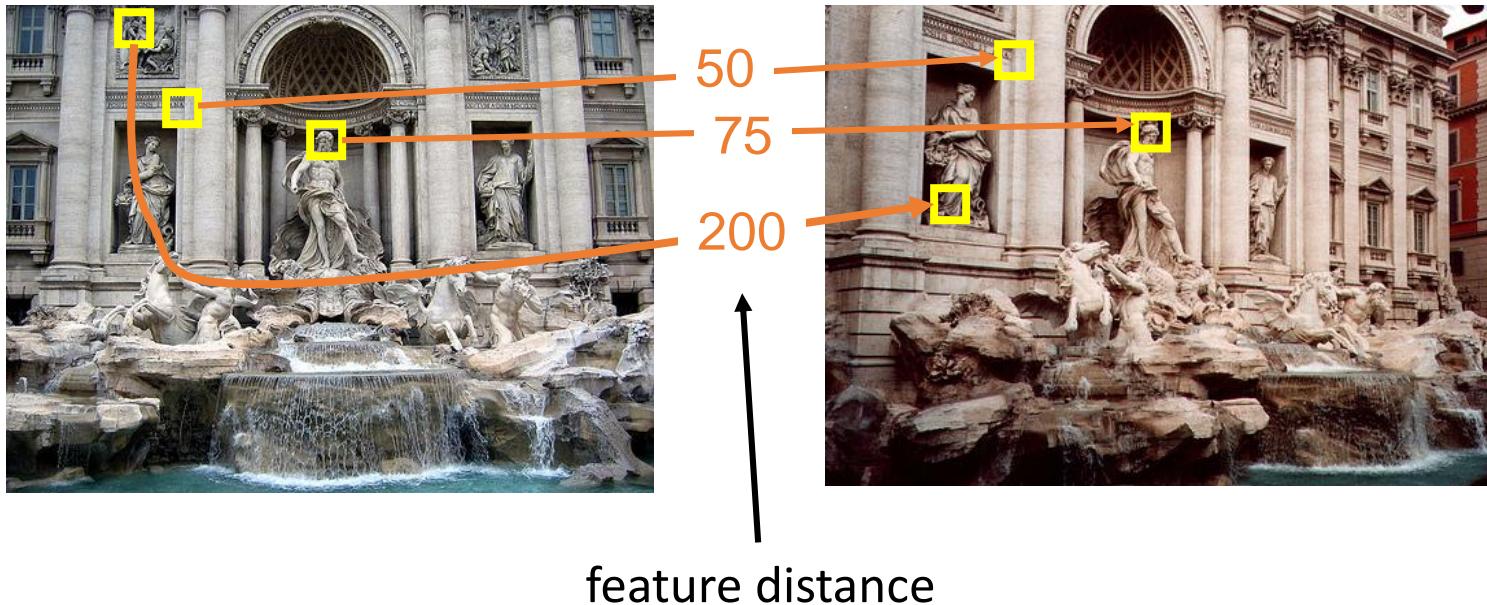
# Feature Matching Example



58 matches

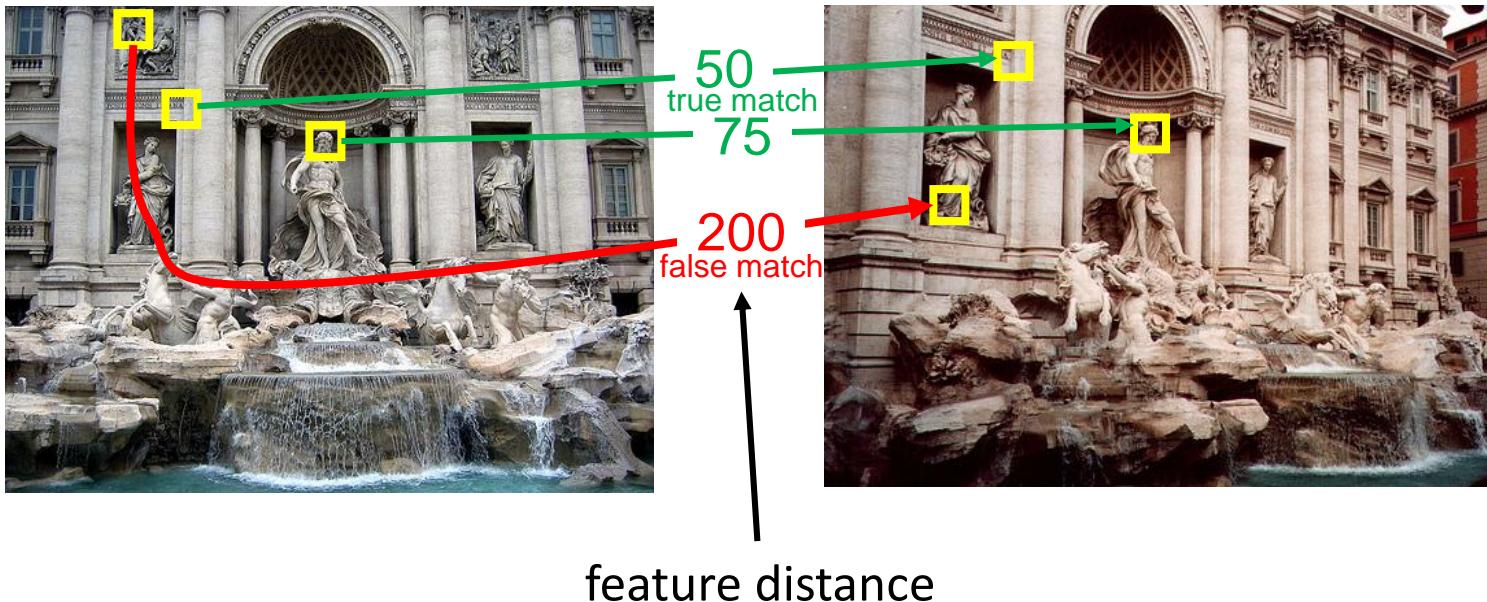
# Evaluating the results

How can we measure the performance of a feature matcher?



# True/False Positives

How can we measure the performance of a feature matcher?

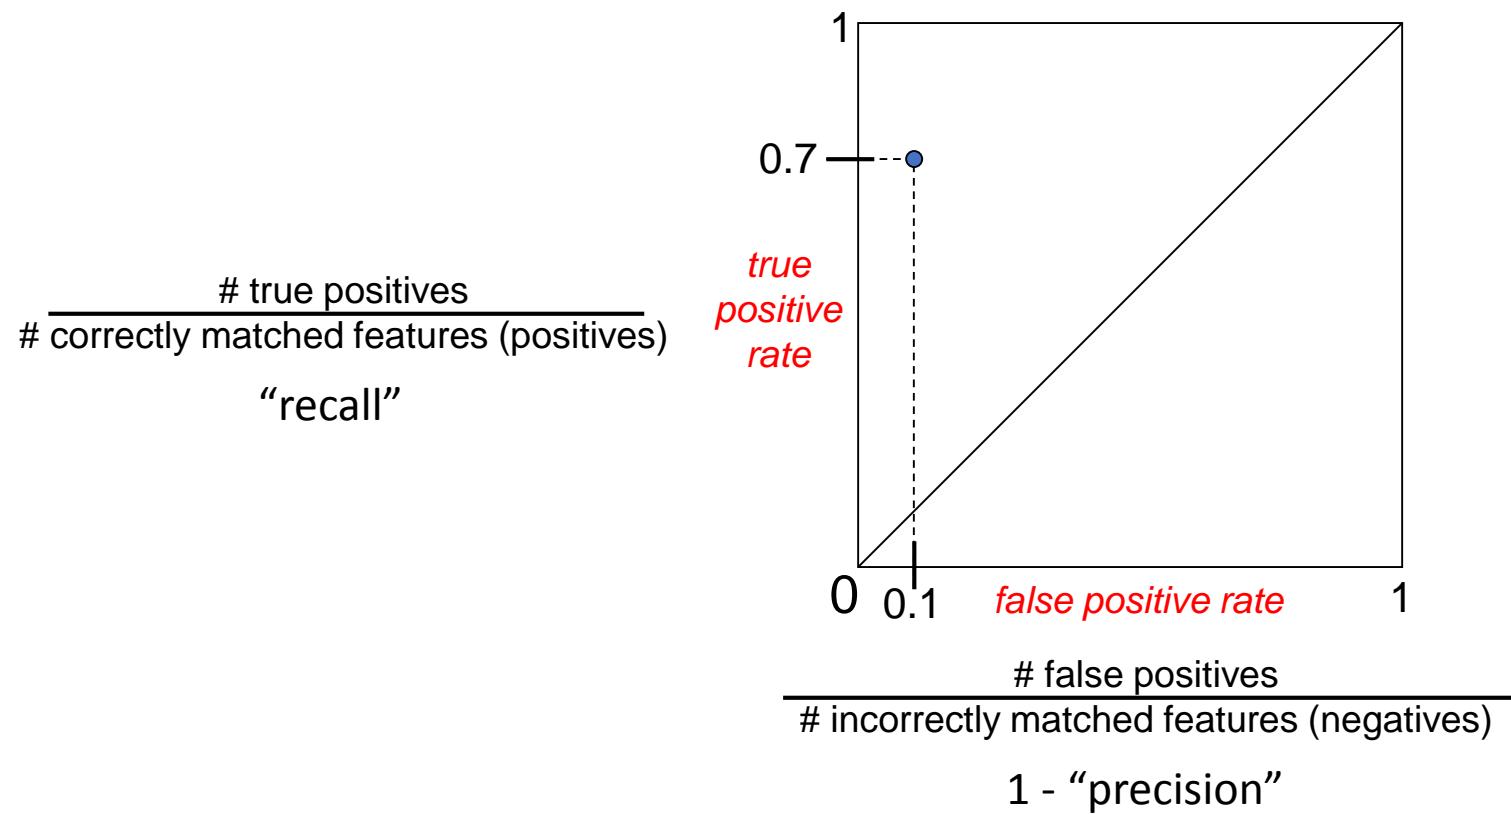


The distance threshold affects performance

- True positives = # of detected matches that are correct
  - Suppose we want to maximize these—how to choose threshold?
- False positives = # of detected matches that are incorrect
  - Suppose we want to minimize these—how to choose threshold?

# Evaluating the results

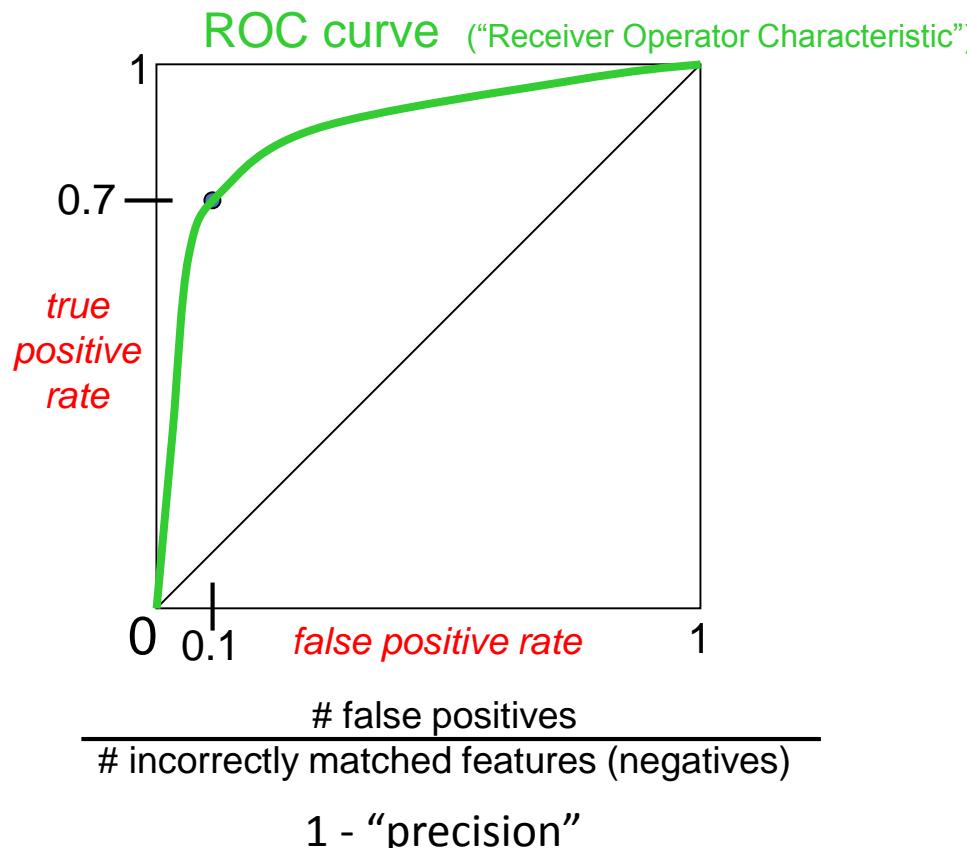
How can we measure the performance of a feature matcher?



# Evaluating the results

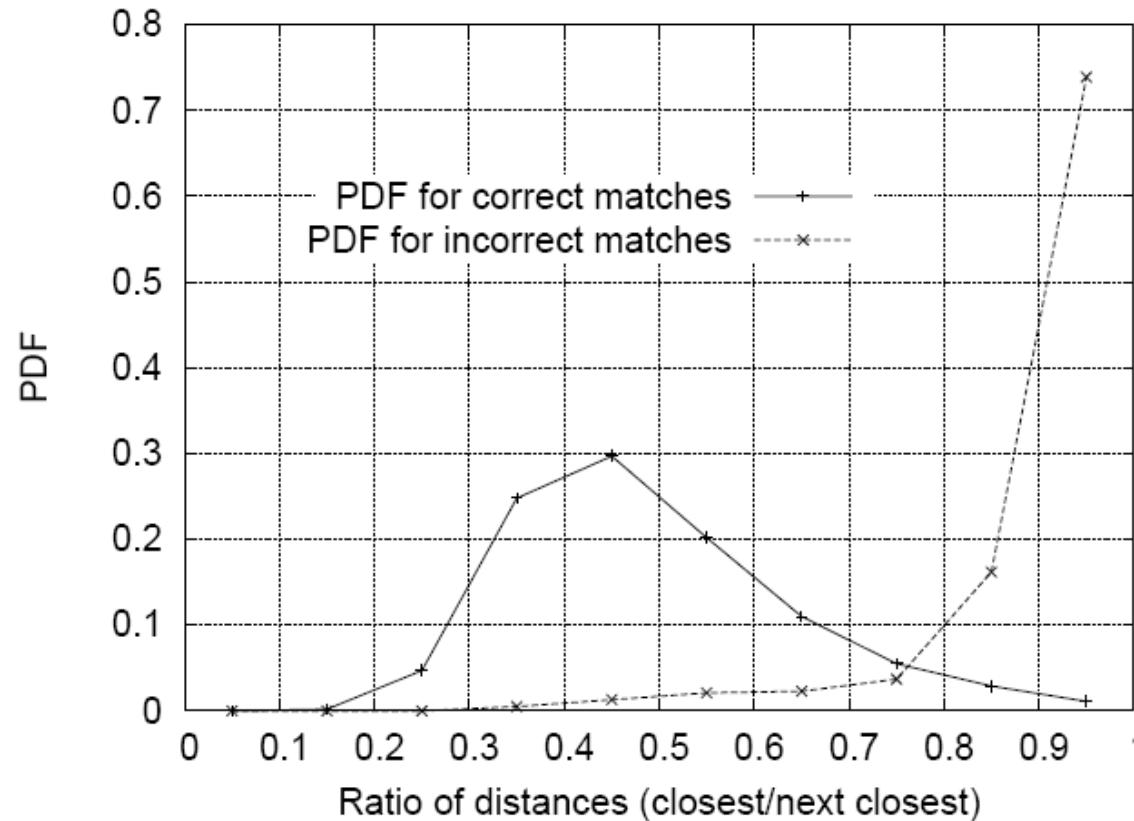
How can we measure the performance of a feature matcher?

$$\frac{\# \text{ true positives}}{\# \text{ correctly matched features (positives)}} \quad \text{"recall"}$$



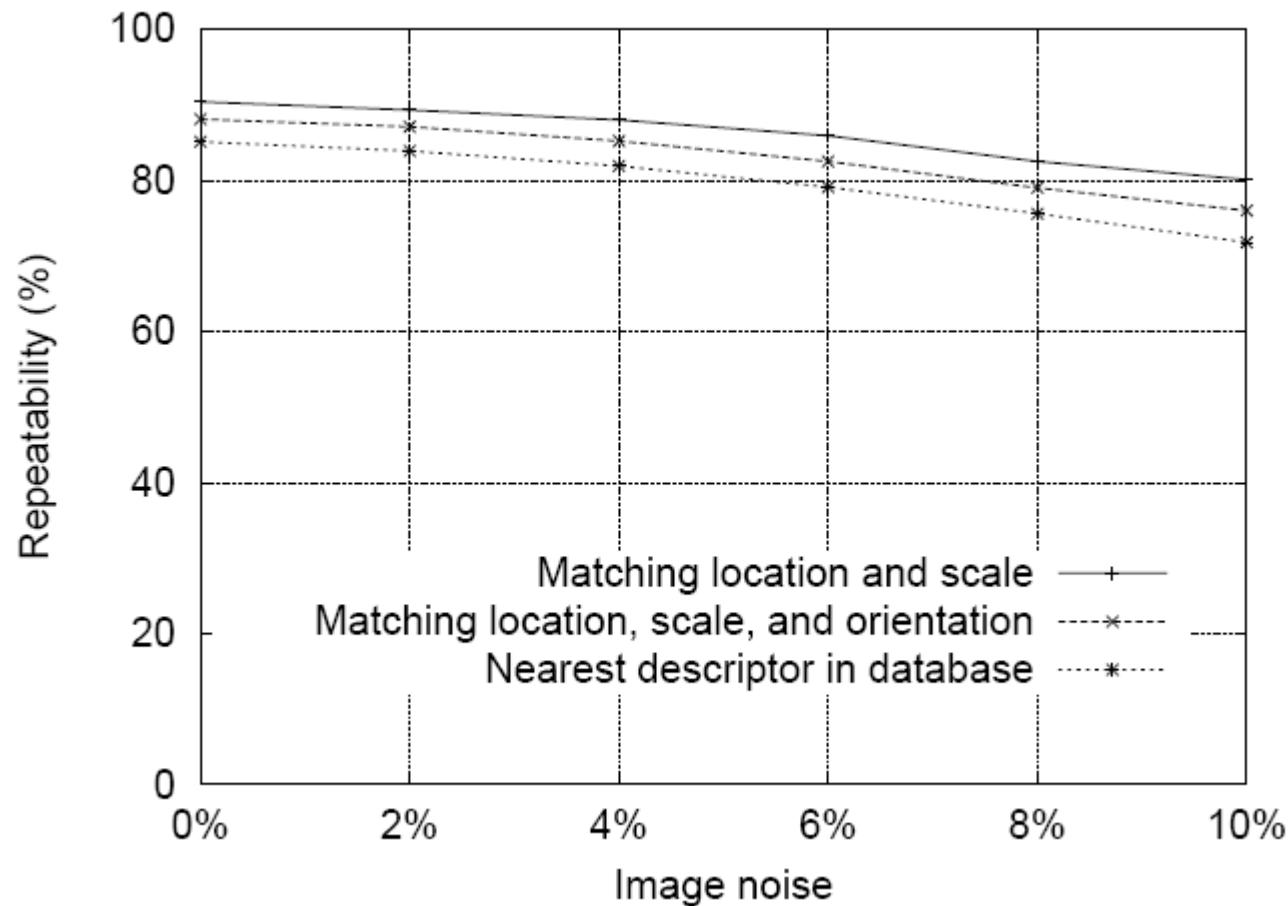
# Matching SIFT Descriptors

- Nearest neighbor (Euclidean distance)
- Threshold ratio of nearest to 2<sup>nd</sup> nearest descriptor



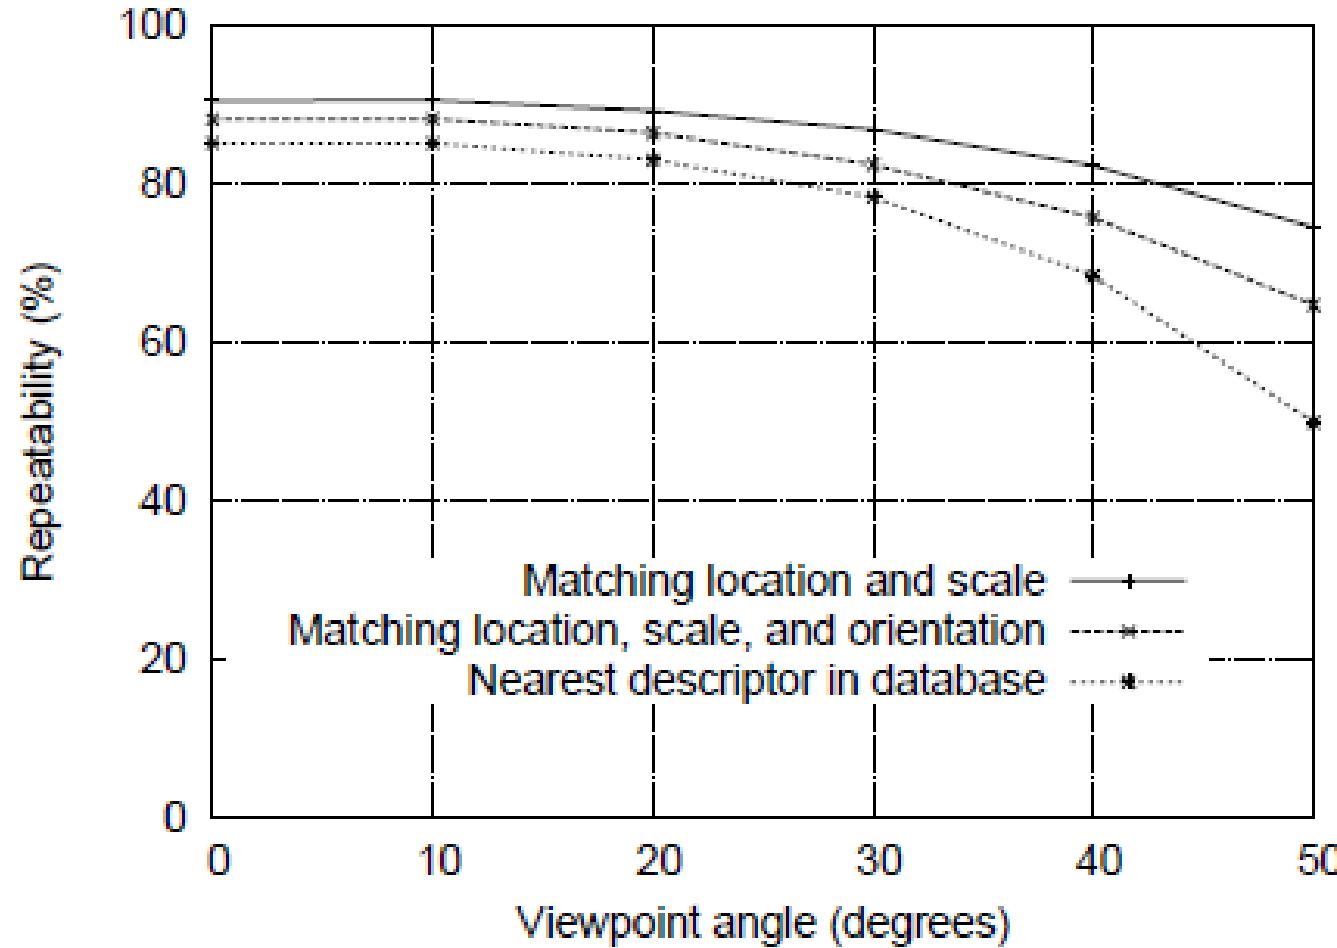
Lowe IJCV 2004

# SIFT Repeatability

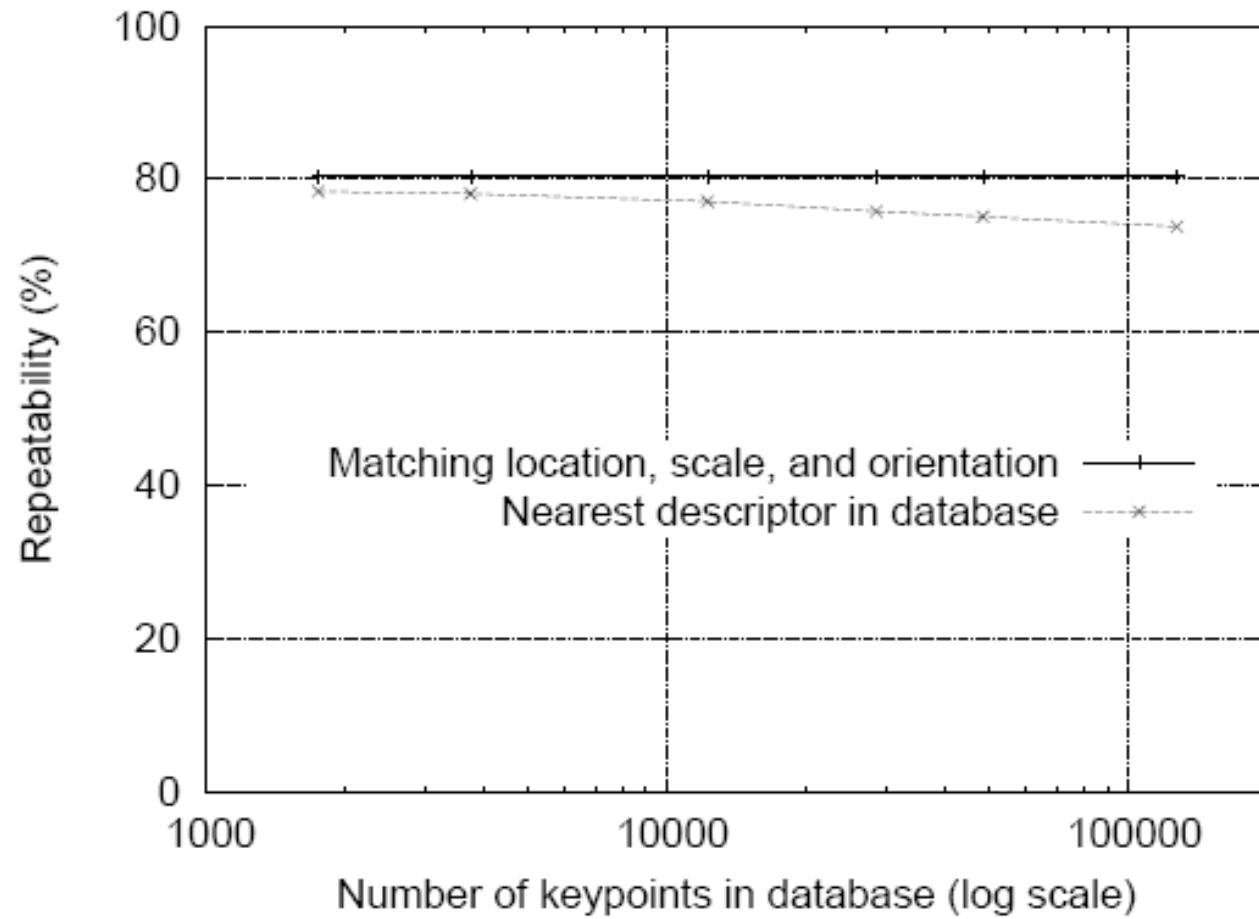


Lowe IJCV 2004

# SIFT Repeatability



# SIFT Repeatability

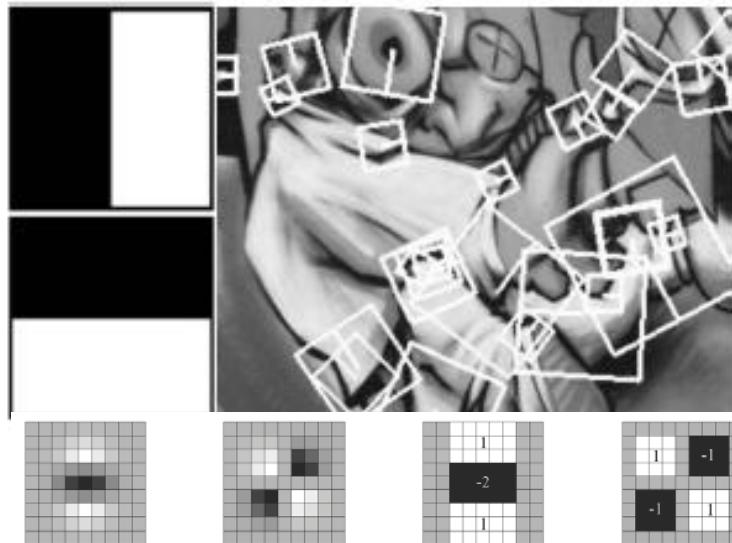


Lowe IJCV 2004

# SIFT Feature Matching Python

```
1 # import required libraries
2 import numpy as np
3 import cv2
4 import matplotlib.pyplot as plt
5
6 # read two input images as grayscale
7 img1 = cv2.imread('0.jpg',0)
8 img2 = cv2.imread('65.jpg',0)
9
10 # Initiate SIFT detector
11 sift = cv2.SIFT_create()
12
13 # detect and compute the keypoints and descriptors with SIFT
14 kp1, des1 = sift.detectAndCompute(img1,None)
15 kp2, des2 = sift.detectAndCompute(img2,None)
16
17 # create BFMatcher object
18 bf = cv2.BFMatcher()
19
20 # Match descriptors.
21 matches = bf.match(des1,des2)
22
23 # sort the matches based on distance
24 matches = sorted(matches, key=lambda val: val.distance)
25
26 # Draw first 50 matches.
27 out = cv2.drawMatches(img1, kp1, img2, kp2, matches[:50], None, flags=2)
28 plt.imshow(out), plt.show()
```





Many other efficient descriptors  
are also available

[Bay, ECCV'06], [Cornelis, CVGPU'08]

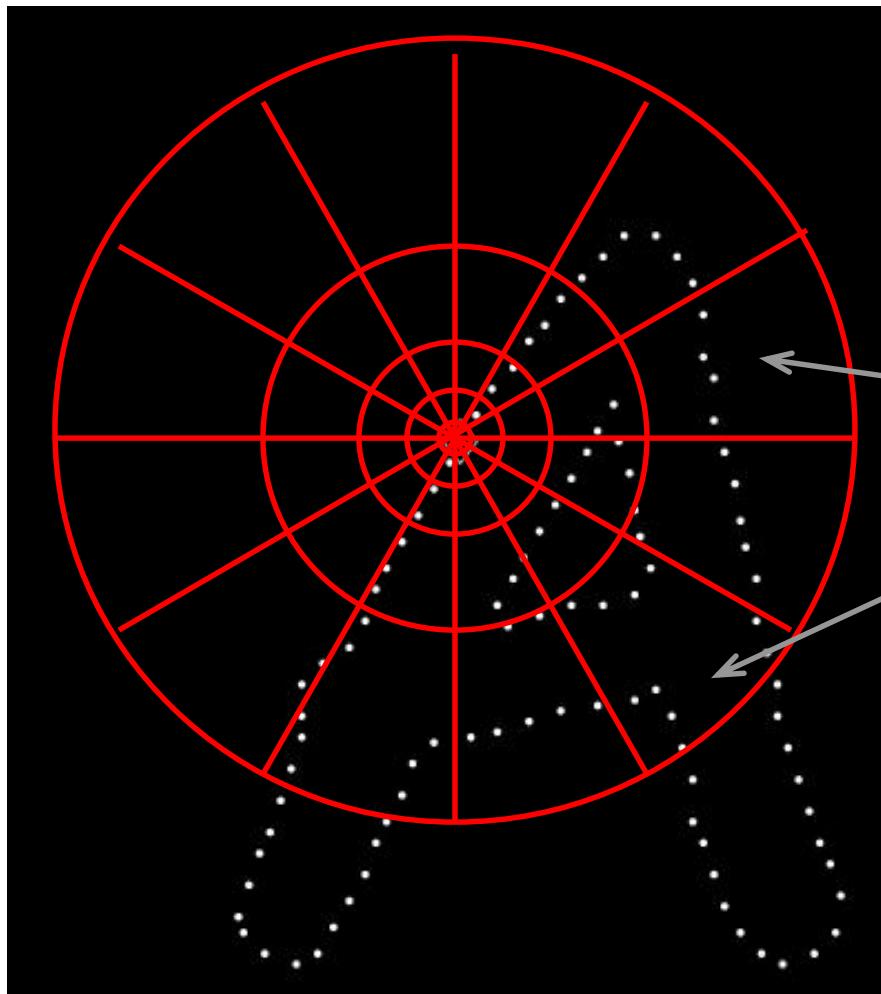
## Fast approximation of SIFT idea

Efficient computation by 2D box filters &  
integral images  
⇒ 6 times faster than SIFT  
Equivalent quality for object  
identification

## GPU implementation available

Feature extraction @ 200Hz  
(detector + descriptor, 640×480 img)  
<http://www.vision.ee.ethz.ch/~surf>

# Local Descriptors: Shape Context



Count the number of points inside each bin, e.g.:

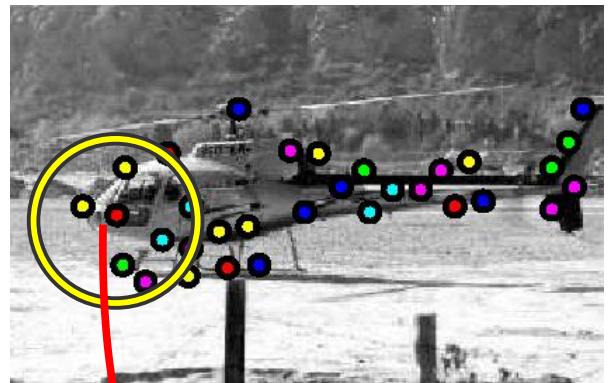
Count = 4

⋮

Count = 10

Log-polar binning: more precision for nearby points, more flexibility for farther points.

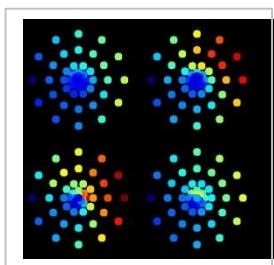
# Local Descriptors: Geometric Blur



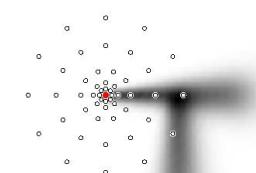
Compute edges

at four  
orientations

Extract a patch  
in each channel



Example descriptor



(Idealized signal)

Apply spatially varying  
blur and sub-sample

Berg & Malik, CVPR 2001

K. Grauman, B. Leibe

# Choosing a Detector

- What do you want it for?
  - Precise localization in x-y: Harris
  - Good localization in scale: Difference of Gaussian
  - Flexible region shape: MSER
- Best choice often application dependent
  - Harris-/Hessian-Laplace/DoG work well for many natural categories
  - MSER works well for buildings and printed things
- Why choose?
  - Get more points with more detectors
- There have been extensive evaluations/comparisons
  - [Mikolajczyk et al., IJCV'05, PAMI'05]
  - All detectors/descriptors shown here work well



# Comparison of Key Point Detectors

Table 7.1 Overview of feature detectors.

Feature Detector	Corner	Blob	Region	Rotation invariant	Scale invariant	Affine invariant	Repeatability	Localization accuracy	Robustness	Efficiency
Harris	√			√			+++	+++	+++	++
Hessian		√		√			++	++	++	+
SUSAN	√			√			++	++	++	+++
Harris-Laplace	√	(√)		√	√		+++	+++	++	+
Hessian-Laplace	(√)	√		√	√		+++	+++	+++	+
DoG	(√)	√		√	√		++	++	++	++
SURF	(√)	√		√	√		++	++	++	+++
Harris-Affine	√	(√)		√	√	√	+++	+++	++	++
Hessian-Affine	(√)	√		√	√	√	+++	+++	+++	++
Salient Regions	(√)	√		√	√	(√)	+	+	++	+
Edge-based	√			√	√	√	+++	+++	+	+
MSER		√		√	√	√	+++	+++	++	+++
Intensity-based		√		√	√	√	++	++	++	++
Superpixels		√		√	(√)	(√)	+	+	+	+

Tuytelaars Mikolajczyk 2008

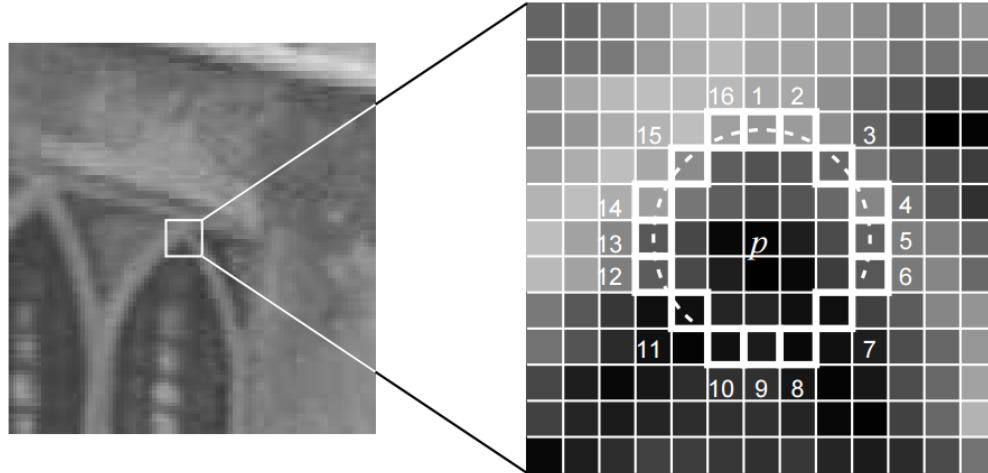


# Choosing a Descriptor

- Again, need not stick to one
- For object instance recognition or stitching, SIFT or variant is a good choice



# Recent advances in Interest Points



[Features from Accelerated Segment Test](#), ECCV 06

## Binary feature descriptors

- [BRIEF: Binary Robust Independent Elementary Features](#), ECCV 10
- [ORB \(Oriented FAST and Rotated BRIEF\)](#), CVPR 11
- [BRISK: Binary robust invariant scalable keypoints](#), ICCV 11
- [Freak: Fast retina keypoint](#), CVPR 12
- [LIFT: Learned Invariant Feature Transform](#), ECCV 16

# Things to Remember

- Keypoint detection: repeatable and distinctive
  - Corners, blobs, stable regions
  - Harris, DoG
- Descriptors: robust and selective
  - spatial histograms of orientation
  - SIFT

